

# Operating Systems

## [ 1. Introduction ]

Chung-Wei Lin

[cwlin@csie.ntu.edu.tw](mailto:cwlin@csie.ntu.edu.tw)

CSIE Department

National Taiwan University

# Objectives

- ❑ Describe the general organization of a computer system and the role of interrupts
- ❑ Describe the components in a modern multiprocessor computer system
- ❑ Illustrate the transition from user mode to kernel mode
- ❑ Discuss how operating systems are used in various computing environments

# Outline

- ❑ **What Operating Systems Do**

- User View, System View, Defining Operating Systems

- ❑ Computer-System Organization

- ❑ Computer-System Architecture

- ❑ Operating-System Operations

- ❑ Resource Management, Security and Protection, Virtualization

- ❑ Distributed Systems

- ❑ Kernel Data Structures

- ❑ Computing Environments

- ❑ Free and Open-Source Operating Systems

# What is an Operating System?

## ❑ An operating system is software

- Manage a computer's hardware
- Provide a basis for application program
- Act as an intermediary between the computer user and the computer hardware

## ❑ Operating systems are everywhere

- Cars
- Home appliances
- Smart phones
- Personal computers
- Enterprise computers
- Cloud computing environments

# Components of Computer Systems

## ❑ Hardware

- Provide the basic computing resources for the system

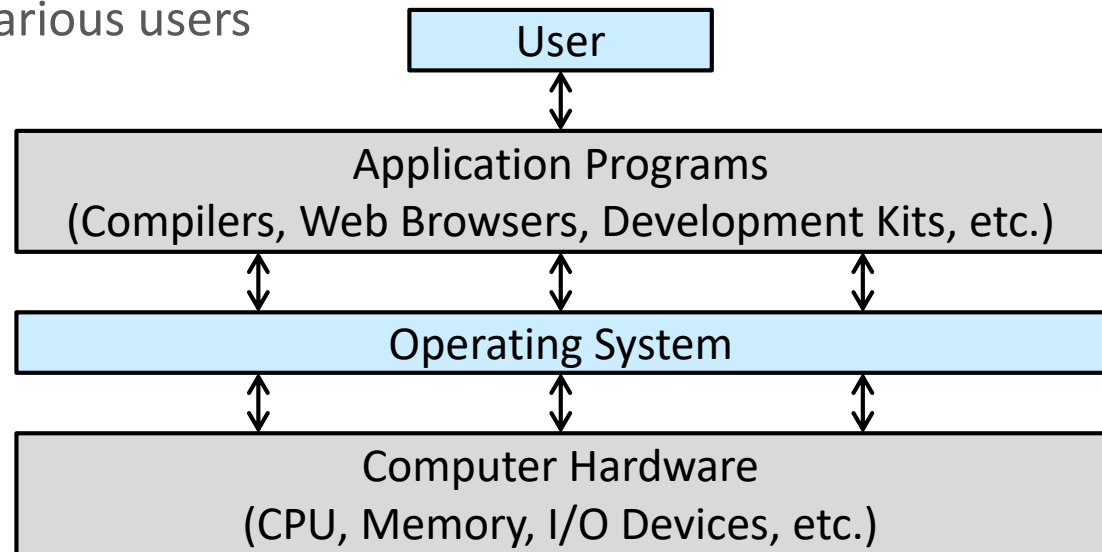
## ❑ Application programs

- Define the ways in which these resources are used to solve users computing problems

## ❑ Operating system

- Control and coordinate the hardware among the various application programs for the various users

## ❑ User



# Outline

## ☐ What Operating Systems Do

➤ User View, System View, and Defining Operating Systems

☐ Computer-System Organization

☐ Computer-System Architecture

☐ Operating-System Operations

☐ Resource Management, Security and Protection, Virtualization

☐ Distributed Systems

☐ Kernel Data Structures

☐ Computing Environments

☐ Free and Open-Source Operating Systems

# User View

## ❑ Many computer users sit with a laptop or in front of a PC

- The operating system is designed mostly for **ease of use**
- Some attention is paid to performance and security
- None (or not too much) is paid to **resource utilization**: how various hardware and software resources are shared

## ❑ Increasing users interact with mobile devices

- Connected to networks through cellular or other wireless technologies
- Featured a touch screen and/or a voice recognition interface

## ❑ Some computers have little or no user view

- They are designed primarily to run without user intervention
- Examples: embedded computers in home devices and automobiles

# Outline

## ☐ What Operating Systems Do

➤ User View, System View, Defining Operating Systems

☐ Computer-System Organization

☐ Computer-System Architecture

☐ Operating-System Operations

☐ Resource Management, Security and Protection, Virtualization

☐ Distributed Systems

☐ Kernel Data Structures

☐ Computing Environments

☐ Free and Open-Source Operating Systems



# System View

## ❑ A resource allocator

- The operating system is the program most intimately involved with the hardware
- The operating system must allocate resources to specific programs and users to make the computer system efficiently and fairly
  - Requests for resources are numerous and possibly conflicting

## ❑ A control program

- A slightly different view emphasizes the need to control the various I/O devices and user programs
- The operating system manages the execution of user programs to prevent errors and improper use of the computer

# Outline

## ☐ What Operating Systems Do

➤ User View, System View, Defining Operating Systems

☐ Computer-System Organization

☐ Computer-System Architecture

☐ Operating-System Operations

☐ Resource Management, Security and Protection, Virtualization

☐ Distributed Systems

☐ Kernel Data Structures

☐ Computing Environments

☐ Free and Open-Source Operating Systems

# Defining Operating Systems (1/3)

- ❑ History of computers
- ❑ No completely adequate definition of an operating system
  - Operating systems exist because they offer a reasonable way to solve the problem of creating a usable computing system
    - Bare hardware alone is not particularly easy to use
    - Application programs require certain common operations
- ❑ No universally accepted definition of what is part of the operating system

# Defining Operating Systems (2/3)

- ❑ A more common definition
- ❑ The one program running at all times on the computer, usually called the kernel
- ❑ Two other types of programs along with the kernel
  - System programs
    - Be associated with the operating system but are not necessarily part of the kernel
  - Application programs
    - Include all programs not associated with the operation of the system

# Defining Operating Systems (3/3)

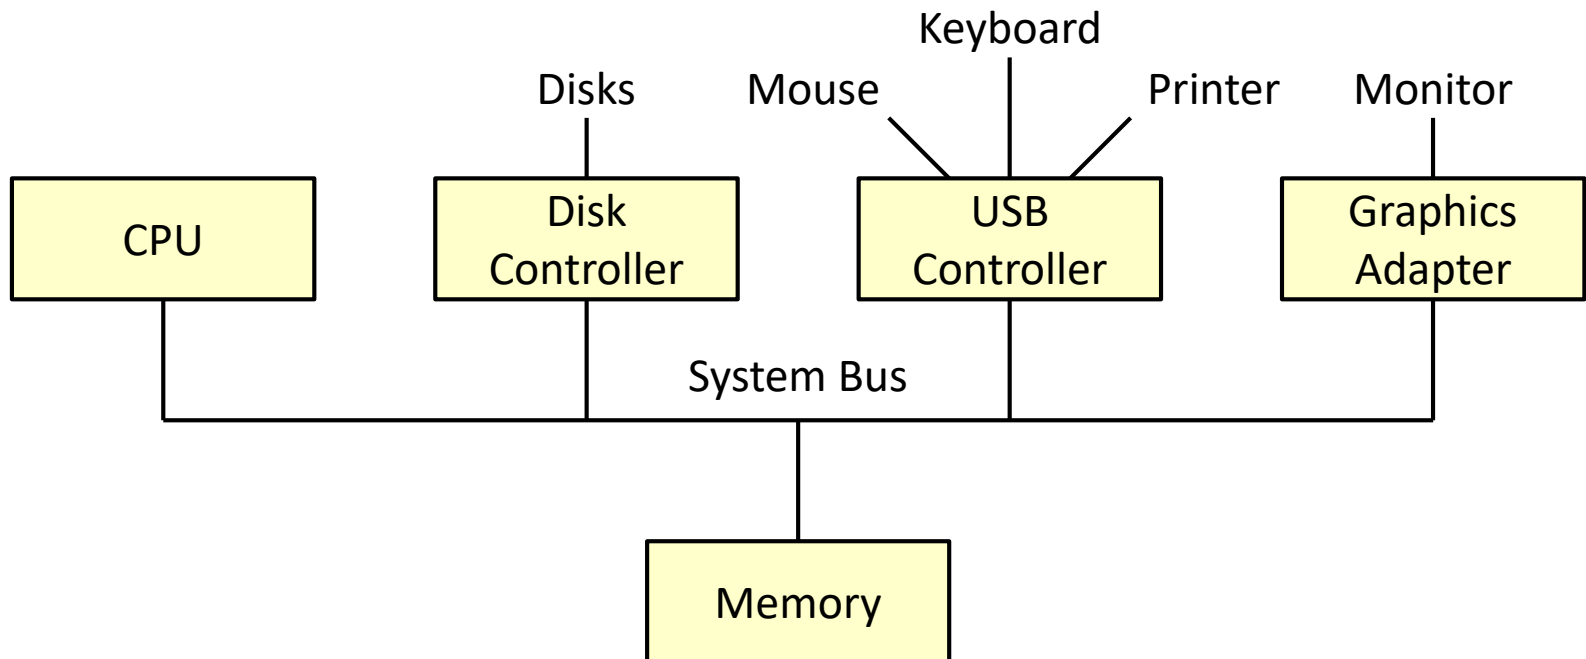
- ❑ Today, mobile operating systems often include not only a core kernel but also middleware
  - A set of software frameworks that provide additional services to application developers
  - Examples: databases, multimedia, graphics

# Outline

- ❑ What Operating Systems Do
- ❑ **Computer-System Organization**
  - Interrupts, Storage Structure, I/O Structure
- ❑ Computer-System Architecture
- ❑ Operating-System Operations
- ❑ Resource Management, Security and Protection, Virtualization
- ❑ Distributed Systems
- ❑ Kernel Data Structures
- ❑ Computing Environments
- ❑ Free and Open-Source Operating Systems

# Computer-System Organization

- ❑ A **bus** connects one or more CPUs and device controllers and provides access between components and shared memory
  - Operating systems have a device driver for each device controller
  - CPUs and device controllers can execute in parallel
  - To ensure orderly access to the shared memory, a memory controller synchronizes access to the memory



# Outline

- ❑ What Operating Systems Do
- ❑ **Computer-System Organization**
  - Interrupts, Storage Structure, I/O Structure
- ❑ Computer-System Architecture
- ❑ Operating-System Operations
- ❑ Resource Management, Security and Protection, Virtualization
- ❑ Distributed Systems
- ❑ Kernel Data Structures
- ❑ Computing Environments
- ❑ Free and Open-Source Operating Systems



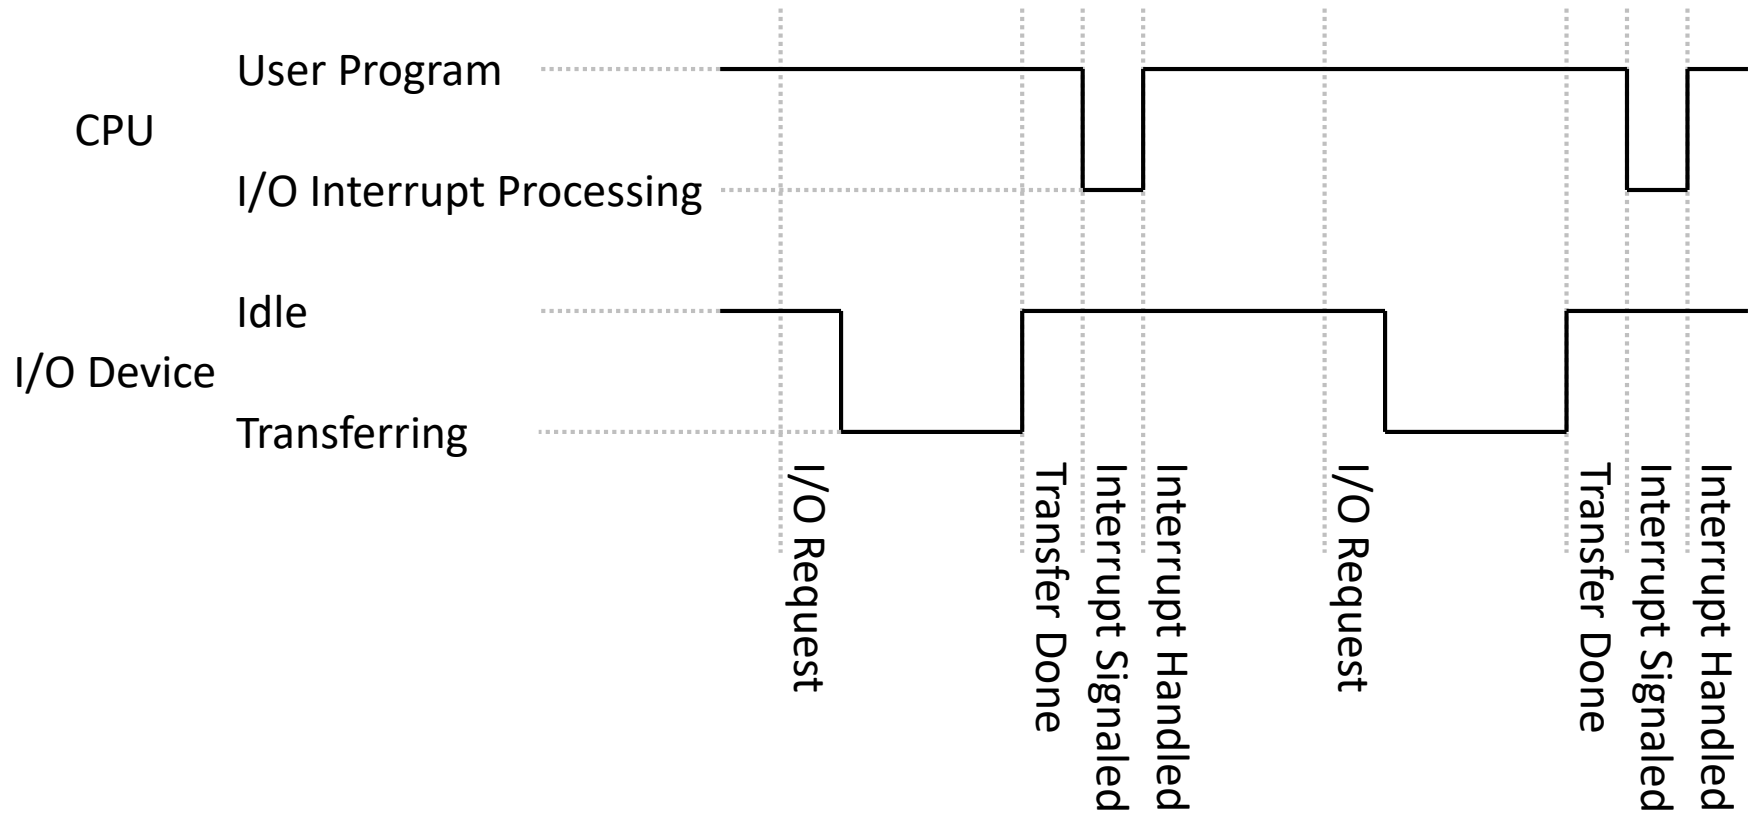
# A Program Performing I/O

- ❑ The device driver loads the appropriate registers in the device controller
- ❑ The device controller
  - Examine the contents of the registers and determine what action to take
    - Example: "read a character from the keyboard"
  - Start the transfer of data from the device to its local buffer
  - Inform the device driver once the transfer of data is complete
- ❑ The device driver gives control to other parts of the operating system
  - Return the data or a pointer to the data (for a read operation)
  - Return status information (for other operations)
- ❑ How does the controller inform the device driver that it has finished its operation?

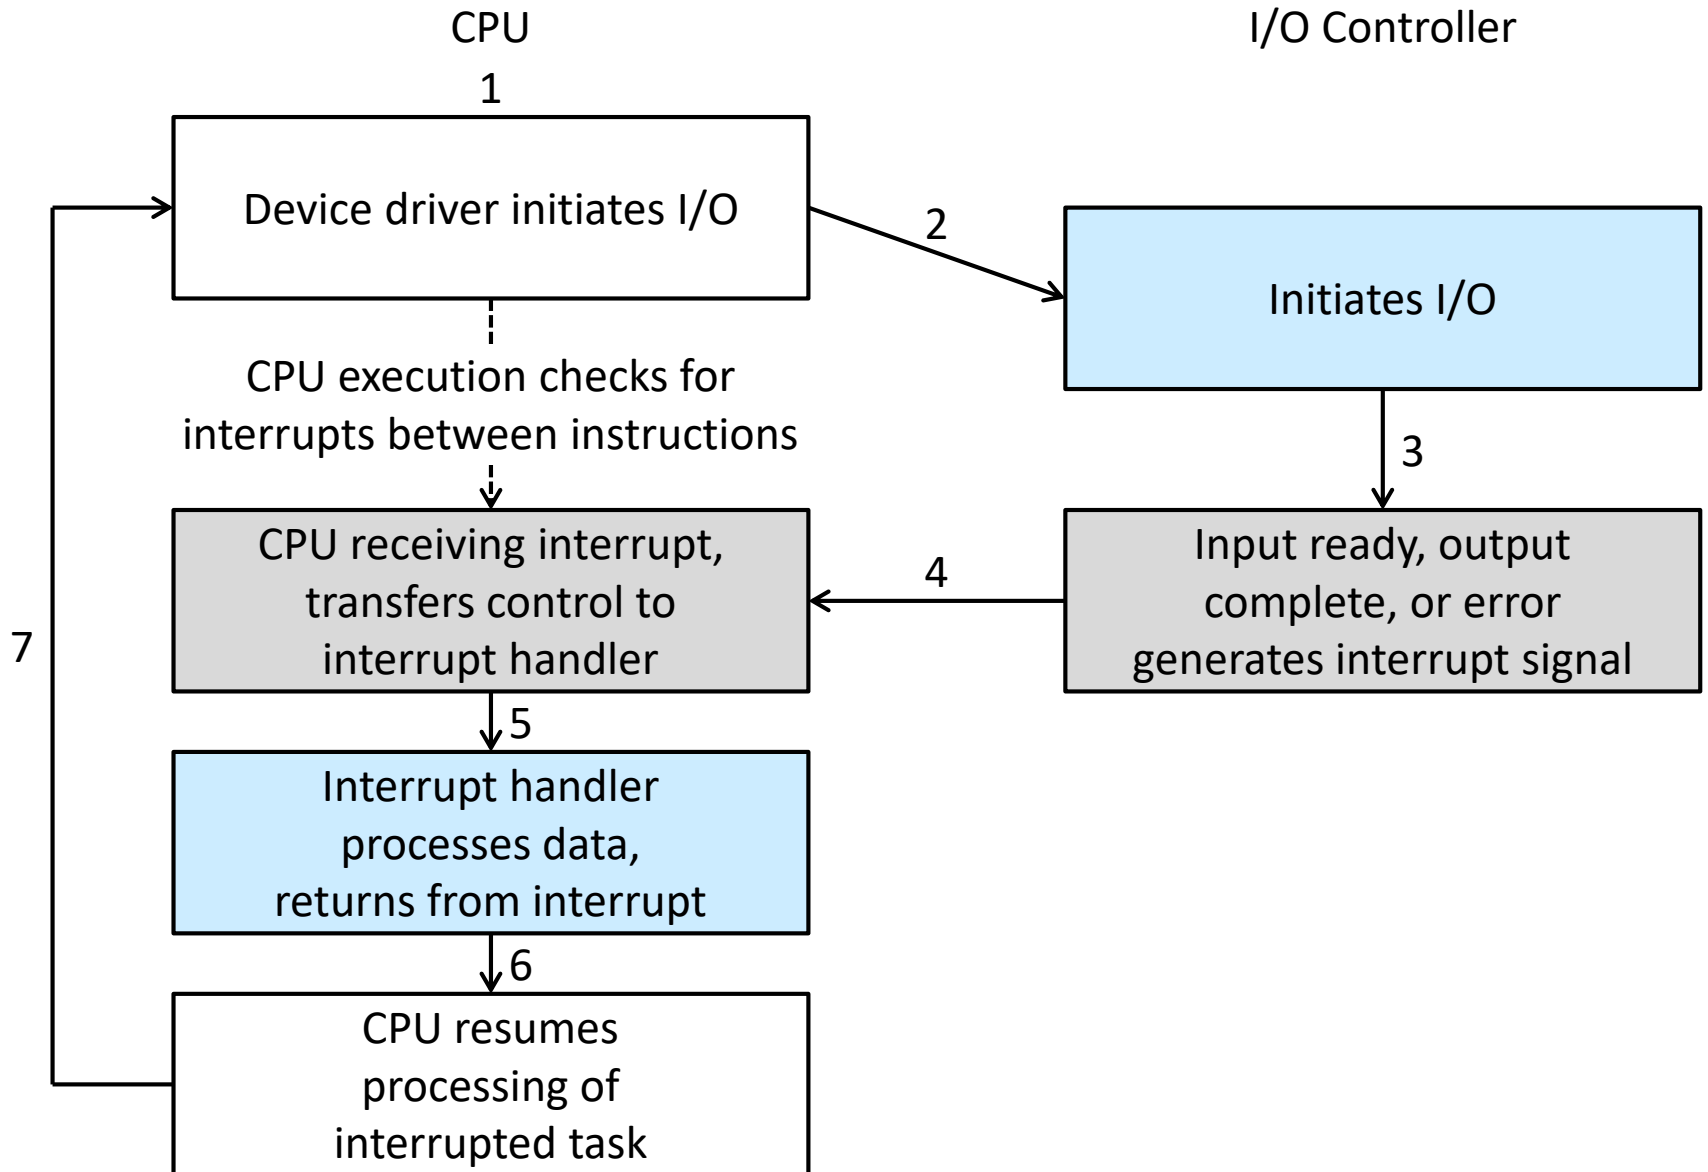
# Overview of Interrupts

- ❑ **Hardware** may trigger an interrupt at any time by sending a signal to the CPU (usually by way of the system bus)
- ❑ **Steps**
  - When the CPU is interrupted, it stops what it is doing and immediately transfers execution to a fixed location
  - The fixed location usually contains the starting address where the service routine for the interrupt is located
  - The interrupt service routine executes
  - On completion, the CPU resumes the interrupted computation
- ❑ **Interrupts must be handled quickly, as they occur very frequently**
  - **Interrupt vector**
    - An array (table) of pointers to interrupt routines

# Interrupt Timeline



# Interrupt-Driven I/O Cycle



# Implementation of Interrupts

## ❑ Interrupt-handler routine

## ❑ Desired features

- The ability to defer interrupt handling during critical processing
  - Nonmaskable and maskable interrupt request lines
- An efficient way to dispatch to the proper interrupt handler for a device
  - Interrupt vector and interrupt chaining
- Multilevel interrupts to distinguish between high- and low-priority interrupts and respond with the appropriate degree of urgency
  - Interrupt priority levels

# Intel Processor Event-Vector Table

Vector Number	Description
0	Divide Error
1	Debug Exception
2	Null Interrupt
3	Breakpoint
4	INTO-Detected Overflow
5	Bound Range Exception
6	Invalid Opcode
7	Device Not Available
8	Double Fault
9	Coprocessor Segment Overrun (Reserved)
10	Invalid Task State Segment
11	Segment Not Present
12	Stack Fault
13	General Protection
14	Page Fault
15	(Intel Reserved, Do Not Use)
16	Floating-Point Error
17	Alignment Check
18	Machine Check
19--31	(Intel Reserved, Do Not Use)
32--255	Maskable Interrupts

# Outline

- ❑ What Operating Systems Do
- ❑ **Computer-System Organization**
  - Interrupts, Storage Structure, I/O Structure
- ❑ Computer-System Architecture
- ❑ Operating-System Operations
- ❑ Resource Management, Security and Protection, Virtualization
- ❑ Distributed Systems
- ❑ Kernel Data Structures
- ❑ Computing Environments
- ❑ Free and Open-Source Operating Systems

# Main Memory

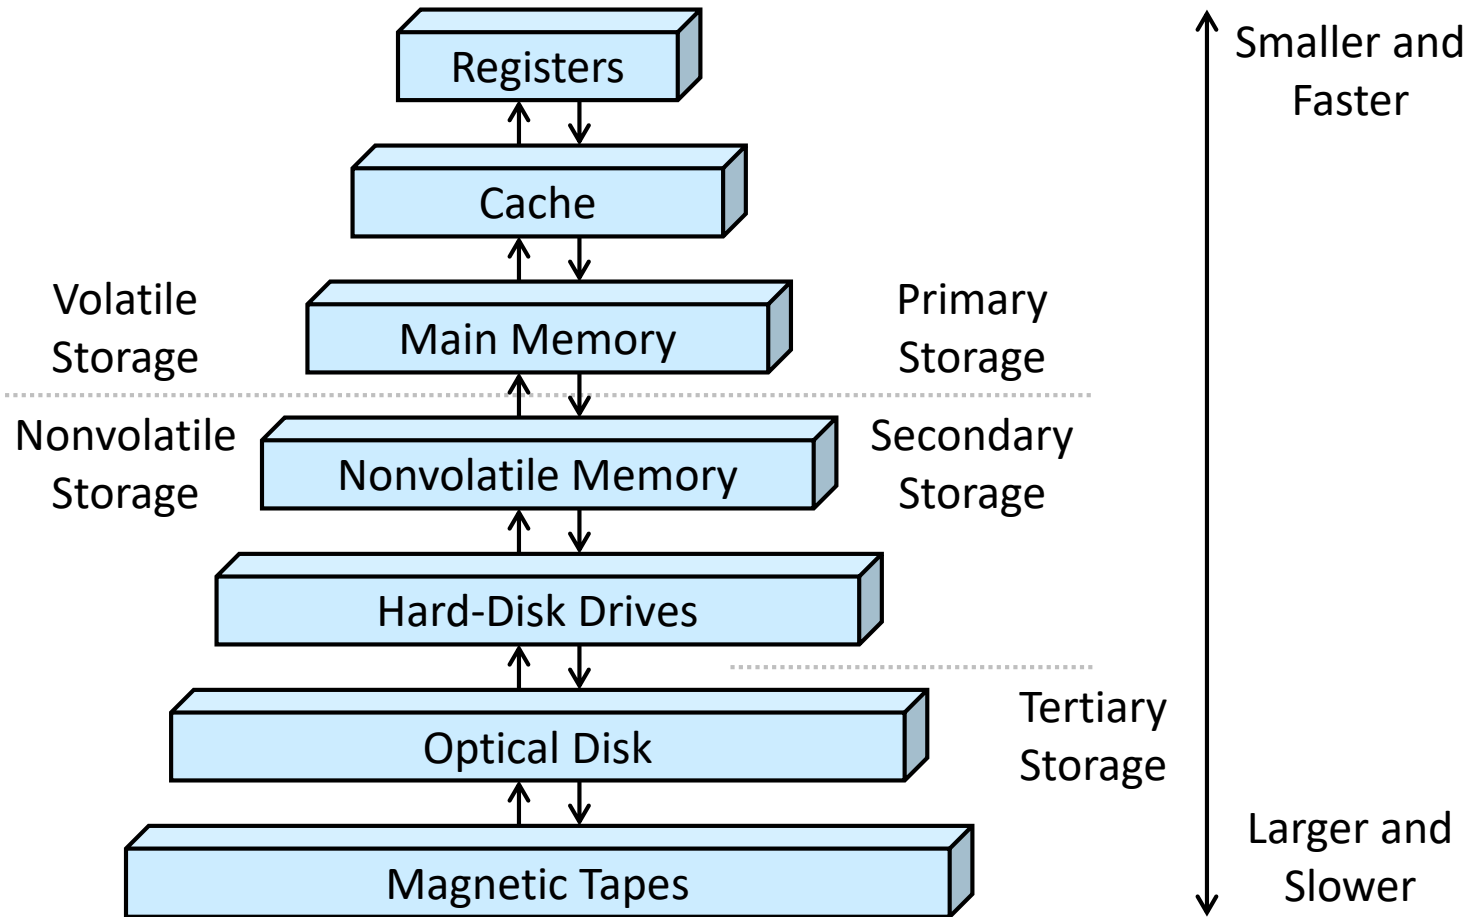
- ❑ General-purpose computers run most of their programs from rewritable memory
  - Called main memory, also called random-access memory (RAM)
  - Typically dynamic random-access memory (DRAM)
- ❑ The first program to run on computer power-on is a bootstrap program which then loads the operating system
  - Since RAM is volatile, the computer uses storage that is infrequently written to and is nonvolatile
  - Example: electrically erasable programmable read-only memory (EEPROM)
- ❑ Reference
  - von Neumann architecture



# Secondary and Tertiary Storage

- ❑ The programs and data cannot reside in main memory permanently
  - Main memory is usually too small to store all needed programs and data
  - Main memory is volatile
- ❑ Secondary storage
  - Hard disk drives (HDD)
  - Nonvolatile memory (NVM) devices
- ❑ Tertiary storage
  - Examples: CD-ROM, blu-ray, magnetic tapes
- ❑ The main differences among the various storage systems lie in speed, size, and volatility

# Storage-Device Hierarchy



# Outline

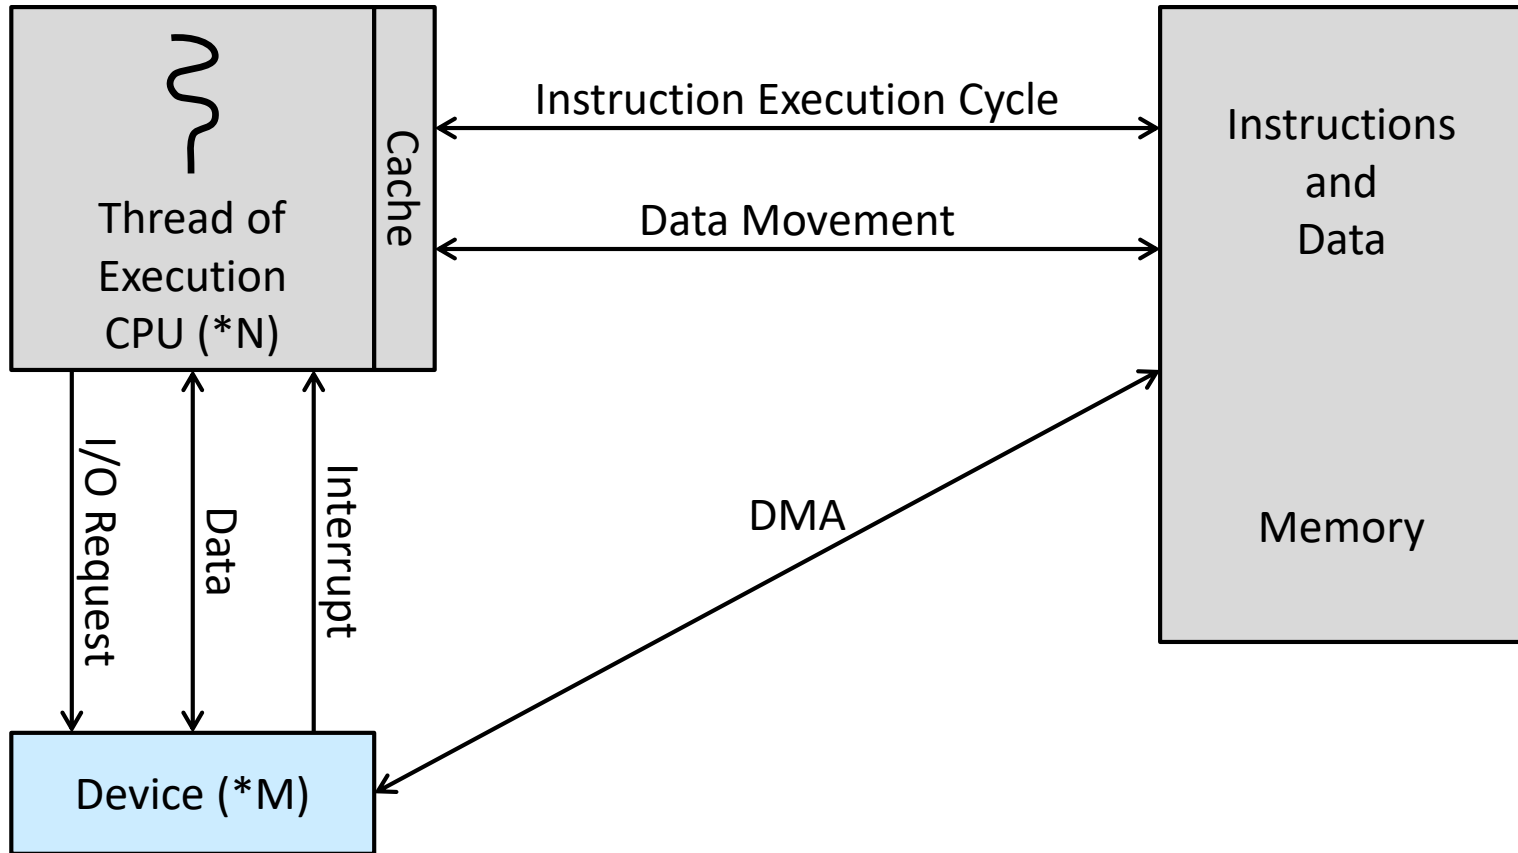
- ❑ What Operating Systems Do
- ❑ **Computer-System Organization**
  - Interrupts, Storage Structure, I/O Structure
- ❑ Computer-System Architecture
- ❑ Operating-System Operations
- ❑ Resource Management, Security and Protection, Virtualization
- ❑ Distributed Systems
- ❑ Kernel Data Structures
- ❑ Computing Environments
- ❑ Free and Open-Source Operating Systems

# Direct Memory Access (1/2)

- ❑ A large portion of operating system code is dedicated to managing I/O
  - Interrupt-driven I/O is fine for moving small amounts of data, but it can produce high overhead for bulk data movement
- ❑ Direct memory access (DMA)
  - The device controller transfers an entire block of data directly to or from the device and main memory
    - No CPU intervention: the CPU is available to accomplish other work
  - Only one interrupt is generated per block to tell the device driver that the operation has completed
- ❑ Some high-end systems use switch (not bus) architecture
  - Multiple components can talk to others concurrently
  - DMA is even more effective in this case

# Direct Memory Access (2/2)

## □ How a modern computer system works



# Outline


- ❑ What Operating Systems Do
- ❑ Computer-System Organization
- ❑ **Computer-System Architecture**
- ❑ Operating-System Operations
- ❑ Resource Management
- ❑ Security and Protection
- ❑ Virtualization
- ❑ Distributed Systems
- ❑ Computing Environments
- ❑ Kernel Data Structures, Free and Open-Source Operating Systems

# Computer-System Architecture

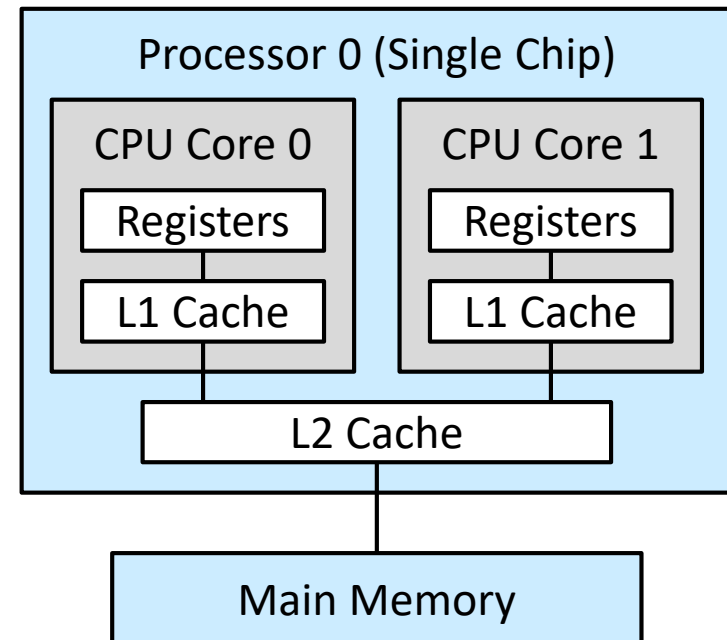
## ❑ Single-processor systems

## ❑ Multiprocessor systems

➤ Growing in use and importance



Symmetric Multiprocessing Architecture



Dual-Core on the Same Chip

## ❑ Clustered systems

➤ Two or more individual systems (or nodes) joined together

# Outline

- ❑ What Operating Systems Do
- ❑ Computer-System Organization
- ❑ Computer-System Architecture
- ❑ **Operating-System Operations**
  - Multiprogramming and Multitasking, Dual-Mode and Multimode Operation, Timer
- ❑ Resource Management, Security and Protection, Virtualization
- ❑ Distributed Systems
- ❑ Kernel Data Structures
- ❑ Computing Environments
- ❑ Free and Open-Source Operating Systems



# Operating-System Operations

## ❑ Bootstrap program

- Initialize all aspects of the system
- Locate the operating-system kernel and load it into memory

## ❑ System daemons

- Services provided outside of the kernel by system programs, running the entire time the kernel is running

## ❑ Events are almost always signaled by the occurrence of an interrupt

- Hardware interrupt (described already)
- Software interrupt (exception or trap)
  - Software error (e.g., division by zero)
  - A specific request from a user program that an operating-system service be performed by executing a special operation called a system call

# Outline

- ❑ What Operating Systems Do
- ❑ Computer-System Organization
- ❑ Computer-System Architecture
- ❑ **Operating-System Operations**
  - **Multiprogramming and Multitasking**, Dual-Mode and Multimode Operation, Timer
- ❑ Resource Management, Security and Protection, Virtualization
- ❑ Distributed Systems
- ❑ Kernel Data Structures
- ❑ Computing Environments
- ❑ Free and Open-Source Operating Systems


# Multiprogramming and Multitasking

❑ One of the most important aspects of operating systems is the ability to run multiple programs

- A program in execution is termed a **process**
- A set of processes is kept in memory
- When a process has to wait, operating systems switch to another process

❑ **Multitasking** is a logical extension of multiprogramming

- The CPU executes multiple processes by switching among them
- The switches occur frequently and provide fast response times



Memory Layout

# Multiprogramming and Multitasking

## ❑ Connections to following chapters

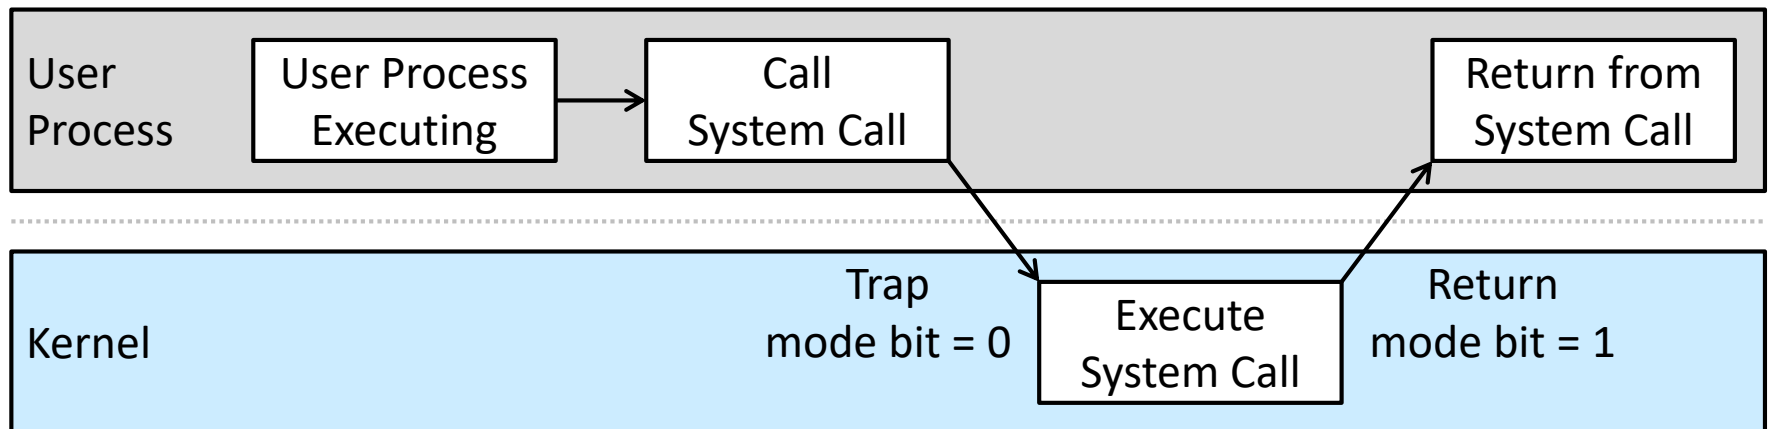
- [9, 10] Having several processes in memory at the same time requires memory management
- [5] The system must choose which process will run next
- [13, 14, 15] Multiprogramming and multitasking systems must provide a file system
- [11] The file system resides on a secondary storage which needs storage management
- [17] A system must protect resources from inappropriate use
- [6, 7] The system must ensure orderly execution for processes
- [8] It may ensure that processes do not get stuck in a deadlock

# Outline

- ❑ What Operating Systems Do
- ❑ Computer-System Organization
- ❑ Computer-System Architecture
- ❑ **Operating-System Operations**
  - Multiprogramming and Multitasking, **Dual-Mode and Multimode Operation**, Timer
- ❑ Resource Management, Security and Protection, Virtualization
- ❑ Distributed Systems
- ❑ Kernel Data Structures
- ❑ Computing Environments
- ❑ Free and Open-Source Operating Systems

# Dual-Mode and Multimode Operation

- ❑ Distinguish between the execution of operating-system code and user-defined code
  - User mode and kernel mode
- ❑ A mode bit provided by hardware
  - A system call changes it to "kernel" and its return resets it to "user"



- ❑ Hardware allows privileged instructions to be executed only in kernel mode

# Outline

- ❑ What Operating Systems Do
- ❑ Computer-System Organization
- ❑ Computer-System Architecture
- ❑ **Operating-System Operations**
  - Multiprogramming and Multitasking, Dual-Mode and Multimode Operation, **Timer**
- ❑ Resource Management, Security and Protection, Virtualization
- ❑ Distributed Systems
- ❑ Kernel Data Structures
- ❑ Computing Environments
- ❑ Free and Open-Source Operating Systems

# Timer

- ❑ We cannot allow a user program to
  - Get stuck in an infinite loop
  - Fail to call system services and never return control to the operating system
- ❑ A timer can be set to interrupt the computer after a specified period
  - Before turning over control to the user, the operating system ensures that the timer is set to interrupt



# Outline

- ❑ What Operating Systems Do
- ❑ Computer-System Organization
- ❑ Computer-System Architecture
- ❑ Operating-System Operations
- ❑ **Resource Management**
- ❑ Security and Protection
- ❑ Virtualization
- ❑ Distributed Systems
- ❑ Kernel Data Structures
- ❑ Computing Environments
- ❑ Free and Open-Source Operating Systems

# Resource Management

## ❑ Process management

- Chapters 3, 4, 5, 6, and 7

## ❑ Memory management

- Chapters 9 and 10

## ❑ File-system management

- Chapters 13, 14, and 15

## ❑ Mass-storage management

- Chapter 11

## ❑ Cache management

## ❑ I/O system management

- Chapter 12

# Outline

- ❑ What Operating Systems Do
- ❑ Computer-System Organization
- ❑ Computer-System Architecture
- ❑ Operating-System Operations
- ❑ Resource Management
- ❑ **Security and Protection**
- ❑ Virtualization
- ❑ Distributed Systems
- ❑ Kernel Data Structures
- ❑ Computing Environments
- ❑ Free and Open-Source Operating Systems

# Protection and Security

## □ Protection

- Any mechanism for controlling the access of processes or users to the resources defined by a computer system
- Chapter 17

## □ Security

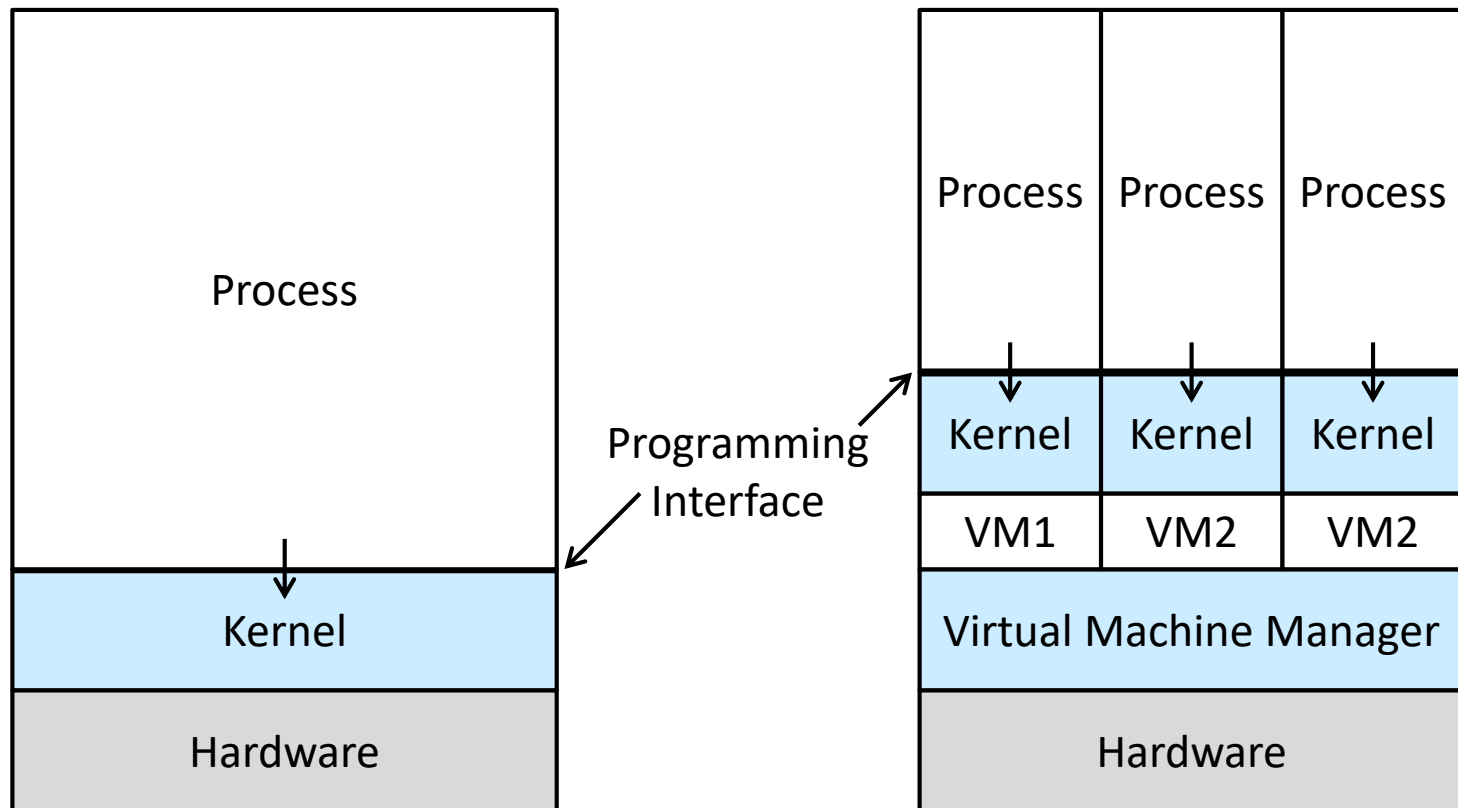
- Defend a system from external and internal attacks
  - Viruses and worms
  - Denial-of-service attacks
  - Identity theft
  - Theft of service
- Chapter 16

# Outline

- ❑ What Operating Systems Do
- ❑ Computer-System Organization
- ❑ Computer-System Architecture
- ❑ Operating-System Operations
- ❑ Resource Management
- ❑ Security and Protection
- ❑ **Virtualization**
- ❑ Distributed Systems
- ❑ Kernel Data Structures
- ❑ Computing Environments
- ❑ Free and Open-Source Operating Systems

# Virtualization

- ❑ Abstract the hardware of a single computer into several different execution environments
  - Create the illusion that each separate environment is running on its own private computer



# Outline

- ❑ What Operating Systems Do
- ❑ Computer-System Organization
- ❑ Computer-System Architecture
- ❑ Operating-System Operations
- ❑ Resource Management
- ❑ Security and Protection
- ❑ Virtualization
- ❑ **Distributed Systems**
- ❑ Kernel Data Structures
- ❑ Computing Environments
- ❑ Free and Open-Source Operating Systems

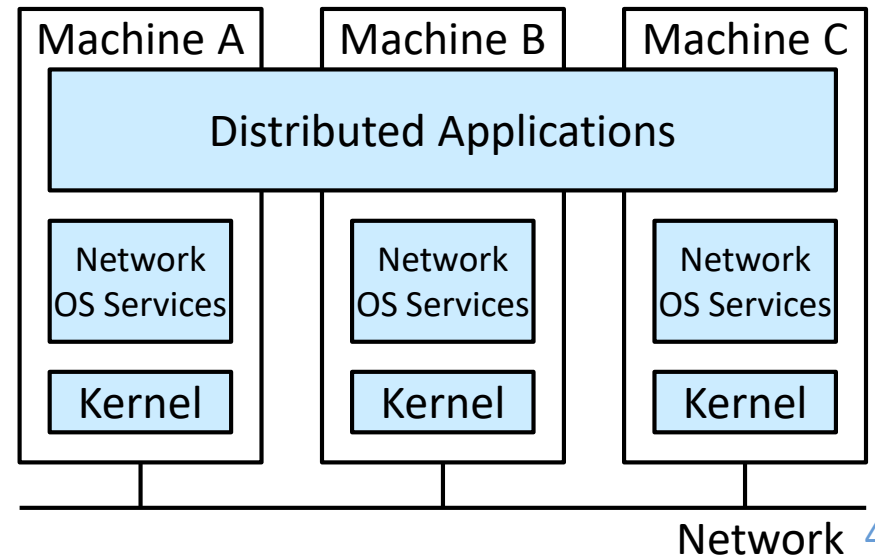
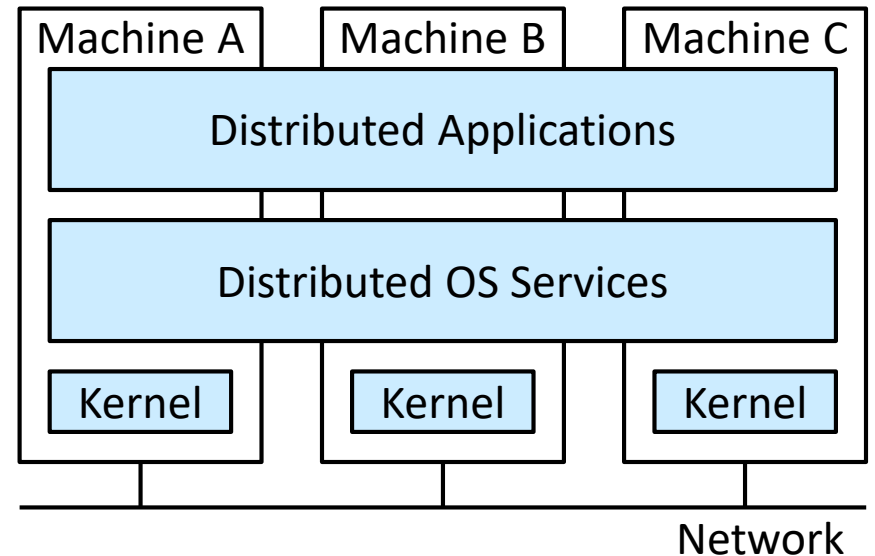
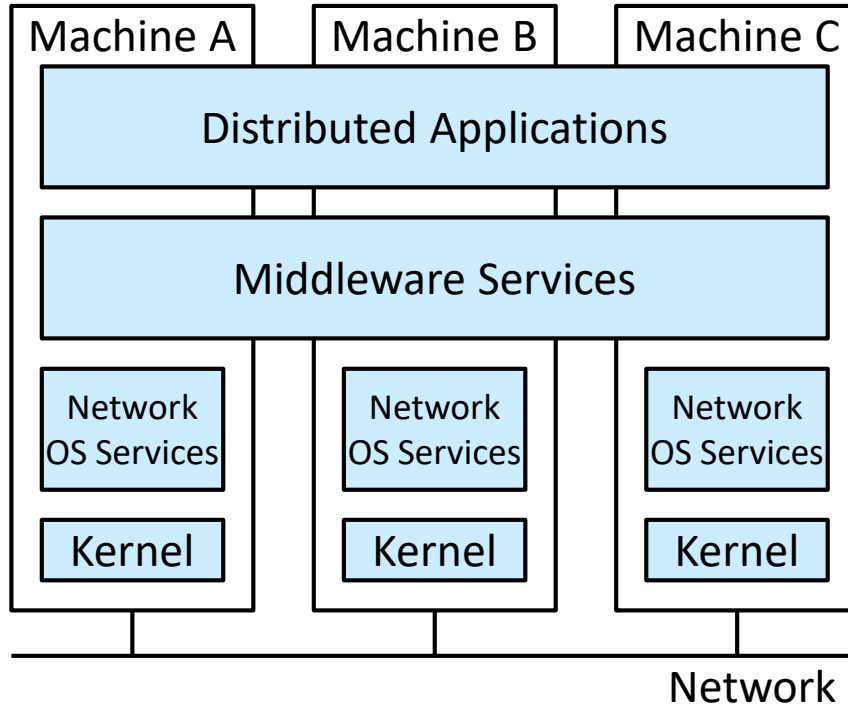
# Distributed Systems

- ❑ A collection of networked, separate, and possibly heterogeneous computer systems
- ❑ A **network** is a communication path between two or more systems
  - **TCP/IP** is the most common network protocol
  - Personal/Local/Wide/Metropolitan Area Network (PAN/LAN/WAN/PAN)



# Distributed Systems

- ❑ Distributed OSs →
- ❑ Network OSs ↘
- ❑ Middleware-based systems  
↓



# Outline

- ❑ What Operating Systems Do
- ❑ Computer-System Organization
- ❑ Computer-System Architecture
- ❑ Operating-System Operations
- ❑ Resource Management
- ❑ Security and Protection
- ❑ Virtualization
- ❑ Distributed Systems
- ❑ **Kernel Data Structures**
- ❑ Computing Environments
- ❑ Free and Open-Source Operating Systems

# Kernel Data Structures

- ❑ Lists, stacks, and queues
- ❑ Trees
- ❑ Hash functions and maps
- ❑ Bitmaps

# Outline

- ❑ What Operating Systems Do
- ❑ Computer-System Organization
- ❑ Computer-System Architecture
- ❑ Operating-System Operations
- ❑ Resource Management
- ❑ Security and Protection
- ❑ Virtualization
- ❑ Distributed Systems
- ❑ Kernel Data Structures
- ❑ **Computing Environments**
- ❑ Free and Open-Source Operating Systems

# Computing Environments

- ❑ Traditional computing
- ❑ Mobile computing
  - Apple iOS and Google Android
- ❑ Client-server computing
- ❑ Peer-to-peer computing
  - Not distinguish clients and servers
- ❑ Cloud computing
  - A logical extension of virtualization
  - Public cloud, private cloud, and hybrid cloud
  - Software as a service (SaaS), platform as a service (PaaS), and infrastructure as a service (IaaS)
- ❑ Real-time embedded systems

# Computing Modes

## ❑ Batch mode for efficiency

- Programs run independent of the environment
- Results are written to files or output devices
- Examples: compiler, machine learning

## ❑ Online mode for responsiveness

- Users may interact with program
- No time constraint is involved
- Examples: spreadsheet, game, browsing

## ❑ Real-time mode for predictability

- Programs interact strongly with the environment
- Time constraints are imposed
- Examples: autonomous vehicles, robot control

# Outline

- ❑ What Operating Systems Do
- ❑ Computer-System Organization
- ❑ Computer-System Architecture
- ❑ Operating-System Operations
- ❑ Resource Management
- ❑ Security and Protection
- ❑ Virtualization
- ❑ Distributed Systems
- ❑ Kernel Data Structures
- ❑ Computing Environments
- ❑ **Free and Open-Source Operating Systems**

# Free and Open-Source OSs

- ❑ GNU/Linux
- ❑ BSD UNIX
- ❑ Solaris



# Objectives

- ❑ Describe the general organization of a computer system and the role of interrupts
- ❑ Describe the components in a modern multiprocessor computer system
- ❑ Illustrate the transition from user mode to kernel mode
- ❑ Discuss how operating systems are used in various computing environments

# Q&A