

# Operating Systems

## [ 3. Process ]

Chung-Wei Lin

[cwlin@csie.ntu.edu.tw](mailto:cwlin@csie.ntu.edu.tw)

CSIE Department

National Taiwan University

# Objectives

- ❑ Identify the separate components of a process and illustrate how they are represented and scheduled in an operating system
- ❑ Describe how processes are created and terminated in an operating system, including developing programs using the appropriate system calls that perform these operations
- ❑ Describe and contrast interprocess communication using shared memory and message passing
- ❑ Design programs that use pipes and POSIX shared memory to perform interprocess communication
- ❑ Describe client-server communication using sockets and remote procedure calls

# Outline

## ☐ **Process Concept**

➤ The Process, Process State, Process Control Block, Threads

## ☐ Process Scheduling

## ☐ Operations on Processes

## ☐ Interprocess Communication

## ☐ IPC in Shared-Memory Systems

## ☐ IPC in Message-Passing Systems

## ☐ Examples of IPC Systems

## ☐ Communication in Client-Server Systems

# Process Concept

## ❑ Process: a program in execution

- A program is a passive entity
  - A file (executable file) containing a list of instructions stored on disk
- A process is an active entity with
  - A program counter specifying the next instruction to execute
  - A set of associated resources
- A program becomes a process when an executable file is loaded into memory
  - Double-click an icon representing the executable file
  - Enter the name of the executable file on the command line
- Although two processes may be associated with the same program, each of these is a separate process
  - It is also common to have a process that spawns many processes as it runs

# Memory Layout of a Process

## ❑ Text section

- The executable code

## ❑ Data section

- Global variables

## ❑ Heap section


- Dynamically allocated during program run time

## ❑ Stack section

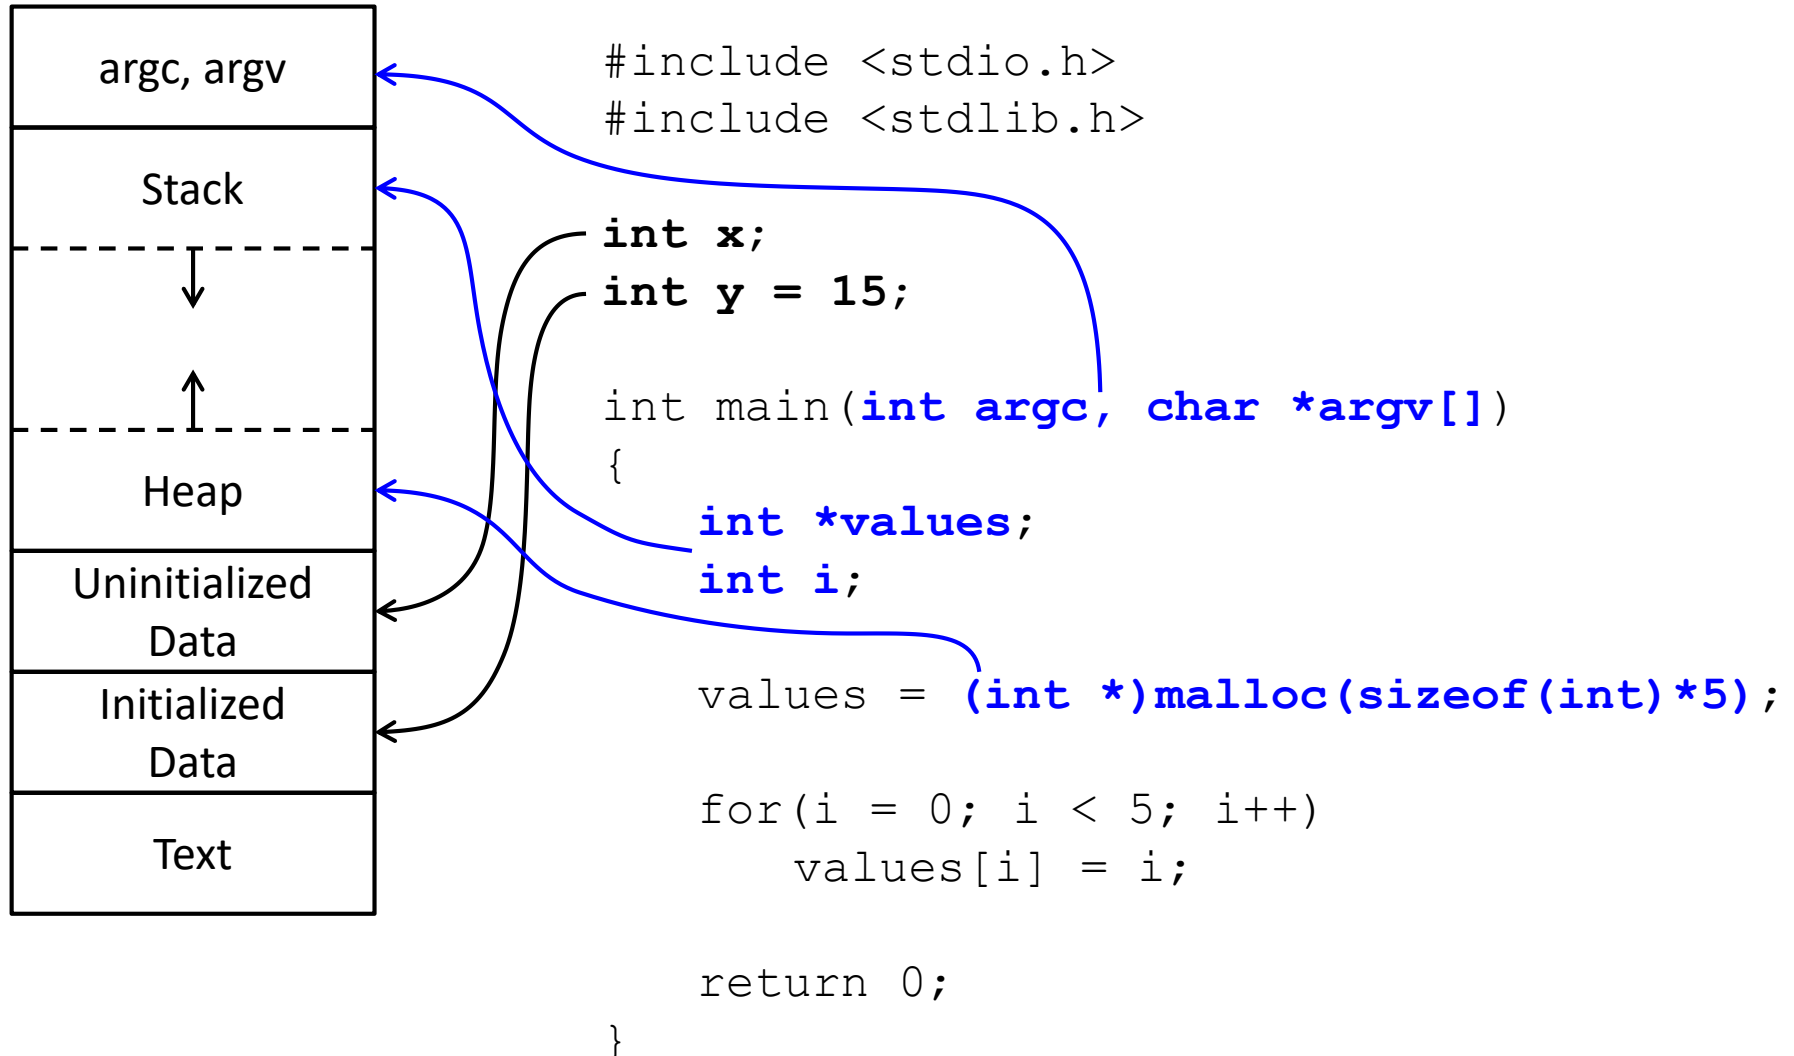
- Temporary data storage when invoking functions
- Examples: function parameters, return addresses, local variables

## ❑ The stack and heap sections can shrink and grow dynamically during program execution

- The operating system must ensure they do not overlap one another



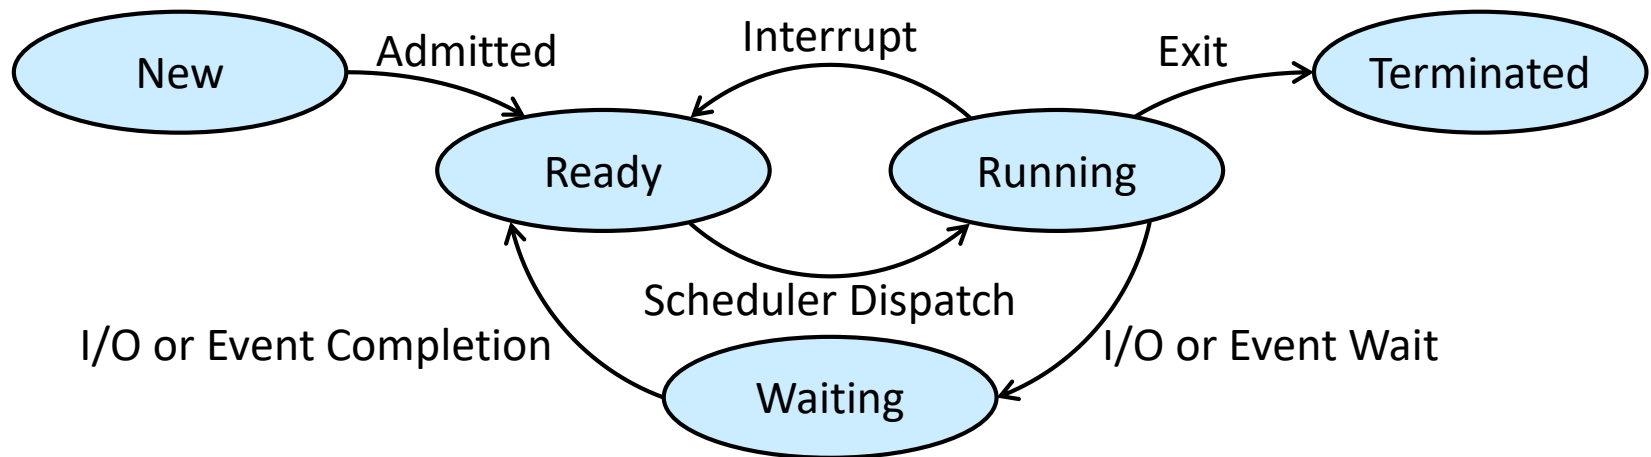
# Memory Layout of a C Program



# Process State

## ❑ As a process executes, it changes state

- **New**: the process is being created
- **Ready**: the process is waiting to be assigned to a processor
- **Running**: instructions are being executed
  - Only one process can be running on any processor core at any instant
- **Waiting**: the process is waiting for some event to occur
  - Examples: I/O completion, reception of a signal
- **Terminated**: the process has finished execution



# Process Control Block

- ❑ Each process is represented in the operating system by a **process control block** (PCB), also called a **task control block**
  - Process state
  - Process number
  - Program counter
    - Address of the next instruction to be executed for this process
  - CPU registers
    - Register set where process needs to be stored for execution for running state
  - CPU-scheduling information: priority, queue pointers, etc. [Chapter 5]
  - Memory-management information: memory limits, etc. [Chapter 9]
  - Accounting information
    - Amount of CPU and real time used, time limits, account numbers, etc.
  - I/O status information
    - List of I/O devices allocated to the process, list of open files, etc.



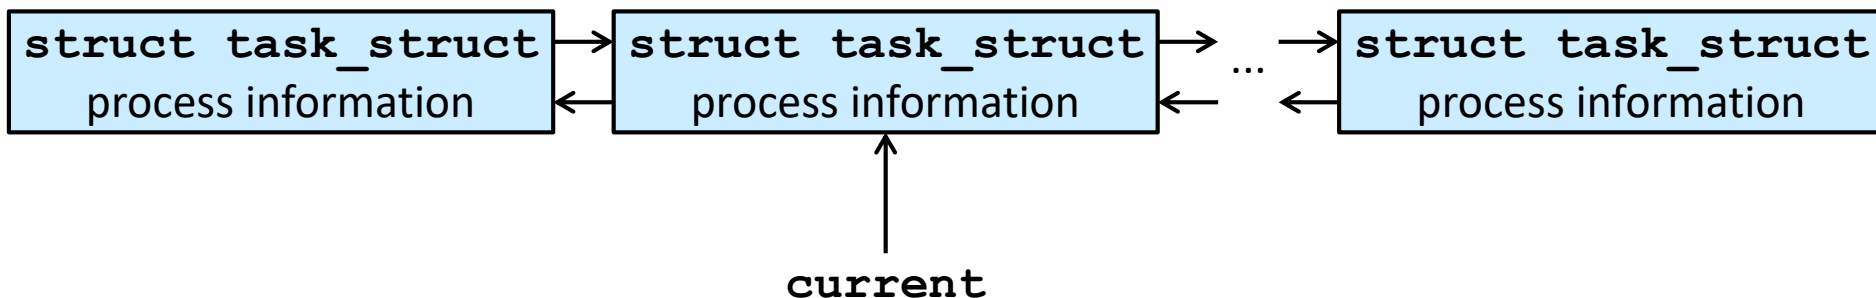
# Process Representation in Linux

## ❑ Represented by the C structure **task\_struct**

```
pid t_pid;                /* process identifier */
long state;               /* state of the process */
unsigned int time_slice   /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children */
struct files_struct *files; /* list of open files */
struct mm_struct *mm;     /* address space of this process */
```

## ❑ Within the Linux kernel, all active processes are represented using a doubly linked list of **task\_struct**

- The kernel maintains a pointer **current** to the process currently executing on the system



# Threads

- ❑ The process model so far has implied that a process is a program that performs a single thread of execution
  - The process performs only one task at a time
  - Example: the user cannot simultaneously type in characters and run the spell checker
- ❑ Most modern operating systems allow a process to have multiple threads of execution
  - The process performs more than one task at a time
    - Especially beneficial on multicore systems where multiple threads can run in parallel
  - The PCB is expanded to include information for each thread
    - Other changes throughout the system are also needed [Chapter 4]

# Outline

- ❑ Process Concept

- ❑ **Process Scheduling**

  - Scheduling Queues, CPU Scheduling, Context Switch

- ❑ Operations on Processes

- ❑ Interprocess Communication

- ❑ IPC in Shared-Memory Systems

- ❑ IPC in Message-Passing Systems

- ❑ Examples of IPC Systems

- ❑ Communication in Client-Server Systems

# Process Scheduling (1/2)

## ❑ Objective of multiprogramming

- To have some process running at all times and maximize CPU utilization

## ❑ Objective of time sharing

- To switch a CPU core among processes so frequently that users can interact with each program while it is running

# Process Scheduling (2/2)

- ❑ To meet the objectives, the process scheduler selects an available process for program execution on a core
  - Each CPU core can run one process at a time
  - If there are more processes than cores, excess processes will have to wait until a core is free and can be rescheduled
  - The number of processes currently in memory is known as the degree of multiprogramming
- ❑ To meet the objectives, it also requires taking the general behavior of a process into account
  - An I/O-bound process uses more of its time doing I/O than computations
  - A CPU-bound process uses more of its time doing computations than I/O

# Scheduling Queues

## ❑ The system includes some queues

### ➤ Ready queue

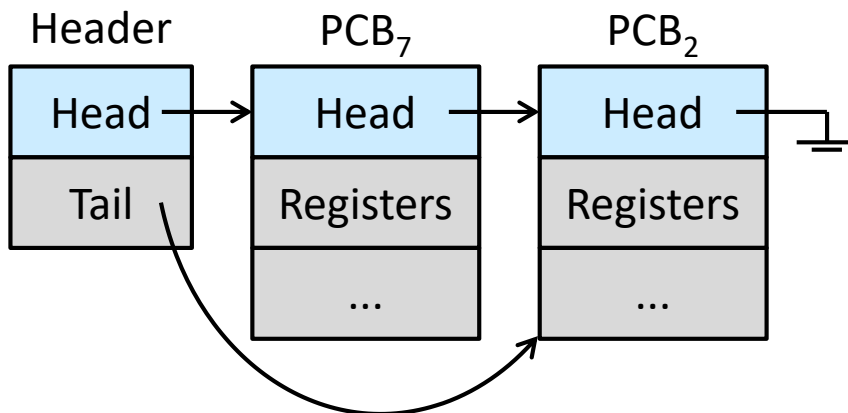
- The set of processes ready and waiting to execute on a CPU's core

### ➤ Wait queues

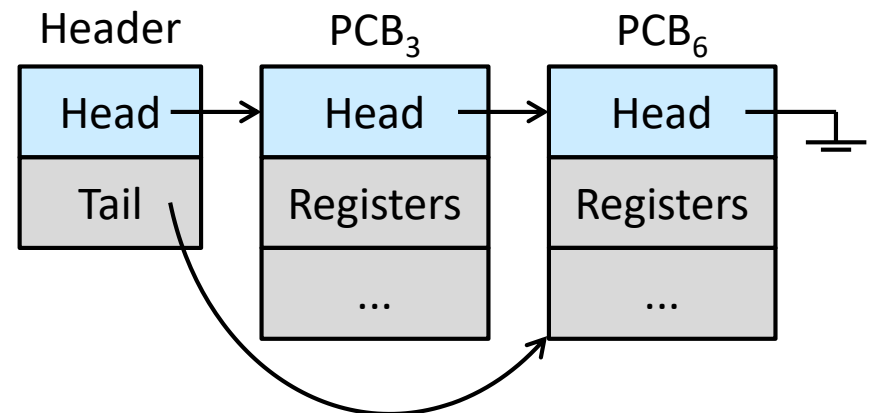
- The set of processes waiting for a certain event (e.g., completion of I/O) to occur

### ➤ A queue is generally stored as a linked list

Ready Queue

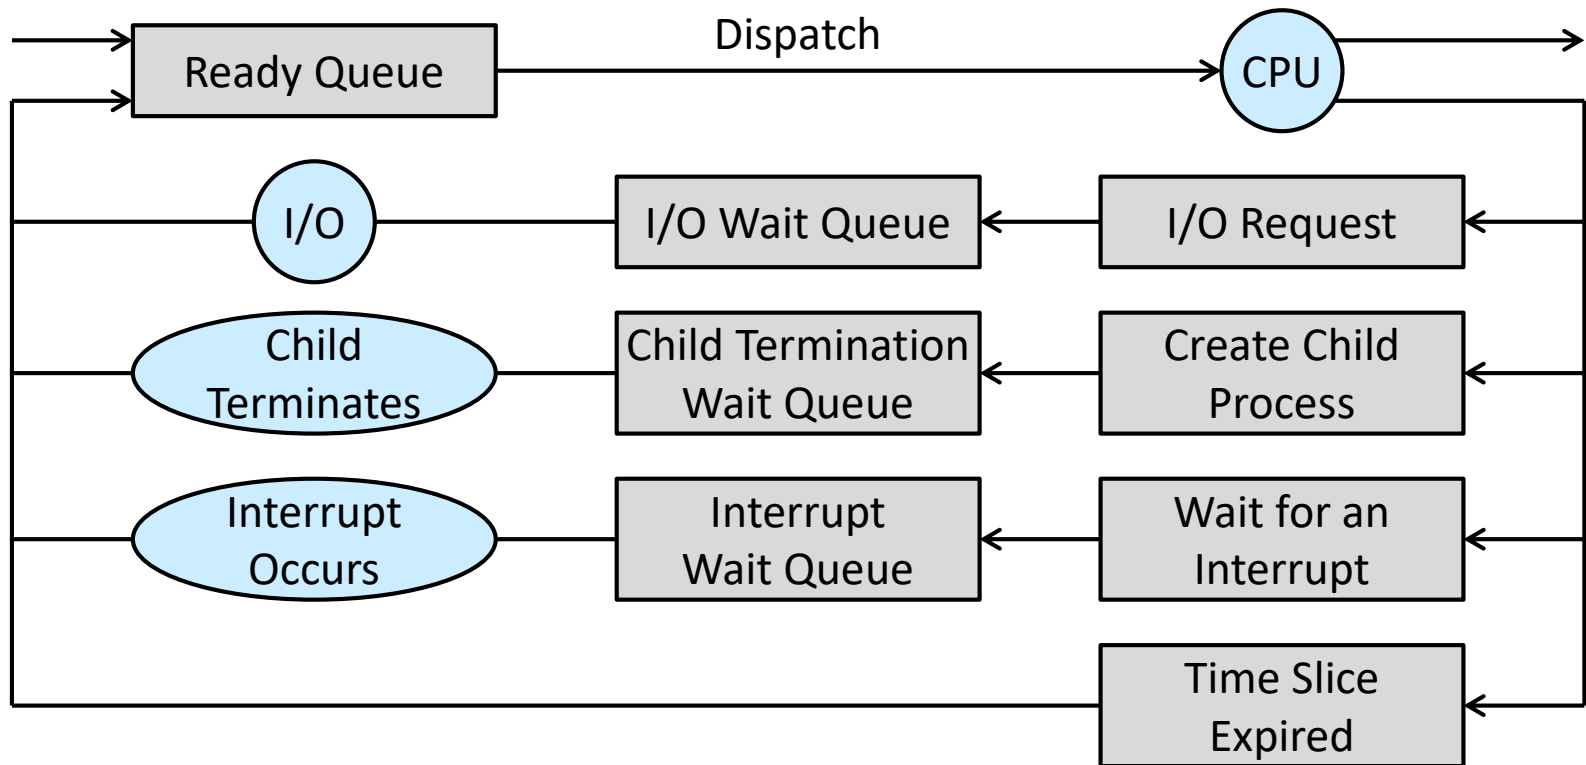


Wait Queue



# Queueing Diagram

- ❑ A ready queue and a set of wait queues
- ❑ Circles represent the resources that serve the queues
- ❑ Arrows indicate the flow of processes in the system



# CPU Scheduling

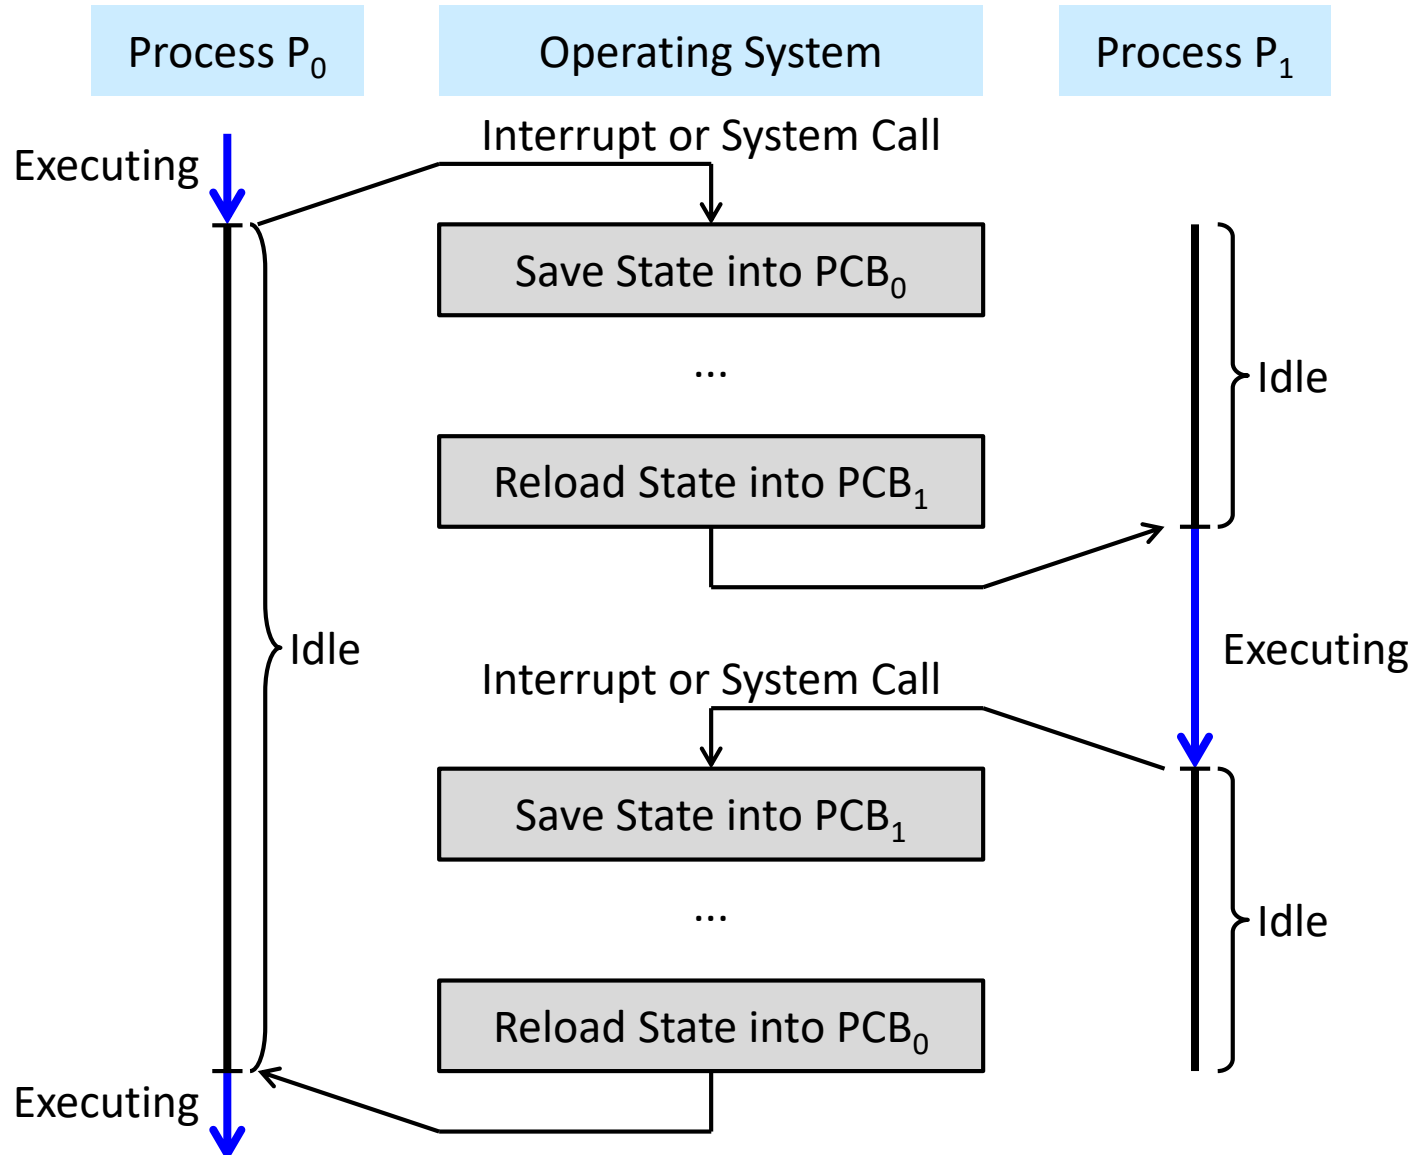
- ❑ A process migrates among the ready queue and various wait queues throughout its lifetime
- ❑ A **CPU scheduler** executes at least once every 100 milliseconds, although typically much more frequently
  - An I/O-bound process may execute for only a few milliseconds before waiting for an I/O request
  - A CPU-bound process will require a CPU core for longer durations, but the scheduler is unlikely to grant the core to it for an extended period
- ❑ **Swapping** [Chapter 9]
  - Remove a process from memory (i.e., active contention for CPU) to disk
  - Reduce the degree of multiprogramming
  - Later, reintroduce the process into memory and continue its execution



# Context Switch (1/2)

- ❑ Switching the CPU core from one process to another
- ❑ View of the current process
  - A state save of the current state of the CPU core and then a state restore to resume operations
- ❑ View of the system
  - Context switch: a state save of the current process and then a state restore of a different process
- ❑ The context of a process is represented in its PCB
  - The value of the CPU registers, the process state, and memory-management information, etc.

# Diagram Showing Context Switch



# Context Switch (2/2)

- ❑ Context switch time is pure overhead
  - The system does no useful work while switching
- ❑ Switching speed (usually several microseconds) depends on
  - Memory speed
  - Number of registers that must be copied
  - Existence of special instructions and hardware support
    - Example: a single instruction to load or store all registers
    - Example: multiple sets of registers, where a context switch here simply requires changing the pointer to the current register set
- ❑ The more complex the operating system, the greater the amount of work that must be done

# Multitasking in Mobile Systems


- ❑ Beginning with iOS 4, iOS allows a single **foreground** application to run with multiple **background** applications
  - Early versions of iOS does not provide user-application multitasking
    - Only one application runs in the foreground
    - All other user applications are suspended
  - **Split-screen**
    - A larger screen allows running two foreground applications at the same time
- ❑ Android runs foreground and background with fewer limits
  - A **service** runs on behalf of the background process
    - The service runs even if the background application is suspended
    - Services do not have a user interface and have a small memory footprint

# Outline

- ❑ Process Concept
- ❑ Process Scheduling
- ❑ **Operations on Processes**
  - **Process Creation**, Process Termination
- ❑ Interprocess Communication
- ❑ IPC in Shared-Memory Systems
- ❑ IPC in Message-Passing Systems
- ❑ Examples of IPC Systems
- ❑ Communication in Client-Server Systems

# Process Creation (1/2)

- ❑ A parent process creates children processes, which in turn create other processes, forming a tree of processes
  - Generally, processes are identified via a process identifier (`pid`)
- ❑ A tree of processes on a typical Linux system



# Process Creation (2/2)

## ❑ Resource (memory, files, etc.) sharing options

- The child shares all of the parent's resources
- The child shares a subset of the parent's resources
  - Prevent any process from overloading the system by creating too many children
- The child shares none of the parent's resources

## ❑ Execution options

- The parent continues to execute concurrently with its children
- The parent waits until some or all of its children have terminated

## ❑ Address-space options

- The child is a duplicate of the parent process
  - It has the same program and data as the parent
- The child has a new program loaded into it

# Process Creation Using UNIX `fork()`

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
int main() {
    pid_t pid;
    /* fork a child process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }
    return 0;
}
```



# Process Creation Using UNIX `fork()`

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
int main() {
    pid_t pid;
    /* fork a child process */
    pid = fork() ;
```

## ❑ `fork()` creates a new process

- The child (new) process consists of a copy of the address space of the parent (original) process
- Both processes continue execution at the instruction after `fork()` with one difference
  - The return code for `fork()` is nonzero (`pid` of the child) for the parent
  - The return code for `fork()` is zero for the child

```
        printf("Child Complete");
    }
    return 0;
}
```

# Process Creation Using UNIX `fork()`

❑ `exec()` loads a binary file into memory and starts execution

- `execlp()` in the example code
- It destroys the original memory image

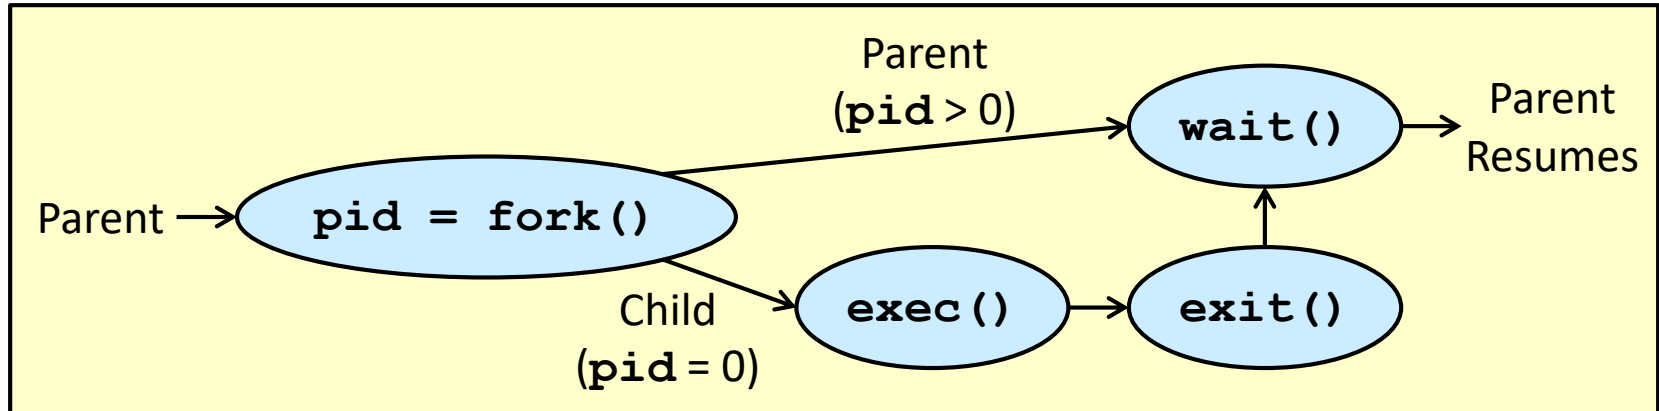
```
/* fork a child process */
pid = fork();
if (pid < 0) { /* error occurred */
    fprintf(stderr, "Fork Failed");
    return 1;
}
else if (pid == 0) { /* child process */
    execlp("/bin/ls", "ls", NULL);
}
else { /* parent process */
    /* parent will wait for the child to complete */
    wait(NULL);
    printf("Child Complete");
}
return 0;
}
```

# Process Creation Using UNIX `fork()`

- ❑ The parent calls `wait()` to move itself off the ready queue until the termination of the child
  - The child terminates by either implicitly or explicitly invoking `exit()`
    - Implicitly: the C run-time library (added to UNIX executable files) includes `exit()` by default
    - Explicitly: directly use `exit()`

```
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }
    return 0;
}
```

# Process Creation Using UNIX `fork()`



```
        fprintf(stderr, "Fork Failed");  
        return 1;  
    }  
    else if (pid == 0) { /* child process */  
        execlp("/bin/ls", "ls", NULL);  
    }  
    else { /* parent process */  
        /* parent will wait for the child to complete */  
        wait(NULL);  
        printf("Child Complete");  
    }  
    return 0;  
}
```

# Process Creation Using Windows API

- ❑ **CreateProcess ()** requires loading a specified program into the address space of the child process at process creation
  - UNIX example: **fork ()** has the child process inheriting the address space of its parent
- ❑ **CreateProcess ()** expects no fewer than ten parameters
  - UNIX example: **fork ()** is passed no parameter
- ❑ **WaitForSingleObject ()** waits for the child process to complete
  - UNIX example: **wait ()** is equivalent
- ❑ **Note: address-space options**
  - The child is a duplicate of the parent process
    - It has the same program and data as the parent
  - The child has a new program loaded into it

# Process Creation Using Windows API

```
int main(VOID){
    STARTUPINFO si;
    PROCESS_INFORMATION pi;
    /* allocate memory */
    ZeroMemory(&si, sizeof(si));
    si.cb = sizeof(si);
    ZeroMemory(&pi, sizeof(pi));
    /* create child process */
    if (!CreateProcess(NULL, /* use command line */
        "C:\\WINDOWS\\system32\\mspaint.exe", /* command */
        NULL, /* don't inherit process handle */
        NULL, /* don't inherit thread handle */
        FALSE, /* disable handle inheritance */
        0, /* no creation flags */
        NULL, /* use parent's environment block */
        NULL, /* use parent's existing directory */
        &si, &pi)){
        fprintf(stderr, "Create Process Failed");
        return -1;
    }
    /* parent will wait for the child to complete */
    WaitForSingleObject(pi.hProcess, INFINITE);
    printf("Child Complete");
    /* close handles */
    CloseHandle(pi.hProcess);
    CloseHandle(pi.hThread);
}
```

# Outline

- ❑ Process Concept
- ❑ Process Scheduling
- ❑ **Operations on Processes**
  - Process Creation, **Process Termination**
- ❑ Interprocess Communication
- ❑ IPC in Shared-Memory Systems
- ❑ IPC in Message-Passing Systems
- ❑ Examples of IPC Systems
- ❑ Communication in Client-Server Systems

# Process Termination (1/3)

- ❑ A process finishes executing its final statement and asks the operating system to delete it by using the **exit()** system call
  - A status value (typically an integer) is returned to its waiting parent
  - All the resources are deallocated and reclaimed by the operating system

- ❑ **exit()** provides an exit status

```
/* exit with status 1 */  
exit(1);
```

- ❑ **wait()** returns the process identifier of the terminated child

```
pid_t pid;  
int status;  
pid = wait(&status);
```

- What if there are multiple children or there is no child?
  - <http://www.csl.mtu.edu/cs4411.ck/www/NOTES/process/fork/wait.html>



# Process Termination (2/3)

- ❑ A parent may terminate the execution of one of its children
  - The child has exceeded its usage of some of the resources that it has been allocated
    - The parent must have a mechanism to inspect the state of its children
  - The task assigned to the child is no longer required
  - The parent is exiting, and the operating system does not allow a child to continue if its parent terminates
    - Cascading termination: if a process terminates (either normally or abnormally), then all its children must also be terminated
    - It is normally initiated by the operating system

# Process Termination (3/3)

- ❑ A **zombie** is a process that has terminated, but whose parent has not yet called **wait()**
  - When a process terminates, its resources are deallocated by the OS
  - However, its entry in the process table (containing the exit status) must remain until the parent calls **wait()**
  - Once the parent calls **wait()**, the zombie's process identifier and process-table entry are released
- ❑ An **orphan** is a process that its parent did not invoke **wait()** and terminated
  - UNIX assigns the **systemd** (or **init**) process as the new parent
  - The **systemd** (or **init**) process periodically invokes **wait()**
    - Allow the exit status of any orphan to be collected
    - Release the orphan's process identifier and process-table entry

# Android Process Hierarchy

- ❑ Importance hierarchy from most to least important
  - Foreground process
  - Visible process
  - Service process
  - Background process
  - Empty process
- ❑ If system resources must be reclaimed, Android will first terminate a least-important process

# Chrome Browser

## ❑ Multiprocess architecture: three different types of processes

- The **browser** process manages the user interface, disk and network I/O
- **Renderer** processes handles HTML, Javascript, images, and so forth
  - A new renderer process is created for each website opened in a new tab
- A **plug-in** process is created for each type of plug-in in use
  - Examples: QuickTime or Flash

## ❑ Each tab represents a separate process

- Better isolation
- Better security
  - Renderer processes run in a **sandbox** which means that access to disk and network I/O is restricted

# Outline


- ❑ Process Concept
- ❑ Process Scheduling
- ❑ Operations on Processes
- ❑ **Interprocess Communication**
- ❑ IPC in Shared-Memory Systems
- ❑ IPC in Message-Passing Systems
- ❑ Examples of IPC Systems
- ❑ Communication in Client-Server Systems

# Interprocess Communication (1/2)


- ❑ A process may be either independent or cooperating
  - A process is **independent** if it does not share data with any other processes executing in the system
  - A process is **cooperating** if it can affect or be affected by the other processes executing in the system
    - Any process that shares data with other processes is a cooperating process
- ❑ Reasons for process cooperation
  - Information sharing, computation speedup, modularity
- ❑ Cooperating processes require an **interprocess communication** (IPC) mechanism for data exchange
  - Model 1: **shared memory**
    - Generally faster (routine memory access once shared memory is established)
  - Model 2: **message passing**
    - Easier to implement in a distributed system

# Interprocess Communication (2/2)

- ❑ Cooperating processes require an interprocess communication (IPC) mechanism for data exchange



Shared Memory



Message Passing

# Outline

- ❑ Process Concept
- ❑ Process Scheduling
- ❑ Operations on Processes
- ❑ Interprocess Communication
- ❑ **IPC in Shared-Memory Systems**
- ❑ IPC in Message-Passing Systems
- ❑ Examples of IPC Systems
- ❑ Communication in Client-Server Systems



# IPC in Shared-Memory Systems

- ❑ Require communicating processes to establish a region of shared memory
  - Typically, the region resides in the address space of the process creating the shared-memory segment
  - Other processes attach the segment to their address space
- ❑ Normally, the operating system tries to prevent one process from accessing another process's memory
  - Shared memory requires that two or more processes agree to remove this restriction
  - The processes are also responsible for
    - The form of the data
    - The location of the data
    - Not writing to the same location simultaneously

# Producer-Consumer Problem

❑ A producer process produces information that is consumed by a consumer process

➤ Examples

- A compiler may produce assembly code that is consumed by an assembler
- The assembler may produce object modules that are consumed by the loader
- A web server produces web content such as HTML files and images, which are consumed by the client web browser requesting the resource

❑ Shared memory solution

- Have a buffer of items that can be filled (produced) by the producer and emptied (consumed) by the consumer
  - This buffer resides in the shared memory
- Synchronize the producer and the consumer [Chapters 6 and 7]
  - The consumer does not try to consume an item that has not yet been produced

# Types of Buffers

- ❑ The **unbounded buffer** places no practical limit on the size of the buffer
  - The producer can always produce new items
  - The consumer may have to wait for new items
- ❑ The **bounded buffer** assumes a fixed buffer size
  - The producer must wait if the buffer is full
  - The consumer must wait if the buffer is empty

# Bounded Buffer

## ❑ The shared **buffer**

```
#define BUFFER_SIZE 10
typedef struct {
    ...
} item;
item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```

- The variable **in** points to the next free position in the buffer
- The variable **out** points to the first full position in the buffer
- The buffer is empty when **in == out**
- The buffer is full when **((in + 1) % BUFFER\_SIZE) == out**
- This scheme allows at most **(BUFFER\_SIZE - 1)** items in the buffer

# Producer Process Using Shared Memory

## ❑ Producer process using shared memory

```
item next_produced;
while (true) {
    /* produce an item in next produced */
    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```

# Consumer Process Using Shared Memory

## ❑ Consumer process using shared memory

```
item next_consumed;  
while (true) {  
    while (in == out)  
        ; /* do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    /* consume the item in next consumed */  
}
```

## ❑ The issue of synchronization is not addressed here

- Both of the producer process and the consumer process attempt to access the shared buffer concurrently [Chapters 6 and 7]

# Outline

- ❑ Process Concept
- ❑ Process Scheduling
- ❑ Operations on Processes
- ❑ Interprocess Communication
- ❑ IPC in Shared-Memory Systems
- ❑ **IPC in Message-Passing Systems**
  - Naming, Synchronization, Buffering
- ❑ Examples of IPC Systems
- ❑ Communication in Client-Server Systems

# IPC in Message-Passing Systems (1/2)

- ❑ A message-passing facility provides at least two operations
  - `send(message)`
  - `receive(message)`
- ❑ Messages can be either fixed or variable in size
  - The system-level implementation is more straightforward with fixed-sized messages
  - The programming tasks are simpler with variable-sized messages



# IPC in Message-Passing Systems (2/2)

- ❑ A communication link must exist between processes sending and receiving messages
- ❑ This link can be implemented in a variety of ways
  - We are concerned here not with the link's physical implementation
    - Examples: shared memory, hardware bus, or network [Chapter 19]
  - We are concerned with its **logical implementation**
    - Direct or indirect communication
    - Synchronous or asynchronous communication
    - Automatic or explicit buffering

# Outline

- ❑ Process Concept
- ❑ Process Scheduling
- ❑ Operations on Processes
- ❑ Interprocess Communication
- ❑ IPC in Shared-Memory Systems
- ❑ **IPC in Message-Passing Systems**
  - **Naming**, Synchronization, Buffering
- ❑ Examples of IPC Systems
- ❑ Communication in Client-Server Systems

# Direct Communication (1/2)

- ❑ Each process that wants to communicate must explicitly name the recipient or sender of the communication

- Symmetry in addressing

- `send(P, message)` : send a message to process `P`
- `receive(Q, message)` : receive a message from process `Q`

- ❑ Properties of communication link

- A link is established automatically between every pair of processes that want to communicate
  - The processes need to know only each other's identity to communicate
- A link is associated with exactly two processes
- Between each pair of processes, there exists exactly one link

# Direct Communication (2/2)

- ❑ Each process that wants to communicate must explicitly name the recipient or sender of the communication
  - Asymmetry in addressing
    - `send(P, message)`: send a message to process `P`
    - `receive(id, message)`: receive a message from any process, where `id` is set to the name of the process with which communication has taken place
- ❑ Disadvantage (both symmetric and asymmetric): limited modularity
  - Changing the identifier of a process may necessitate examining all other process definitions
  - All references to the old identifier must be found and modified to the new identifier

# Indirect Communication (1/4)

- ❑ The messages are sent to and received from mailboxes (or ports)
  - **send(A, message)** : send a message to mailbox **A**
  - **receive(A, message)** : receive a message from mailbox **A**
  - A mailbox can be viewed abstractly as an object into which messages can be placed by processes and from which messages can be removed
  - Each mailbox has a unique identification
  - A process can communicate with another process via a number of different mailboxes
  - Two processes can communicate only if they have a shared mailbox

# Indirect Communication (2/4)

## ❑ Properties of indirect communication link

- A link is established between a pair of processes only if both members of the pair have a shared mailbox
- A link may be associated with more than two processes
- Between each pair of communicating processes, a number of different links may exist, with each link corresponding to one mailbox

## ❑ Properties of direct communication link (for comparison)

- A link is established automatically between every pair of processes that want to communicate
- A link is associated with exactly two processes
- Between each pair of processes, there exists exactly one link

# Indirect Communication (3/4)

## ❑ Processes $P_1$ , $P_2$ , and $P_3$ share mailbox **A**

- Process  $P_1$  sends a message to **A**
- Both  $P_2$  and  $P_3$  execute a **receive()** from **A**
- Which process will receive the message sent by  $P_1$ ?

## ❑ Some solutions

- Allow a link to be associated with two processes at most
- Allow at most one process at a time to execute a **receive()** operation
- Allow the system to select arbitrarily which process will receive the message
  - Define an algorithm for selecting which process will receive the message

# Indirect Communication (4/4)

## ❑ A mailbox owned by a process (the mailbox is part of the address space of the process)

- When a process that owns a mailbox terminates, the mailbox disappears
- Any process that subsequently sends a message to this mailbox must be notified that the mailbox no longer exists

## ❑ A mailbox owned by the operating system

- The operating system must allow a process to
  - Create a new mailbox
  - Send and receive messages through the mailbox
  - Delete a mailbox
- The process that creates a new mailbox is the mailbox's owner
- The ownership and receiving privilege may be passed to other processes through appropriate system calls



# Outline

- ❑ Process Concept
- ❑ Process Scheduling
- ❑ Operations on Processes
- ❑ Interprocess Communication
- ❑ IPC in Shared-Memory Systems
- ❑ **IPC in Message-Passing Systems**
  - Naming, **Synchronization**, Buffering
- ❑ Examples of IPC Systems
- ❑ Communication in Client-Server Systems

# Synchronization

❑ Message passing may be either blocking (synchronous) or nonblocking (asynchronous)

➤ Blocking send

- The sending process is blocked until the message is received by the receiving process or by the mailbox

➤ Nonblocking send

- The sending process sends the message and resumes operation

➤ Blocking receive

- The receiver blocks until a message is available

➤ Nonblocking receive

- The receiver retrieves either a valid message or a null

❑ Different combinations of **send()** and **receive()** are possible

➤ When both send() and receive() are blocking, we have a rendezvous

# Blocking Producer and Consumer

- ❑ The solution is trivial with blocking **send()** and **receive()**

- ❑ The producer process

```
message next_produced;  
while (true) {  
    /* produce an item in next_produced */  
    send(next_produced);  
}
```

- ❑ The consumer process

```
message next_consumed;  
while (true) {  
    receive(next_consumed);  
    /* consume the item in next_consumed */  
}
```

# Outline

- ❑ Process Concept
- ❑ Process Scheduling
- ❑ Operations on Processes
- ❑ Interprocess Communication
- ❑ IPC in Shared-Memory Systems
- ❑ **IPC in Message-Passing Systems**
  - Naming, Synchronization, **Buffering**
- ❑ Examples of IPC Systems
- ❑ Communication in Client-Server Systems

# Buffering

- ❑ Messages exchanged by communicating processes reside in a temporary queue
  - No matter communication is direct or indirect
- ❑ Basically, such queues can be implemented in three ways
  - Zero capacity
    - The sender must block until the recipient receives the message
  - Bounded capacity
    - If the link is full, the sender must block until space is available in the queue
  - Unbounded capacity
    - The sender never blocks

# Outline

- ❑ Process Concept
- ❑ Process Scheduling
- ❑ Operations on Processes
- ❑ Interprocess Communication
- ❑ IPC in Shared-Memory Systems
- ❑ IPC in Message-Passing Systems
- ❑ **Examples of IPC Systems**
  - **POSIX Shared Memory**, Mach Message Passing, Windows, Pipes
- ❑ Communication in Client-Server Systems

# POSIX Shared Memory

## ❑ Portable Operating System Interface (POSIX)

- A family of standards specified by the IEEE Computer Society for maintaining compatibility between operating systems [Wikipedia]
  - Variants of Unix and other operating systems

## ❑ Several IPC mechanisms are available for POSIX systems

## ❑ Here, we explore the POSIX API for shared memory

- Create a shared-memory object
  - `shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);`
- Configure the size of the object in bytes
  - `ftruncate(shm_fd, 4096);`
- Establish a memory-mapped file containing the shared-memory object and return a pointer to the memory-mapped file
  - `mmap()`

[https://man7.org/linux/man-pages/man3/shm\\_open.3.html](https://man7.org/linux/man-pages/man3/shm_open.3.html)

<https://man7.org/linux/man-pages/man2/mmap.2.html>

# Producer with POSIX Shared Memory

```
int main() {
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* strings written to shared memory */
    const char *message_0 = "Hello";
    const char *message_1 = "World!";
    /* shared memory file descriptor */
    int fd;
    /* pointer to shared memory object */
    char *ptr;
    /* create the shared memory object */
    fd = shm_open(name, O_CREAT | O_RDWR, 0666);
    /* configure the size of the shared memory object */
    ftruncate(fd, SIZE);
    /* memory map the shared memory object */
    ptr = (char *) mmap(0, SIZE, PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0);
    /* write to the shared memory object */
    sprintf(ptr,"%s",message_0);
    ptr += strlen(message_0);
    sprintf(ptr,"%s",message_1);
    ptr += strlen(message_1);
    return 0;
}
```



# Consumer with POSIX Shared Memory

```
int main() {
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* shared memory file descriptor */
    int fd;
    /* pointer to shared memory object */
    char *ptr;
    /* open the shared memory object */
    fd = shm_open(name, O_RDONLY, 0666);
    /* memory map the shared memory object */
    ptr = (char *) mmap(0, SIZE, PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0);
    /* read from the shared memory object */
    printf("%s", (char *)ptr);
    /* remove the shared memory object */
    shm_unlink(name);
    return 0;
}
```

# Outline

- ❑ Process Concept
- ❑ Process Scheduling
- ❑ Operations on Processes
- ❑ Interprocess Communication
- ❑ IPC in Shared-Memory Systems
- ❑ IPC in Message-Passing Systems
- ❑ **Examples of IPC Systems**
  - POSIX Shared Memory, **Mach Message Passing**, Windows, Pipes
- ❑ Communication in Client-Server Systems

# Mach Message Passing (1/3)

- ❑ Mach was included in the macOS and iOS operating systems
  - Mach was initially designed for distributed systems
- ❑ Most communication in Mach, including all inter-task communication, is carried out by messages
  - Create a new **port** (mailbox) and allocate space for its message queue
    - `mach_port_allocate()`
  - Send and receive messages
    - `mach_msg()`
- ❑ Two special ports are created with a task (similar to a process)
  - The kernel has receive rights to the **Task Self** port
    - The task can send messages to the kernel
  - The task has receive rights to the **Notify** port
    - The kernel can send notification of event occurrences to the task

# Mach Message Passing (2/3)

## ❑ Two message fields

- A fixed-size message header and a variable-sized body

```
struct message {  
    mach_msg_header_t header;  
    int data;  
};
```

## ❑ The send and receive operations themselves are flexible

- Example: when a message is sent to a port, if the port's queue is full, the sender has several options (specified via parameters to `mach_msg()`)
  - Wait indefinitely until there is room in the queue
  - Wait at most *n* milliseconds
  - Do not wait at all but rather return immediately
  - Temporarily cache a message

# Mach Message Passing (3/3)

```
mach_port_t client;
mach_port_t server;

/* Client Code */
struct message message;
// construct the header
message.header.msgh_size = sizeof(message);
message.header.msgh_remote_port = server;
message.header.msgh_local_port = client;
// send the message
mach_msg(&message.header, // message header
        MACH_SEND_MSG, // sending a message
        sizeof(message), // size of message sent
        0, // maximum size of received message - unnecessary
        MACH_PORT_NULL, // name of receive port - unnecessary
        MACH_MSG_TIMEOUT_NONE, // no time outs
        MACH_PORT_NULL // no notify port
);

/* Server Code */
struct message message;
// receive the message
mach_msg(&message.header, // message header
        MACH_RCV_MSG, // sending a message --> receiving a message
        0, // size of message sent
        sizeof(message), // maximum size of received message
        server, // name of receive port
        MACH_MSG_TIMEOUT_NONE, // no time outs
        MACH_PORT_NULL // no notify port
);
```

# Outline

- ❑ Process Concept
- ❑ Process Scheduling
- ❑ Operations on Processes
- ❑ Interprocess Communication
- ❑ IPC in Shared-Memory Systems
- ❑ IPC in Message-Passing Systems
- ❑ **Examples of IPC Systems**
  - POSIX Shared Memory, Mach Message Passing, **Windows**, Pipes
- ❑ Communication in Client-Server Systems


# Windows

## ❑ Support multiple operating environments (subsystems)

- Application programs communicate with these subsystems via a message-passing mechanism
- Application programs can be considered clients of a subsystem server

## ❑ Advanced local procedure call (ALPC) facility

- The client opens a handle (abstract reference) to the server's **connection-port** object and sends a connection request to that port
- The server creates a channel and returns a handle to the client
  - The channel consists of two private **communication ports** for client-server messages and for server-client messages



# Outline

- ❑ Process Concept
- ❑ Process Scheduling
- ❑ Operations on Processes
- ❑ Interprocess Communication
- ❑ IPC in Shared-Memory Systems
- ❑ IPC in Message-Passing Systems
- ❑ **Examples of IPC Systems**
  - POSIX Shared Memory, Mach Message Passing, Windows, **Pipes**
- ❑ Communication in Client-Server Systems



# Pipes

## ❑ A conduit allowing two processes to communicate

- Pipes were one of the first IPC mechanisms in early UNIX systems

## ❑ Simple but with some issues

- Bidirectional or unidirectional?
- If bidirectional, half duplex (only one direction at a time) or full duplex (both directions at a time)?
- Must a relationship (such as parent-child) exist between the communicating processes?
- Can the pipes communicate over a network, or must the communicating processes reside on the same machine?

## ❑ Two common types used on both UNIX and Windows systems

- Ordinary pipes
- Named pipes

# Ordinary Pipes

- ❑ Ordinary pipes allow two processes to communicate in standard producer-consumer fashion
  - The producer writes to one end of the pipe (the write end)
  - The consumer reads from the other end (the read end)
- ❑ An ordinary pipe is unidirectional
- ❑ An ordinary pipe cannot be accessed from outside the process that created it
- ❑ Ordinary pipes on Windows are termed anonymous pipes

# Ordinary Pipes in UNIX

```
#define BUFFER_SIZE 25
#define READ_END 0
#define WRITE_END 1
int main(void) {
    char write_msg[BUFFER_SIZE] = "Greetings";
    char read_msg[BUFFER_SIZE];
    int fd[2];
    pid_t pid;
    pipe(fd); /* create the pipe */
    /* omit codes if pipe failed */
    pid = fork(); /* fork a child process */
    /* omit codes if fork failed */
    if (pid > 0) { /* parent process */
        close(fd[READ_END]); /* close the unused end of the pipe */
        write(fd[WRITE_END], write_msg, strlen(write_msg)+1); /* write to the pipe */
        close(fd[WRITE_END]); /* close the write end of the pipe */
    }
    else { /* child process */
        close(fd[WRITE_END]); /* close the unused end of the pipe */
        read(fd[READ_END], read_msg, BUFFER_SIZE); /* read from the pipe */
        printf("read %s", read_msg);
        close(fd[READ_END]); /* close the read end of the pipe */
    }
    return 0;
}
```

# Named Pipes


- ❑ Named pipes provide a much more powerful communication tool
  - Communication can be bidirectional
  - No parent-child relationship is required
  - Several processes can use a named pipe for communication
    - In a typical scenario, a named pipe has several writers
  - Named pipes continue to exist after communicating processes have finished
- ❑ Both UNIX and Windows systems support named pipes
  - Half-duplex in UNIX
  - Full-duplex in Windows

# Outline

- ❑ Process Concept
- ❑ Process Scheduling
- ❑ Operations on Processes
- ❑ Interprocess Communication
- ❑ IPC in Shared-Memory Systems
- ❑ IPC in Message-Passing Systems
- ❑ Examples of IPC Systems
- ❑ **Communication in Client-Server Systems**
  - **Sockets**, Remote Procedure Calls

# Sockets

- ❑ A **socket** is defined as an endpoint for communication
  - A socket is identified by an IP address concatenated with a port number
- ❑ Sockets use a client-server architecture
  - A server waits for client requests by listening to a specified port
  - Once a request is received, the server accepts a connection from the client socket to complete the connection
- ❑ All ports below 1024 are considered **well-known**
  - SSH: port 22; FTP: port 21; HTTP: port 80
- ❑ Communication using a pair of sockets
  - (146.86.5.20:1625) on host X and (161.25.19.8:80) on the web server



# Sockets in Java

## ❑ Three types of sockets

- Connection-oriented (TCP) sockets
  - `Socket` class
- Connectionless (UDP) sockets
  - `DatagramSocket` class
- Multicast sockets, allowing data to be sent to multiple recipients
  - `MulticastSocket` class, a subclass of the `DatagramSocket` class

# Date Server in Java

```
import java.net.*;
import java.io.*;
public class DateServer
{
    public static void main(String[] args) {
        try {
            ServerSocket sock = new ServerSocket(6013);
            /* now listen for connections */
            while (true) {
                Socket client = sock.accept();
                PrintWriter pout = new
                    PrintWriter(client.getOutputStream(), true);
                /* write the Date to the socket */
                pout.println(new java.util.Date().toString());
                /* close the socket and resume */
                /* listening for connections */
                client.close();
            }
        }
        catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```



# Date Server in Java

- ❑ The server listens to the port with the **accept()** method
  - The server blocks on the **accept()** method waiting for a client to request a connection
  - When a connection request is received, **accept()** returns a socket that the server can use to communicate with the client

```
        while (true) {
            Socket client = sock.accept();
            PrintWriter pout = new
                PrintWriter(client.getOutputStream(), true);
            /* write the Date to the socket */
            pout.println(new java.util.Date().toString());
            /* close the socket and resume */
            /* listening for connections */
            client.close();
        }
    }
    catch (IOException ioe) {
        System.err.println(ioe);
    }
}
```

# Date Client in Java

```
import java.net.*;
import java.io.*;
public class DateClient
{
    public static void main(String[] args) {
        try {
            /* make connection to server socket */
            Socket sock = new Socket("127.0.0.1",6013);
            InputStream in = sock.getInputStream();
            BufferedReader bin = new
                BufferedReader(new InputStreamReader(in));
            /* read the date from the socket */
            String line;
            while ( (line = bin.readLine()) != null)
                System.out.println(line);
            /* close the socket connection*/
            sock.close();
        }
        catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```

# Outline

- ❑ Process Concept
- ❑ Process Scheduling
- ❑ Operations on Processes
- ❑ Interprocess Communication
- ❑ IPC in Shared-Memory Systems
- ❑ IPC in Message-Passing Systems
- ❑ Examples of IPC Systems
- ❑ **Communication in Client-Server Systems**
  - Sockets, Remote Procedure Calls

# Remote Procedure Calls (1/2)

- ❑ Remote procedure call (RPC) abstracts the procedure-call mechanism between systems with network connections
  - Allow a client to invoke a procedure on a remote host as it would invoke a procedure locally
- ❑ The RPC system provides a stub on the client side
  - When the client invokes a remote procedure, the RPC system calls the appropriate stub (with parameters)
  - This stub locates the port on the server and marshals the parameters
  - The stub transmits a message to the server using message passing
  - A similar stub on the server side receives this message and invokes the procedure on the server
    - If necessary, return values are passed back to the client using the same technique
- ❑ Microsoft Interface Definition Language (MIDL) on windows

# Remote Procedure Calls (2/2)

## ❑ Many RPC systems define a machine-independent representation of data

### ➤ Example: external data representation (XDR)

- On the client side, parameter marshaling involves converting the machine-dependent data into XDR
- On the server side, the XDR data are unmarshaled and converted to the machine-dependent representation

## ❑ Remote communication has more failure scenarios

### ➤ Messages can be acted on exactly once, rather than at most once

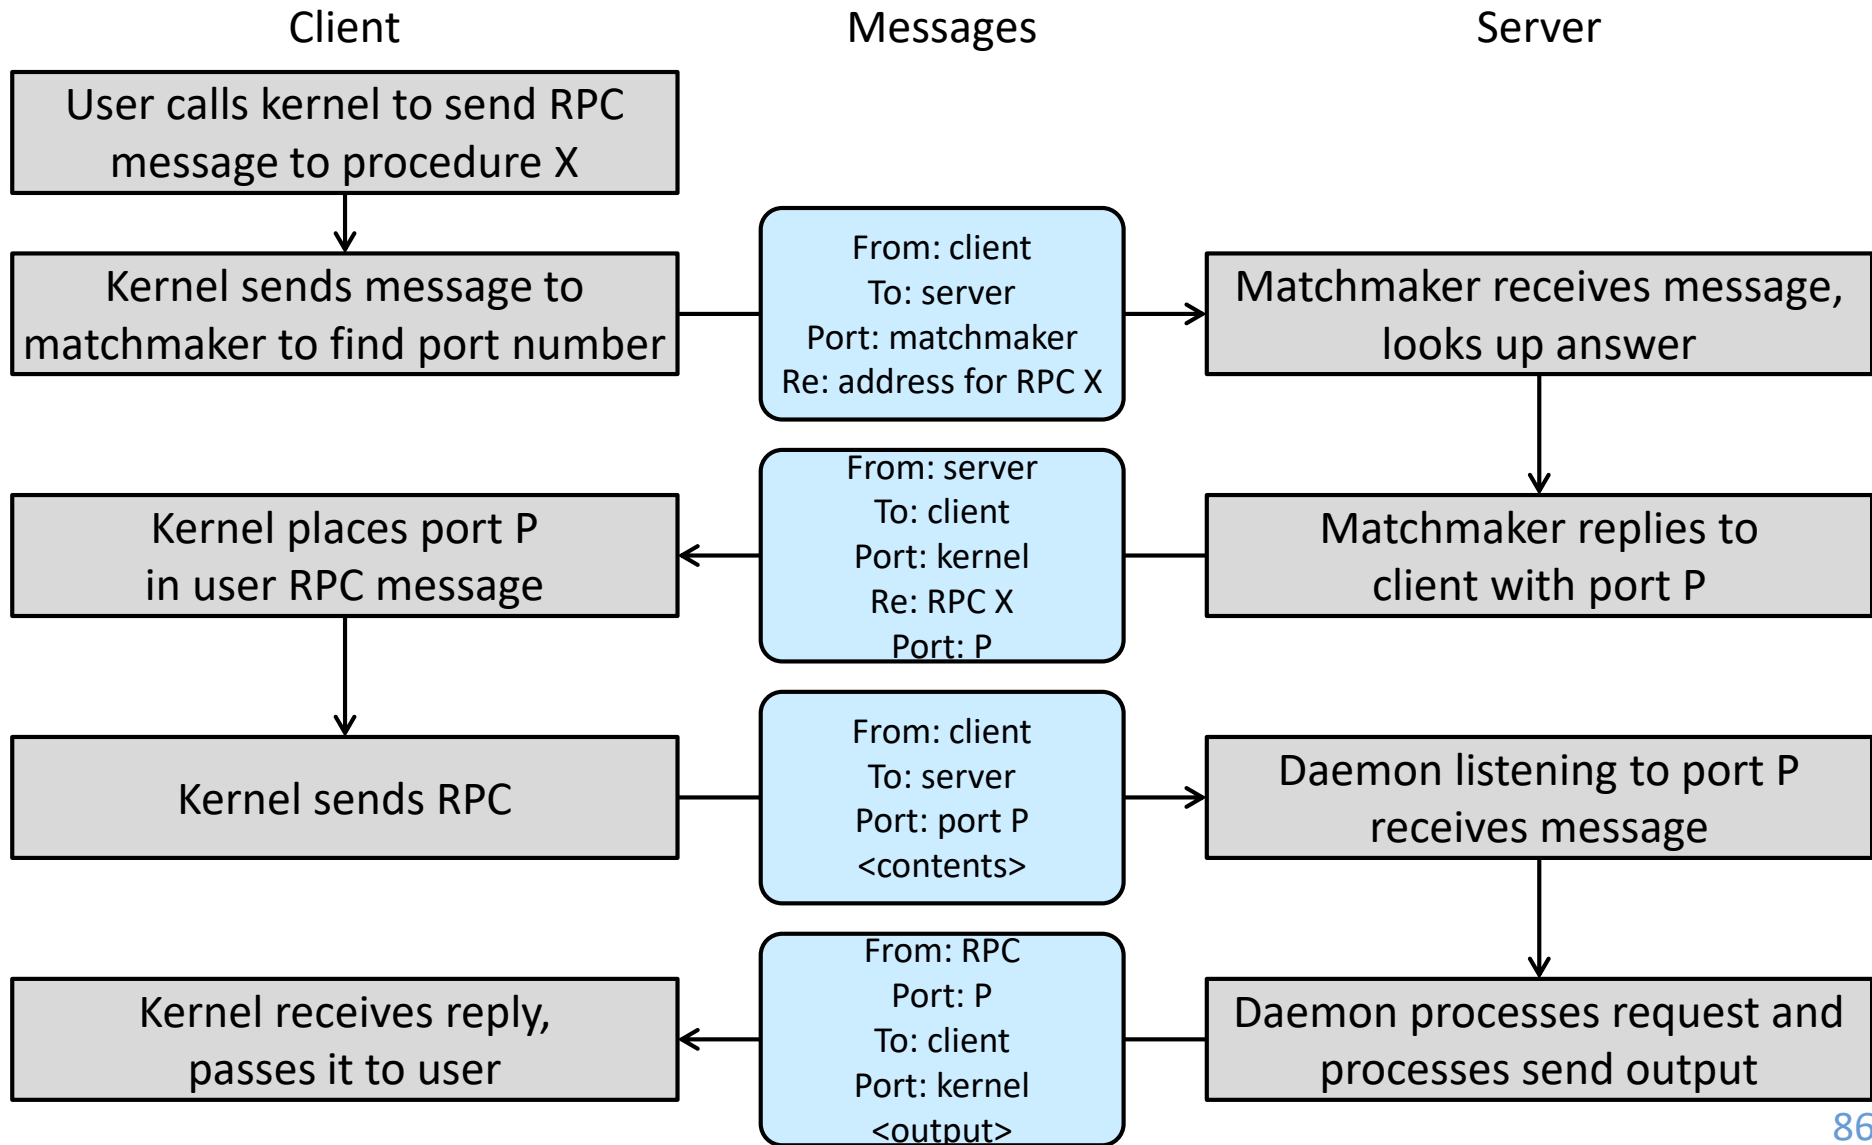
## ❑ Binding

### ➤ Predetermined

### ➤ Dynamic

- An operating system provides a rendezvous (also called a matchmaker) daemon on a fixed RPC port

# Dynamic Binding and Execution of RPC



# Objectives

- ❑ Identify the separate components of a process and illustrate how they are represented and scheduled in an operating system
- ❑ Describe how processes are created and terminated in an operating system, including developing programs using the appropriate system calls that perform these operations
- ❑ Describe and contrast interprocess communication using shared memory and message passing
- ❑ Design programs that use pipes and POSIX shared memory to perform interprocess communication
- ❑ Describe client-server communication using sockets and remote procedure calls

# Q&A