

Operating Systems

[6. Synchronization Tools]

Chung-Wei Lin

cwlin@csie.ntu.edu.tw

CSIE Department

National Taiwan University

Objectives

- ❑ Describe the critical-section problem and illustrate a race condition
- ❑ Illustrate hardware solutions to the critical-section problem using memory barriers, compare-and-swap operations, and atomic variables
- ❑ Demonstrate how mutex locks, semaphores, monitors, and condition variables can be used to solve the critical section problem
- ❑ Evaluate tools that solve the critical-section problem in low-, moderate-, and high-contention scenarios

Outline

- ❑ **Background**
- ❑ The Critical-Section Problem
- ❑ Peterson's Solution
- ❑ Hardware Support for Synchronization
- ❑ Mutex Locks
- ❑ Semaphores
- ❑ Monitors
- ❑ Liveness
- ❑ Evaluation

Background

❑ Processes can execute concurrently

- They may be interrupted at any time, partially completing execution

❑ Concurrent access to shared data may result in data inconsistency

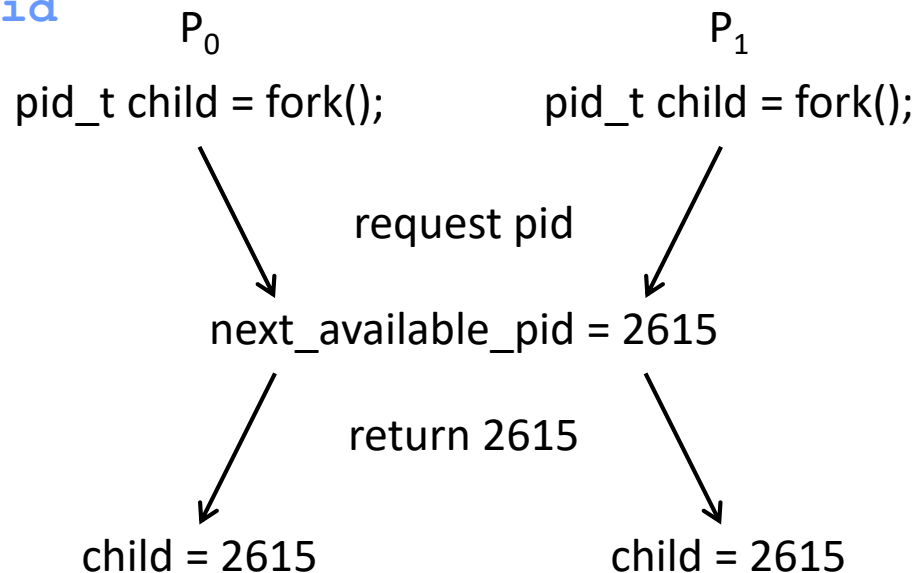
- Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes

❑ The "bounded buffer" problem

- Use of a counter that is updated concurrently by the producer and consumer leads to race condition
- Process 1: **count++**
 - `reg1 = count; reg1 = reg1 + 1; count = reg1;`
- Process 2: **count--**
 - `reg2 = count; reg2 = reg2 - 1; count = reg2;`
- What if the lower-level statements are interleaved in some order?

Race Condition

- ❑ Processes access and manipulate the same data concurrently and the outcome depends on the accessing order
 - Processes P_0 and P_1 are creating child processes using `fork()`
 - Race condition on kernel variable `next_available_pid` which represents the next available process identifier `pid`
 - The same `pid` could be assigned to two different processes
 - Unless there is a mechanism to prevent P_0 and P_1 from accessing the variable `next_available_pid`



Outline

- ❑ Background
- ❑ **The Critical-Section Problem**
- ❑ Peterson's Solution
- ❑ Hardware Support for Synchronization
- ❑ Mutex Locks
- ❑ Semaphores
- ❑ Monitors
- ❑ Liveness
- ❑ Evaluation

Critical-Section Problem

- ❑ Consider a system of n processes $\{p_0, p_1, \dots, p_{n-1}\}$
- ❑ Each process has critical section segment of code
 - The process may be accessing and updating data that is shared with at least one other process
 - When one process is executing in its critical section, no other process is allowed to execute in its critical section
- ❑ Critical section problem is to design protocol to solve this
- ❑ Each process
 - Must ask permission to enter critical section in entry section
 - May follow critical section with exit section and then remainder section

General Structure of Process

```
Do {  
    entry section  
    critical section  
    exit section  
    remainder section  
} while (true)
```


Solution Requirements (1/2)

❑ Requirement 1: Mutual Exclusion

- If a process is executing in its critical section
- Then no other processes can be executing in their critical sections

❑ Requirement 2: Progress

- If
 - No process is executing in its critical section, and
 - Some processes wish to enter their critical sections
- Then
 - Only those processes that are not executing in their remainder sections can participate in deciding which will enter its critical section next, and
 - This selection cannot be postponed indefinitely

Solution Requirements (2/2)

❑ Requirement 3: Bounded Waiting

- There exists a bound, or limit, on the number of times that other processes are allowed to enter their critical sections
 - After a process has made a request to enter its critical section, and
 - Before that request is granted

❑ Assumption

- We assume that each process is executing at a nonzero speed
- However, we can make no assumption concerning the relative speed of the n processes

Interrupt-Based Solution

☐ Entry section

- Disable interrupts

☐ Exit section

- Enable interrupts

☐ Will this solve the problem?

- What if the critical section is code that runs for an hour?
- Can some processes starve (never enter their critical section)?
- What if there are two CPUs?

Software Solution (1/2)

❑ Two-process solution

- Assume that the **load** and **store** machine-language instructions are atomic (cannot be interrupted)
- The two processes share **turn** to indicate whose turn

❑ Algorithm for process P_i

```
while (true) {  
    turn = i;  
    while (turn == j)  
        ;  
    /* critical section */  
    turn = j;  
    /* remainder section */  
}
```

Software Solution (2/2)

```
while (true){  
    while (turn == j)  
        ;  
    /* critical section */  
    turn = j;  
    /* remainder section */  
}
```

☐ Mutual exclusion is preserved?

➤ **turn** cannot be both 0 and 1 at the same time, but ...

☐ What about the progress requirement?

☐ What about the bounded-waiting requirement?

☐ What if we remove "**turn = i**"?

Outline

- ❑ Background
- ❑ The Critical-Section Problem
- ❑ **Peterson's Solution**
- ❑ Hardware Support for Synchronization
- ❑ Mutex Locks
- ❑ Semaphores
- ❑ Monitors
- ❑ Liveness
- ❑ Evaluation

Peterson's Solution (1/2)

❑ Two-process solution

- Assume that the **load** and **store** machine-language instructions are atomic (cannot be interrupted)
- The two processes share: **int turn** and **boolean flag[2]**
 - **turn** indicates whose turn it is to enter the critical section
 - **flag** array indicates if a process is ready to enter the critical section
 - **flag[i] = true** implies that process P_i is ready

Peterson's Solution (2/2)

❑ Algorithm for process P_i

```
while (true){
```

```
    flag[i] = true;
```

```
    turn = j;
```

```
    while (flag[j] && turn == j)
```

```
        ;
```

```
    /* critical section */
```

```
    flag[i] = false;
```

```
    /* remainder section */
```

```
}
```

For Comparison

```
while (true){
```

```
    turn = i;
```

```
    while (turn == j)
```

```
        ;
```

```
    /* critical section */
```

```
    turn = j;
```

```
    /* remainder section */
```

```
}
```


Peterson's Solution: Correctness

```
while (true){  
    flag[i] = true;  
    turn = j;  
    while (flag[j] && turn == j)  
        ;  
    /* critical section */  
    flag[i] = false;  
    /* remainder section */  
}
```

□ Provable that the three requirements are met

- Mutual exclusion is preserved
- The progress requirement is satisfied
- The bounded-waiting requirement is met

Modern Architecture

❑ Peterson's solution is not guaranteed to work on modern computer architectures

- To improve system performance, processors and/or compilers may reorder read and write operations that have no dependencies
 - For a single-threaded application, the reordering is fine as the final values are consistent with what is expected
 - For a multi-threaded application with shared data, the reordering may render inconsistent or unexpected results

Modern Architecture: Example (1/2)

- ❑ Two threads share the data

```
boolean flag = false;  
int x = 0;
```

- ❑ Thread 1 performs

```
while (!flag)  
;  
print x
```

- ❑ Thread 2 performs

```
x = 100;  
flag = true;
```

- ❑ What is the expected output?

100

Modern Architecture: Example (2/2)

- ❑ However, since the variables **flag** and **x** are independent of each other, the following instructions for Thread 2 may be reordered


```
flag = true;
```

```
x = 100;
```

- ❑ If this occurs, the output may be 0

Peterson's Solution Revisited

❑ Instruction reordering in Peterson's Solution



- ❑ This allows both processes to be in their critical section at the same time

Outline

- ❑ Background
- ❑ The Critical-Section Problem
- ❑ Peterson's Solution
- ❑ **Hardware Support for Synchronization**
 - **Memory Barriers**
 - Hardware Instructions
 - Atomic Variables
- ❑ Mutex Locks
- ❑ Semaphores
- ❑ Monitors
- ❑ Liveness
- ❑ Evaluation

Memory Barriers (1/2)

❑ Memory model

- How a computer architecture determines what memory guarantees it will provide to an application program

❑ Memory models may be either

➤ Strongly ordered

- A memory modification on one processor is immediately visible to all other processors

➤ Weakly ordered

- Modifications to memory on one processor may not be immediately visible to other processors

❑ A memory barrier is an instruction that forces any change in memory to be propagated (made visible) to all other processors

Memory Barriers (2/2)

❑ When a memory barrier instruction is performed

- The system ensures that all loads and stores are completed before any subsequent load or store operations are performed

❑ Therefore, even with instructions reordering, the memory barrier ensures that

- The store operations are completed in memory and visible to other processors before future load or store operations are performed

❑ Note that

- Memory barriers are considered very low-level operations
 - They are typically only used by kernel developers when writing specialized code that ensures mutual exclusion

Memory Barriers: Example

- ❑ Add memory barriers to the previous example

- ❑ Thread 1 performs

```
while (!flag)
    memory_barrier();
print x
```

- ❑ Thread 2 performs

```
x = 100;
memory_barrier();
flag = true
```

- ❑ We are guaranteed that

- For Thread 1, the value of **flag** is loaded before the value of **x**
- For Thread 2, the assignment to **x** occurs before the assignment to **flag**

Outline

- ❑ Background
- ❑ The Critical-Section Problem
- ❑ Peterson's Solution
- ❑ **Hardware Support for Synchronization**
 - Memory Barriers
 - **Hardware Instructions**
 - Atomic Variables
- ❑ Mutex Locks
- ❑ Semaphores
- ❑ Monitors
- ❑ Liveness
- ❑ Evaluation

Hardware Instructions

- ❑ Many modern computer systems provide special hardware instructions that allow us either
 - To test and modify the content of a word, or
 - `test_and_set`
 - To swap the contents of two words atomically (uninterruptedly)
 - `compare_and_swap` (CAS)

test_and_set Instruction

□ Definition

```
boolean test_and_set (boolean *target) {  
    boolean rv = *target;  
    *target = true;  
    return rv;  
}
```

□ Properties

- Be executed atomically
- Return the original value of the passed parameter
- Set the value of the passed parameter to **true**

Solution Using `test_and_set`

- ❑ A shared boolean variable `lock`, initialized to `false`

- ❑ Solution

```
do {  
    while (test_and_set(&lock))  
        ; /* do nothing */  
    /* critical section */  
    lock = false;  
    /* remainder section */  
} while (true);
```

- ❑ Does it solve the critical-section problem?

compare_and_swap Instruction

□ Definition

```
int compare_and_swap(int *value, int expected,
int new_value) {
    int temp = *value;
    if (*value == expected)
        *value = new_value;
    return temp;
}
```

□ Properties

- Be executed atomically
- Return the original value of the passed parameter **value**
- If ***value == expected**, set the value of the passed variable **value** to the passed parameter **new_value**
 - "Swap" takes place only under this condition

Solution Using `compare_and_swap`

- ❑ A shared integer `lock`, initialized to 0

- ❑ Solution

```
while (true) {  
    while (compare_and_swap(&lock, 0, 1) != 0)  
        ; /* do nothing */  
    /* critical section */  
    lock = 0;  
    /* remainder section */  
}
```

- ❑ Does it solve the critical-section problem?

Bounded-Waiting with CAS

```
while (true) {
    waiting[i] = true;
    key = 1;
    while (waiting[i] && key == 1)
        key = compare_and_swap(&lock, 0, 1);
    waiting[i] = false;
    /* critical section */
    j = (i + 1) % n;
    while ((j != i) && !waiting[j])
        j = (j + 1) % n;
    if (j == i)
        lock = 0;
    else
        waiting[j] = false;
    /* remainder section */
}
```


Outline

- ❑ Background
- ❑ The Critical-Section Problem
- ❑ Peterson's Solution
- ❑ **Hardware Support for Synchronization**
 - Memory Barriers
 - Hardware Instructions
 - **Atomic Variables**
- ❑ Mutex Locks
- ❑ Semaphores
- ❑ Monitors
- ❑ Liveness
- ❑ Evaluation

Atomic Variables

- ❑ Typically, instructions such as **compare-and-swap** are used as building blocks for other synchronization tools
 - Not used directly to provide mutual exclusion
- ❑ One such tool is an **atomic variable**
 - Provide atomic (uninterruptible) operations on basic data types such as integers and booleans

Atomic Variables: Example

❑ Example

- Let **sequence** be an atomic variable
- Let **increment()** be an operation on **sequence**
- The command **increment(&sequence)** ensures **sequence** is incremented without interruption

❑ **increment()** function

```
void increment(atomic_int *v) {  
    int temp;  
    do {  
        temp = *v;  
    }  
    while (temp != (compare_and_swap(v, temp, temp+1)) );  
}
```

Outline

- ❑ Background
- ❑ The Critical-Section Problem
- ❑ Peterson's Solution
- ❑ Hardware Support for Synchronization
- ❑ **Mutex Locks**
- ❑ Semaphores
- ❑ Monitors
- ❑ Liveness
- ❑ Evaluation

Mutex (Mutual Exclusion) Locks (1/2)

- ❑ Previous solutions are complicated and generally inaccessible to application programmers
 - OS designers build software tools to solve the critical section problem
- ❑ A mutex lock protects critical sections and thus prevent race conditions
 - First **acquire()** a lock and then **release()** the lock
 - Calls to **acquire()** and **release()** must be atomic; usually implemented via hardware atomic instructions such as compare_and_swap
 - A boolean variable **available** indicating if the lock is available or not
- ❑ This solution requires busy waiting
 - While a process is in its critical section, another process trying to enter its critical section must loop continuously in the call to **acquire()**
 - This lock therefore called a spinlock

Mutex Locks (2/2)

```
while (true){  
    acquire lock  
    /* critical section */  
    release lock  
    /* remainder section */  
}  
  
acquire() {  
    while (!available)  
        ; /* busy wait */  
    available = false;  
}  
  
release() {  
    available = true;  
}
```

Outline

- ❑ Background
- ❑ The Critical-Section Problem
- ❑ Peterson's Solution
- ❑ Hardware Support for Synchronization
- ❑ Mutex Locks
- ❑ **Semaphores**
 - Semaphore Usage
 - Semaphore Implementation
- ❑ Monitors
- ❑ Liveness
- ❑ Evaluation

Semaphore

- ❑ A more robust tool that can
 - Behave similarly to a mutex lock
 - Provide more sophisticated ways for processes to synchronize activities
- ❑ A semaphore S is an integer variable that, apart from initialization, is accessed only through two atomic operations
- ❑ **wait()** operation (originally called **P()**)

```
wait(S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}
```

- ❑ **signal()** operation (originally called **V()**)
 signal(S) { S++; }

Outline

- ❑ Background
- ❑ The Critical-Section Problem
- ❑ Peterson's Solution
- ❑ Hardware Support for Synchronization
- ❑ Mutex Locks
- ❑ **Semaphores**
 - **Semaphore Usage**
 - Semaphore Implementation
- ❑ Monitors
- ❑ Liveness
- ❑ Evaluation

Semaphore Usage

❑ Counting semaphore

- Range over an unrestricted domain

❑ Binary semaphore

- Range only between 0 and 1
 - Same as a mutex lock
- Can implement a counting semaphore S as a binary semaphore

❑ We can solve various synchronization problems with semaphores

Semaphore Usage: Examples

- ❑ Create a semaphore "mutex" initialized to 1

```
wait(mutex) ;  
/* critical section */  
signal(mutex) ;
```

- ❑ Consider P_1 and P_2 with two statements S_1 and S_2 , respectively, and require that S_1 happens before S_2

- Create a semaphore "synch" initialized to 0

```
P1:  
    S1 ;  
    signal(synch) ;
```

```
P2:  
    wait(synch) ;  
    S2 ;
```

Outline

- ❑ Background
- ❑ The Critical-Section Problem
- ❑ Peterson's Solution
- ❑ Hardware Support for Synchronization
- ❑ Mutex Locks
- ❑ **Semaphores**
 - Semaphore Usage
 - **Semaphore Implementation**
- ❑ Monitors
- ❑ Liveness
- ❑ Evaluation

Semaphore Implementation

- ❑ No two processes can execute the **wait()** and **signal()** on the same semaphore at the same time
 - Thus, the implementation becomes the critical section problem where the **wait()** and **signal()** code are placed in the critical section
- ❑ The definitions of the **wait()** and **signal()** semaphore operations have the busy-waiting problem
 - Applications may spend lots of time in critical sections

Semaphore without Busy Waiting (1/4)

❑ Modify **wait()** and **signal()** operations

- When a process executes the **wait()** operation and finds that the semaphore value is not positive, it must wait
 - The process can suspend itself, rather than engaging in busy waiting
 - The suspend operation places a process into a waiting queue associated with the semaphore
 - The state of the process is switched to the waiting state
- Control is transferred to the CPU scheduler (selecting another process)
- A process that is suspended, waiting on a semaphore S, should be restarted when some other process executes a **signal()** operation
 - The process is restarted by a **wakeup()** operation, which changes the process from the waiting state to the ready state
 - The process is then placed in the ready queue

Semaphore without Busy Waiting (2/4)

- Each semaphore has an integer **value** and a list of processes **list**

```
typedef struct {  
    int value;  
    struct process *list;  
} semaphore;
```

Semaphore without Busy Waiting (3/4)

- ❑ When a process must wait on a semaphore, it is added to the list of processes
- ❑ **sleep()** suspends the process that invokes it

```
wait(semaphore *S) {  
    S->value--;  
    if (S->value < 0) {  
        add this process to S->list;  
        sleep();  
    }  
}
```


Semaphore without Busy Waiting (4/4)

- ❑ **signal ()** removes one process from the list of waiting processes and awakens that process
- ❑ **wakeup (P)** resumes the execution of a suspended process P

```
signal(semaphore *S) {  
    S->value++;  
    if (S->value <= 0) {  
        remove a process P from S->list;  
        wakeup(P) ;  
    }  
}
```

Semaphore Discussion

- ❑ Semaphore operations must be executed atomically
- ❑ We can use a FIFO queue to ensure bounded waiting
 - However, in general, the list can use any queuing strategy
- ❑ We have not completely eliminated busy waiting with this definition of the **wait()** and **signal()** operations
 - Rather, we have moved busy waiting from the entry section to the critical sections of application programs
 - We have limited busy waiting to the critical sections of the **wait()** and **signal()** operations
 - These sections are usually short

Outline

- ❑ Background
- ❑ The Critical-Section Problem
- ❑ Peterson's Solution
- ❑ Hardware Support for Synchronization
- ❑ Mutex Locks
- ❑ Semaphores
- ❑ **Monitors**
 - Monitor Usage
 - Implementing a Monitor Using Semaphores
 - Resuming Processes within a Monitor
- ❑ Liveness
- ❑ Evaluation

Monitors

- ❑ Unfortunately, timing errors can still occur when either mutex locks or semaphores are used
- ❑ Examples
 - `signal(mutex); ... critical section ... wait(mutex);`
 - Several processes may be executing in their critical sections simultaneously
 - `wait(mutex); ... critical section ... wait(mutex);`
 - The process will permanently block on the second call to `wait()`
 - Omit the `wait(mutex)`, or the `signal(mutex)`, or both
 - Either mutual exclusion is violated or the process will permanently block
 - They can be generated easily when programmers use semaphores or mutex locks incorrectly to solve the critical-section problem
 - Using them incorrectly can result in timing errors that are difficult to detect

Outline

- ❑ Background
- ❑ The Critical-Section Problem
- ❑ Peterson's Solution
- ❑ Hardware Support for Synchronization
- ❑ Mutex Locks
- ❑ Semaphores
- ❑ **Monitors**
 - **Monitor Usage**
 - Implementing a Monitor Using Semaphores
 - Resuming Processes within a Monitor
- ❑ Liveness
- ❑ Evaluation

Monitor Usage


- ❑ A high-level abstraction that provides a convenient and effective mechanism for process synchronization
- ❑ **Abstract data type (ADT)**
 - Encapsulate data with a set of functions to operate on that data that are independent of any specific implementation of the ADT
- ❑ **Monitor type**
 - An ADT that includes a set of programmer-defined operations that are provided with mutual exclusion within the monitor
 - A function defined within a monitor can access only those variables declared locally within the monitor and its formal parameters
 - The local variables of a monitor can be accessed by only the local functions

Pseudocode Syntax of a Monitor

```
monitor monitor-name {  
    /* shared variable declarations */  
    procedure P1 (...) {...}  
    procedure P2 (...) {...}  
    procedure Pn (...) {...}  
    initialization_code (...) {...}  
}
```

Schematic View of a Monitor

- ❑ The monitor construct ensures that only one process at a time is active within the monitor



Condition Variables (1/4)

❑ `condition x, y;`

❑ The only operations that can be invoked on a condition variable are `wait()` and `signal()`

➤ `x.wait()`

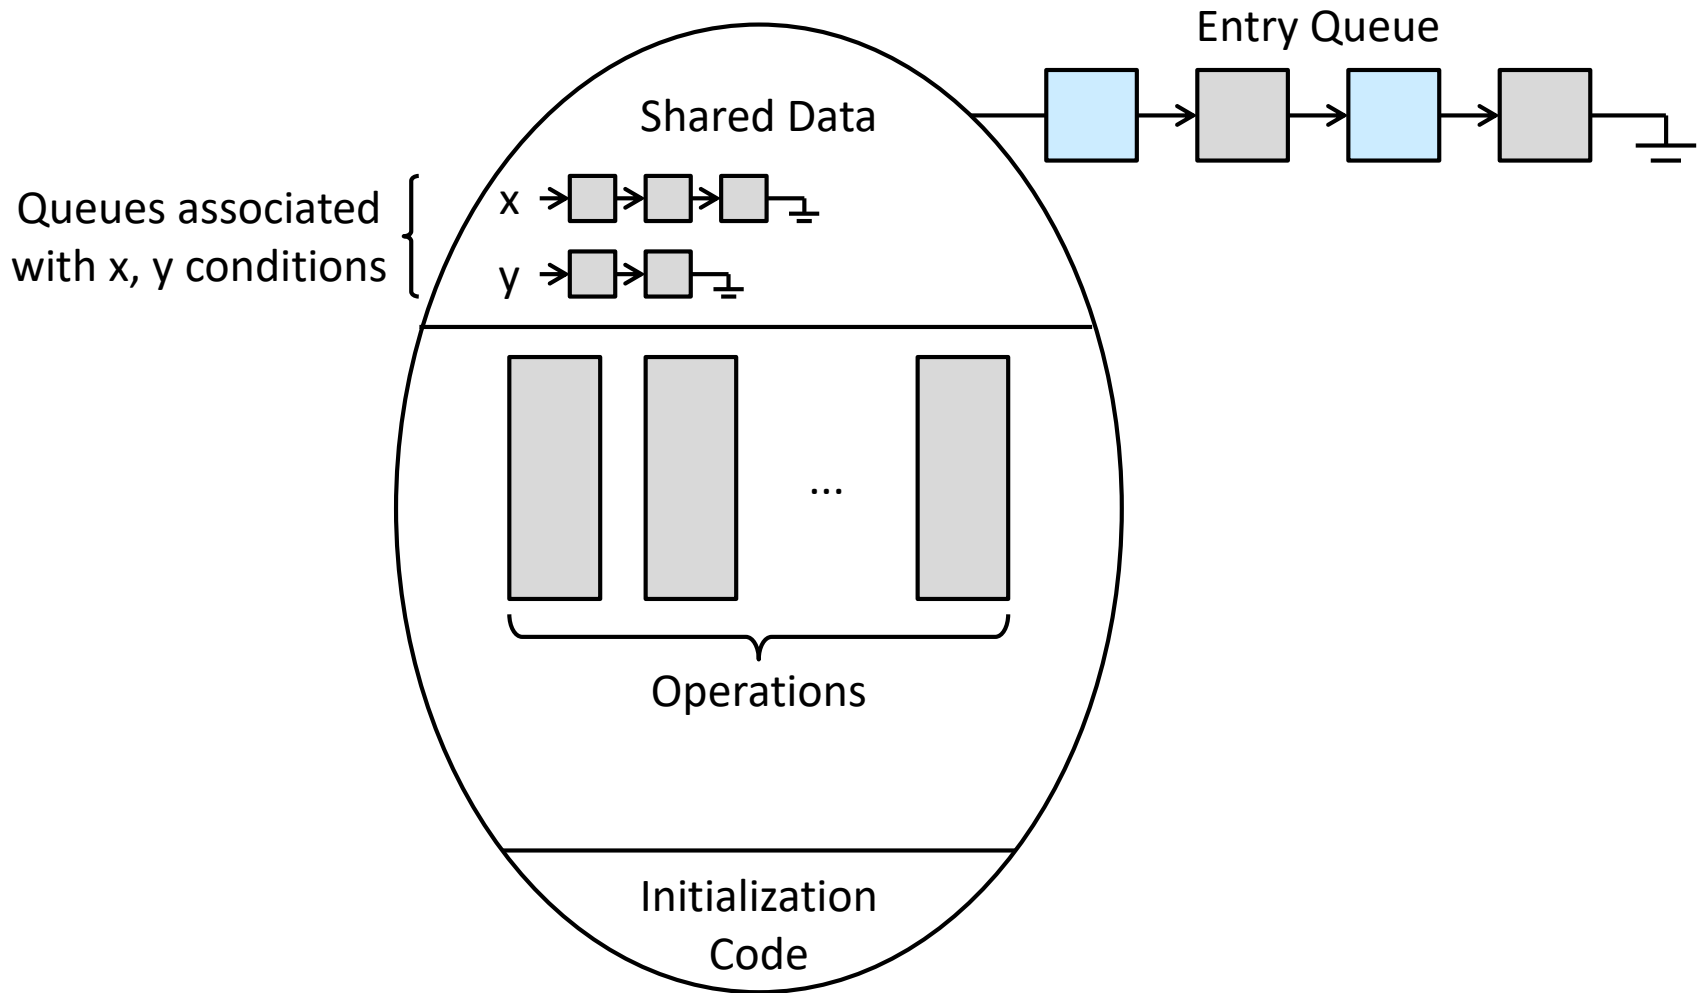
- The process invoking this operation is suspended until another process invokes `x.signal()`

➤ `x.signal()`

- Resume exactly one suspended process (if any)
- If no process is suspended, then it has no effect

Condition Variables (2/4)

Monitor with condition variables



Condition Variables (3/4)

❑ Suppose that when the **`x.signal()`** operation is invoked by a process P, there exists a suspended process Q associated with condition **`x`**

- Conceptually both processes can continue with their execution
- However, both P and Q should not be active simultaneously within the monitor

❑ Two possibilities

- Signal and wait
 - P either waits until Q leaves the monitor or waits for another condition
- Signal and continue
 - Q either waits until P leaves the monitor or waits for another condition

Condition Variables (4/4)

❑ Reasonable arguments in favor of adopting either option

- Since P was already executing in the monitor, the signal-and-continue method seems more reasonable
- If we allow P to continue, then by the time Q is resumed, the logical condition for which Q was waiting may no longer hold

❑ A compromise between the two options

- When thread P executes the **signal** operation, it immediately leaves the monitor
- Hence, Q is immediately resumed

Outline

- ❑ Background
- ❑ The Critical-Section Problem
- ❑ Peterson's Solution
- ❑ Hardware Support for Synchronization
- ❑ Mutex Locks
- ❑ Semaphores
- ❑ **Monitors**
 - Monitor Usage
 - **Implementing a Monitor Using Semaphores**
 - Resuming Processes within a Monitor
- ❑ Liveness
- ❑ Evaluation

Monitor Using Semaphores (1/3)

- ❑ A binary semaphore **mutex** is provided to ensure mutual exclusion

```
semaphore mutex; // (initially = 1)
```

- ❑ A process must
 - Execute **wait(mutex)** before entering the monitor
 - Execute **signal(mutex)** after leaving the monitor

Monitor Using Semaphores (2/3)

❑ Use the signal-and-wait scheme

- An additional binary semaphore **next**
 - The signaling processes can use **next** to suspend themselves
- An integer variable **next_count** counts the number of processes suspended on **next**

```
semaphore mutex; // (initially = 1)
semaphore next;  // (initially = 0)
int next_count = 0;
```

❑ Each external function F is replaced by

```
wait(mutex);
... body of F ...
if (next_count > 0)    signal(next);
else                  signal(mutex);
```

Monitor Using Semaphores (3/3)

- For each condition variable **x**, we have

```
semaphore x_sem; // (initially = 0)
int x_count = 0;
```

- **x.wait()** operation

```
x_count++;
if (next_count > 0)
    signal(next);
else
    signal(mutex);
wait(x_sem);
x_count--;
```

- **x.signal()** operation

```
if (x_count > 0) {
    next_count++;
    signal(x_sem);
    wait(next);
    next_count--;
}
```


Outline

- ❑ Background
- ❑ The Critical-Section Problem
- ❑ Peterson's Solution
- ❑ Hardware Support for Synchronization
- ❑ Mutex Locks
- ❑ Semaphores
- ❑ **Monitors**
 - Monitor Usage
 - Implementing a Monitor Using Semaphores
 - **Resuming Processes within a Monitor**
- ❑ Liveness
- ❑ Evaluation

Resuming Processes within a Monitor

❑ Which process should be resumed?

- If several processes are suspended on condition **x**, and
- An **x.signal()** operation is executed by some process

❑ One simple solution: first-come, first-served (FCFS) ordering

- Not adequate in many circumstances

❑ The conditional-wait construct **x.wait(c)** can be used

- **c** is an integer expression (called a priority number) evaluated when the **wait()** operation is executed
- The value of **c** is then stored with the name of the process that is suspended
- When **x.signal()** is executed, the process with the smallest priority number is resumed next

Single Resource Allocation (1/3)

- ❑ Each process, when requesting an allocation of a resource, specifies the maximum time t it plans to use the resource
- ❑ The monitor allocates the resource to the process that has the shortest time-allocation
- ❑ R is an instance of type **ResourceAllocator**

```
R.acquire(t) ;  
...  
access the resource ;  
...  
R.release ;
```

Single Resource Allocation (2/3)

```
monitor ResourceAllocator {
    boolean busy;
    condition x;
    void acquire(int time) {
        if (busy)
            x.wait(time);
        busy = true;
    }
    void release() {
        busy = false;
        x.signal();
    }
    initialization code() {
        busy = false;
    }
}
```

Single Resource Allocation (3/3)

❑ Problems

- A process might access a resource without first gaining access permission to the resource
- A process might never release a resource once it has been granted access to the resource
- A process might attempt to release a resource that it never requested
- A process might request the same resource twice (without first releasing the resource)

❑ Incorrect use of monitor operations

- **release()** ... **acquire()**
- **acquire()** ... **acquire()**
- Omit **acquire()** and/or **release()**

Outline

- ❑ Background
- ❑ The Critical-Section Problem
- ❑ Peterson's Solution
- ❑ Hardware Support for Synchronization
- ❑ Mutex Locks
- ❑ Semaphores
- ❑ Monitors
- ❑ **Liveness**
 - Deadlock
 - Priority Inversion
- ❑ Evaluation

Liveness

- ❑ Processes may have to wait indefinitely while trying to acquire a synchronization tool such as a mutex lock or semaphore
 - Waiting indefinitely violates the progress and bounded-waiting requirements
- ❑ **Liveness** refers to a set of properties that a system must satisfy to ensure processes make progress
 - Indefinite waiting is an example of a liveness failure
- ❑ **Starvation** (indefinite blocking)
 - A process may never be removed from the semaphore queue in which it is suspended

Outline

- ❑ Background
- ❑ The Critical-Section Problem
- ❑ Peterson's Solution
- ❑ Hardware Support for Synchronization
- ❑ Mutex Locks
- ❑ Semaphores
- ❑ Monitors
- ❑ **Liveness**
 - **Deadlock**
 - Priority Inversion
- ❑ Evaluation

Deadlock

❑ Every process in the set is waiting for an event that can be caused only by another process in the set

❑ Example

- Let S and Q be two semaphores initialized to 1
- P_0 : **wait**(S) ; **wait**(Q) ; ... **signal**(S) ; **signal**(Q) ;
- P_1 : **wait**(Q) ; **wait**(S) ; ... **signal**(Q) ; **signal**(S) ;
- P_0 executes **wait**(S) and P_1 executes **wait**(Q)
- When P_0 executes **wait**(Q), it must wait until P_1 executes **signal**(Q)
- However, P_1 is waiting until P_0 executes **signal**(S)
- Since these **signal**() operations will never be executed, P_0 and P_1 are deadlocked

❑ Chapter 8

Outline

- ❑ Background
- ❑ The Critical-Section Problem
- ❑ Peterson's Solution
- ❑ Hardware Support for Synchronization
- ❑ Mutex Locks
- ❑ Semaphores
- ❑ Monitors
- ❑ **Liveness**
 - Deadlock
 - **Priority Inversion**
- ❑ Evaluation

Priority Inversion

- ❑ Scheduling problem when lower-priority process holds a lock needed by higher-priority process

- Example

- There are three processes L, M, and H with priorities $L < M < H$
- H requires a semaphore S which is currently being accessed by L
- H waits for L to finish using resource S
- M becomes runnable and preempts L

- ❑ **Priority-inheritance protocol**

- All processes accessing resources needed by a higher-priority process inherit the higher priority until they are finished with the resources

Outline

- ❑ Background
- ❑ The Critical-Section Problem
- ❑ Peterson's Solution
- ❑ Hardware Support for Synchronization
- ❑ Mutex Locks
- ❑ Semaphores
- ❑ Monitors
- ❑ Liveness
- ❑ **Evaluation**

Evaluation

- ❑ Trying to identify when to use which tool can be a daunting challenge
 - Low overhead
 - Ability to scale
 - Ease to develop and test
 - Simplicity and ease of use
 - Optimistic approach vs. pessimistic approach
 - Uncontended
 - Moderate contention
 - High contention

Objectives

- ❑ Describe the critical-section problem and illustrate a race condition
- ❑ Illustrate hardware solutions to the critical-section problem using memory barriers, compare-and-swap operations, and atomic variables
- ❑ Demonstrate how mutex locks, semaphores, monitors, and condition variables can be used to solve the critical section problem
- ❑ Evaluate tools that solve the critical-section problem in low-, moderate-, and high-contention scenarios

Q&A