

Operating Systems

[7. Synchronization Examples]

Chung-Wei Lin

cwlin@csie.ntu.edu.tw

CSIE Department

National Taiwan University

Objectives

- ❑ Explain the bounded-buffer, readers-writers, and dining-philosophers synchronization problems
- ❑ Describe specific tools used by Linux and Windows to solve process synchronization problems
- ❑ Illustrate how POSIX and Java can be used to solve process synchronization problems

Outline

☐ Classic Problems of Synchronization

- Bounded-Buffer Problem
- Readers-Writers Problem
- Dining-Philosophers Problem

☐ Synchronization within the Kernel

☐ POSIX Synchronization

☐ Synchronization in Java

☐ Alternative Approaches

Bounded-Buffer Problem

□ n buffers

- Each can hold one item

□ Shared data

- Semaphore **mutex** initialized to the value 1 (binary)
- Semaphore **full** initialized to the value 0
- Semaphore **empty** initialized to the value n

Bounded-Buffer Problem: Producer

- The structure of the producer process

```
while (true) {  
    ...  
    /* produce an item in next_produced */  
    ...  
    wait(empty);  
    wait(mutex);  
    ...  
    /* add next_produced to the buffer */  
    ...  
    signal(mutex);  
    signal(full);  
}
```

Bounded-Buffer Problem: Consumer

- The structure of the consumer process

```
while (true) {  
    wait(full);  
    wait(mutex);  
  
    ...  
    /* remove an item from buffer to next_consumed */  
    ...  
    signal(mutex);  
    signal(empty);  
  
    ...  
    /* consume the item in next_consumed */  
    ...  
}
```

Outline

☐ Classic Problems of Synchronization

- Bounded-Buffer Problem
- Readers-Writers Problem
- Dining-Philosophers Problem

☐ Synchronization within the Kernel

☐ POSIX Synchronization

☐ Synchronization in Java

☐ Alternative Approaches

Readers-Writers Problem (1/2)

❑ A data set is shared among a number of concurrent processes

➤ Readers

- They only read the data set
- They do not perform any update

➤ Writers

- They can both read and write

❑ Problem

- Multiple readers can read at the same time
- Only one single writer can access the shared data at the same time

Readers-Writers Problem (2/2)

- ❑ Several variations of how readers and writers are considered
 - All involve some form of priorities
 - First readers-writers problem
 - No reader be kept waiting unless a writer has already obtained permission to use the shared object
 - Second readers-writers problem
 - Once a writer is ready, that writer perform its write as soon as possible
- ❑ Shared data (for the first readers-writers problem)
 - Data set
 - Semaphore **rw_mutex** initialized to 1 (binary)
 - Semaphore **mutex** initialized to 1 (binary)
 - Integer **read_count** initialized to 0

Readers-Writers Problem: Writer

- The structure of a writer process

```
while (true) {  
    wait(rw_mutex);  
    ...  
    /* writing is performed */  
    ...  
    signal(rw_mutex);  
}
```

Readers-Writers Problem: Reader

□ The structure of a reader process

```
while (true){
    wait(mutex);
    read_count++;
    if (read_count == 1) /* first reader */
        wait(rw_mutex);
    signal(mutex);
    ...
    /* reading is performed */
    ...
    wait(mutex);
    read_count--;
    if (read_count == 0) /* last reader */
        signal(rw_mutex);
    signal(mutex);
}
```

Reader-Writer Locks

- ❑ The readers-writers problem and its solutions have been generalized to provide reader-writer locks on some systems
 - Acquiring a reader-writer lock requires specifying the mode of the lock
 - A process wishing only to read shared data requests the lock in "read" mode
 - A process wishing to modify shared data requests the lock in "write" mode
 - Multiple processes are permitted to concurrently acquire the lock in read mode
 - Only one process may acquire the lock in write mode
- ❑ Reader-writer locks are most useful
 - In applications where it is easy to identify which processes only read shared data and which processes only write shared data
 - In applications that have more readers than writers

Outline

☐ Classic Problems of Synchronization

- Bounded-Buffer Problem
- Readers-Writers Problem
- Dining-Philosophers Problem

☐ Synchronization within the Kernel


☐ POSIX Synchronization

☐ Synchronization in Java

☐ Alternative Approaches

Dining-Philosophers Problem

- ❑ N philosophers sit at a round table with a bowl of rice in the middle
- ❑ They spend their lives alternating thinking and eating
- ❑ They occasionally try to pick up 2 chopsticks (one at a time) to eat from bowl
 - Need both to eat, then release both when done
- ❑ Shared data (in the case of 5 philosophers)
 - Bowl of rice (data set)
 - Semaphore **chopstick**[5] initialized to 1 (binary)



Semaphore Solution

- ❑ The structure of philosopher i

```
while (true){  
    wait (chopstick[i]);  
    wait (chopstick[(i+1)%5]);  
    ... /* eat for awhile */ ...  
    signal (chopstick[i]);  
    signal (chopstick[(i+1)%5]);  
    ... /* think for awhile */ ...  
}
```

- ❑ What is the problem with this algorithm?

Monitor Solution (1/3)

```
monitor DiningPhilosophers
{
    enum {THINKING, HUNGRY, EATING} state[5];
    condition self[5];
    void pickup(int i){
        state[i] = HUNGRY;
        test(i);
        if(state[i] != EATING) self[i].wait();
    }
    void putdown(int i){
        state[i] = THINKING;
        // test left and right neighbors
        test((i + 4) % 5);
        test((i + 1) % 5);
    }
}
```


Monitor Solution (2/3)

```
void test(int i){
    if((state[(i + 4) % 5] != EATING) &&
        (state[i] == HUNGRY) &&
        (state[(i + 1) % 5] != EATING)){
        state[i] = EATING;
        self[i].signal();
    }
}

initialization_code(){
    for(int i = 0; i < 5; i++)
        state[i] = THINKING;
}
}
```

Monitor Solution (3/3)

- ❑ Each philosopher i invokes the operations **pickup()** and **putdown()** in the following sequence:

```
...  
DiningPhilosophers.pickup(i);  
/* eat */  
DiningPhilosophers.putdown(i);  
...
```

- ❑ No deadlock, but starvation is possible

Outline

- ❑ Classic Problems of Synchronization
- ❑ **Synchronization within the Kernel**
 - **Synchronization in Windows**
 - Synchronization in Linux
- ❑ POSIX Synchronization
- ❑ Synchronization in Java
- ❑ Alternative Approaches

Synchronization in Windows (1/3)

❑ When the kernel accesses a global resource on a single-processor system

- It temporarily mask interrupts for all interrupt handlers that may also access the global resource

❑ On a multiprocessor system

- It protects access to global resources using spinlocks
 - The kernel uses spinlocks only to protect short code segments
 - For reasons of efficiency, the kernel ensures that a thread will never be preempted while holding a spinlock

Synchronization in Windows (2/3)


❑ For thread synchronization outside the kernel, Windows provides **dispatcher objects**

- Using a dispatcher object, threads synchronize according to several different mechanisms
 - Mutex locks, semaphores, events, and timers
- The system protects shared data by requiring a thread to
 - Gain ownership of a mutex to access the data
 - Release ownership when it is finished
- Events are similar to condition variables
 - They may notify a waiting thread when a desired condition occurs
- Timers are used to notify one (or more than one) thread that a specified amount of time has expired

Synchronization in Windows (3/3)

❑ Dispatcher objects may be in either a signaled or nonsignaled state

- An object in a signaled state is available
 - A thread will not block when acquiring the object
- An object in a nonsignaled state is not available
 - A thread will block when attempting to acquire the object



Outline

- ❑ Classic Problems of Synchronization
- ❑ **Synchronization within the Kernel**
 - Synchronization in Windows
 - **Synchronization in Linux**
- ❑ POSIX Synchronization
- ❑ Synchronization in Java
- ❑ Alternative Approaches

Synchronization in Linux

- ❑ A nonpreemptive kernel prior to Version 2.6; preemptive now

- ❑ Atomic integers

```
atomic_t counter; (atomic_t is the type for atomic integer)
int value;
atomic_set(&counter, 5); // counter = 5
atomic_add(10, &counter); // counter = counter + 10
atomic_sub(4, &counter); // counter = counter - 4
atomic_inc(&counter); // counter = counter + 1
value = atomic_read(&counter); // value = 12
```

- ❑ Single processor

- Disable/enable kernel preemption

- ❑ Multiple processors

- Acquire/release spinlock

Outline

- ❑ Bounded-Buffer Problem
- ❑ Readers-Writers Problem
- ❑ Dining-Philosophers Problem
- ❑ Synchronization in Windows and Linux
- ❑ **POSIX Synchronization**
 - **POSIX Mutex Locks**
 - POSIX Semaphores
 - POSIX Condition Variables
- ❑ Synchronization in Java
- ❑ Alternative Approaches

POSIX Mutex Locks

❑ Create and initialize the lock

```
#include <pthread.h>
pthread_mutex_t mutex;
/* create and initialize the mutex lock */
pthread_mutex_init(&mutex, NULL);
```

❑ Acquire and release the lock

```
/* acquire the mutex lock */
pthread_mutex_lock(&mutex);
/* critical section */
/* release the mutex lock */
pthread_mutex_unlock(&mutex);
```

Outline

- ❑ Bounded-Buffer Problem
- ❑ Readers-Writers Problem
- ❑ Dining-Philosophers Problem
- ❑ Synchronization in Windows and Linux
- ❑ **POSIX Synchronization**
 - POSIX Mutex Locks
 - **POSIX Semaphores**
 - POSIX Condition Variables
- ❑ Synchronization in Java
- ❑ Alternative Approaches

POSIX Semaphores

- ❑ Semaphores are not part of the POSIX standard and instead belong to the POSIX SEM extension
- ❑ **Named** semaphores
 - Can be used by unrelated processes
- ❑ **Unnamed** semaphores
 - Cannot be used by unrelated processes

POSIX Named Semaphores

- ❑ Create and initialize the semaphore

```
#include <semaphore.h>
sem_t *sem;
/* Create the semaphore & initialize it to 1 */
sem = sem_open("SEM", O_CREAT, 0666, 1);
```

- ❑ Another process can access the semaphore by referring to its name **SEM**

- ❑ Acquire and release the semaphore

```
/* acquire the semaphore */
sem_wait(sem);
/* critical section */
/* release the semaphore */
sem_post(sem);
```

POSIX Unnamed Semaphores

❑ Create and initialize the semaphore

```
#include <semaphore.h>
sem_t sem;
/* Create the semaphore & initialize it to 1 */
sem_init(&sem, 0, 1);
```

❑ Acquire and release the semaphore

```
/* acquire the semaphore */
sem_wait(&sem);
/* critical section */
/* release the semaphore */
sem_post(&sem);
```

Outline

- ❑ Bounded-Buffer Problem
- ❑ Readers-Writers Problem
- ❑ Dining-Philosophers Problem
- ❑ Synchronization in Windows and Linux
- ❑ **POSIX Synchronization**
 - POSIX Mutex Locks
 - POSIX Semaphores
 - **POSIX Condition Variables**
- ❑ Synchronization in Java
- ❑ Alternative Approaches

POSIX Condition Variables (1/2)

- ❑ POSIX condition variables are associated with a POSIX mutex lock to provide mutual exclusion
 - Since POSIX is typically used in C/C++ and these languages do not provide a monitor

- ❑ Create and initialize the condition variable

```
pthread_mutex_t mutex;  
pthread_cond_t cond_var;  
pthread_mutex_init(&mutex, NULL) ;  
pthread_cond_init(&cond_var, NULL) ;
```


POSIX Condition Variables (2/2)

- ❑ Thread waiting for the condition `a == b` to become true

```
pthread_mutex_lock(&mutex);  
while(a != b)  
    pthread_cond_wait(&cond_var, &mutex)  
pthread_mutex_unlock(&mutex);
```

- ❑ Thread signaling another thread waiting on the condition variable

```
pthread_mutex_lock(&mutex);  
a = b;  
pthread_cond_signal(&cond_var);  
pthread_mutex_unlock(&mutex);
```

Outline


- ❑ Bounded-Buffer Problem
- ❑ Readers-Writers Problem
- ❑ Dining-Philosophers Problem
- ❑ Synchronization in Windows and Linux
- ❑ POSIX Synchronization
- ❑ **Synchronization in Java**
 - **Java Monitors**
 - Reentrant Locks
 - Semaphores
 - Condition Variables
- ❑ Alternative Approaches

Java Monitors (1/3)

- ❑ Every object in Java has associated with it a single lock
- ❑ When a method is declared to be synchronized, calling the method requires owning the lock for the object
 - Declare a synchronized method by placing the `synchronized` keyword in the method definition
 - Examples: the `insert()` and `remove()` methods in the `BoundedBuffer` class
- ❑ If the lock is already owned by another thread, the calling thread blocks and is placed in the entry set of the object
 - The entry set represents the set of threads waiting for the lock to become available
- ❑ The lock is released when the thread exits the method

Java Monitors (2/3)

- ❑ A thread that tries to acquire an unavailable lock is placed in the object's entry set
- ❑ Each object also has a wait set
 - When a thread calls `wait()`
 - The thread releases the lock for the object
 - The state of the thread is set to blocked
 - The thread is placed in the wait set for the object



Java Monitors (3/3)

- ❑ A thread typically calls **wait()** when it is waiting for a condition to become true
 - How does the thread get notified?
- ❑ When a thread calls **notify()**
 - Pick an arbitrary thread T from the list of threads in the wait set
 - Move T from the wait set to the entry set
 - Set the state of T from blocked to runnable
- ❑ T is now eligible to compete for the lock with the other threads

Java Monitors: Bounded-Buffer (1/3)

```
public class BoundedBuffer<E> {  
    private static final int BUFFER_SIZE = 5;  
    private int count, in, out;  
    private E[] buffer;  
    public BoundedBuffer() {  
        count = 0;  
        in = 0;  
        out = 0;  
        buffer = (E[]) new Object[BUFFER_SIZE];  
    }  
    /* Producers call this method */  
    public synchronized void insert(E item) {  
        /* See following slides */  
    }  
    /* Consumers call this method */  
    public synchronized E remove() {  
        /* See following slides */  
    }  
}
```

Java Monitors: Bounded-Buffer (2/3)

```
/* Producers call this method */
public synchronized void insert(E item) {
    while (count == BUFFER_SIZE) {
        try {
            wait();
        }
        catch (InterruptedException ie) { }
    }
    buffer[in] = item;
    in = (in + 1) % BUFFER_SIZE;
    count++;
    notify();
}
```

Java Monitors: Bounded-Buffer (3/3)

```
/* Consumers call this method */
public synchronized E remove() {
    E item;
    while (count == 0) {
        try {
            wait();
        }
        catch (InterruptedException ie) { }
    }
    item = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    count--;
    notify();
    return item;
}
```


Outline

- ❑ Bounded-Buffer Problem
- ❑ Readers-Writers Problem
- ❑ Dining-Philosophers Problem
- ❑ Synchronization in Windows and Linux
- ❑ POSIX Synchronization
- ❑ **Synchronization in Java**
 - Java Monitors
 - **Reentrant Locks**
 - Semaphores
 - Condition Variables
- ❑ Alternative Approaches

Java Reentrant Locks

- ❑ **lock()** assigns the invoking thread lock ownership
 - If the lock is available, or
 - If the thread already owns it, which is why it is termed reentrant
- ❑ The **finally** clause ensures that the lock will be released in case an exception occurs in the **try** block

```
Lock key = new ReentrantLock();  
key.lock();  
try {  
    / * critical section */  
}  
finally {  
    key.unlock();  
}
```

Outline

- ❑ Bounded-Buffer Problem
- ❑ Readers-Writers Problem
- ❑ Dining-Philosophers Problem
- ❑ Synchronization in Windows and Linux
- ❑ POSIX Synchronization
- ❑ **Synchronization in Java**
 - Java Monitors
 - Reentrant Locks
 - **Semaphores**
 - Condition Variables
- ❑ Alternative Approaches

Java Semaphores

❑ Constructor

```
Semaphore(int value);
```

❑ Usage

```
Semaphore sem = new Semaphore(1);  
try {  
    sem.acquire();  
    /* critical section */  
}  
catch (InterruptedException ie) {}  
finally {  
    sem.release();  
}
```

Outline

- ❑ Bounded-Buffer Problem
- ❑ Readers-Writers Problem
- ❑ Dining-Philosophers Problem
- ❑ Synchronization in Windows and Linux
- ❑ POSIX Synchronization
- ❑ **Synchronization in Java**
 - Java Monitors
 - Reentrant Locks
 - Semaphores
 - **Condition Variables**
- ❑ Alternative Approaches

Java Condition Variables (1/3)

❑ Create a condition variable by

- Creating a **ReentrantLock**
- Invoking its **newCondition()** method
 - It returns a **Condition** object representing the condition variable for the associated **ReentrantLock**

❑ Usage

```
Lock key = ReentrantLock();  
Condition condVar = key.newCondition();
```

❑ A thread

- Wait by calling the **await()** method
- Signal by calling the **signal()** method

Java Condition Variables (2/3)

❑ Example

- Five threads numbered from 0 to 4
 - Threads share variable `turn` indicating which thread's turn it is
- A thread calls `doWork()` when it wishes to do some work
 - If it is not its turn, wait
 - If it is its turn, do some work for awhile
 - When it completes, it notifies the thread whose turn is next

❑ Necessary data structures

```
Lock lock = new ReentrantLock();  
Condition[] condVars = new Condition[5];  
for (int i = 0; i < 5; i++)  
    condVars[i] = lock.newCondition();
```

Java Condition Variables (3/3)

```
/* threadNumber is the thread wishing to do some work */
public void doWork(int threadNumber) {
    lock.lock();
    try {
        /* If not my turn, then wait until signaled */
        if (threadNumber != turn)
            condVars[threadNumber].await();
        /* Do some work for awhile */
        /* Now signal to the next thread */
        turn = (turn + 1) % 5;
        condVars[turn].signal();
    }
    catch (InterruptedException ie) { }
    finally {
        lock.unlock();
    }
}
```


Outline

- ❑ Bounded-Buffer Problem
- ❑ Readers-Writers Problem
- ❑ Dining-Philosophers Problem
- ❑ Synchronization in Windows and Linux
- ❑ POSIX Synchronization
- ❑ Synchronization in Java
- ❑ **Alternative Approaches**
 - Transactional Memory
 - OpenMP
 - Functional Programming Languages

Transactional Memory

- ❑ A memory transaction is a sequence of read-write operations to memory that are performed atomically
 - A transaction can be completed by adding `atomic{S}` which ensure statements in `S` are executed atomically

```
void update () {  
    atomic {  
        /* modify shared data */  
    }  
}
```

OpenMP

- ❑ OpenMP is a set of compiler directives and API that support parallel programming

```
void update(int value)
{
    #pragma omp critical
    {
        count += value
    }
}
```

- ❑ The code contained within the `#pragma omp critical` directive is
 - Treated as a critical section
 - Performed atomically

Functional Programming Languages

- ❑ Functional programming languages offer a different paradigm than procedural languages in that they do not maintain state
- ❑ Variables are treated as immutable and cannot change state once they have been assigned a value
- ❑ There is increasing interest in functional languages such as Erlang and Scala for their approach in handling data races

Objectives

- ❑ Explain the bounded-buffer, readers-writers, and dining-philosophers synchronization problems
- ❑ Describe specific tools used by Linux and Windows to solve process synchronization problems
- ❑ Illustrate how POSIX and Java can be used to solve process synchronization problems

Q&A