

Operating Systems

[9. Main Memory]

Chung-Wei Lin

cwlin@csie.ntu.edu.tw

CSIE Department

National Taiwan University

Objectives

- ❑ Explain the difference between a logical and a physical address and the role of the memory management unit (MMU) in translating addresses
- ❑ Apply first-, best-, and worst-fit strategies for allocating memory contiguously
- ❑ Explain the distinction between internal and external fragmentation
- ❑ Translate logical to physical addresses in a paging system that includes a translation look-aside buffer (TLB)
- ❑ Describe hierarchical paging, hashed paging, and inverted page tables

Outline

☐ **Background**

- Basic Hardware
- Address Binding
- Logical Versus Physical Address Space
- Dynamic Loading
- Dynamic Linking and Shared Libraries

☐ **Contiguous Memory Allocation**

☐ **Paging**

☐ **Structure of the Page Table**

☐ **Swapping**

☐ **Example: Intel 32- and 64-bit Architecture**

☐ **Example: ARMv8 Architecture**

Background

❑ A typical instruction-execution cycle

- Fetch an instruction from memory
- Decode the instruction
- Fetch operands from memory
- Execute the instruction
- Store back results in memory

❑ We are interested only in the sequence of memory addresses generated by the running program

- The memory unit sees only a stream of memory addresses
- It does not know how they are generated or what they are for (instructions or data)

Outline

☐ **Background**

- **Basic Hardware**
- Address Binding
- Logical Versus Physical Address Space
- Dynamic Loading
- Dynamic Linking and Shared Libraries

☐ **Contiguous Memory Allocation**

☐ **Paging**

☐ **Structure of the Page Table**

☐ **Swapping**

☐ **Example: Intel 32- and 64-bit Architecture**

☐ **Example: ARMv8 Architecture**


Basic Hardware

❑ The general-purpose storage that the CPU can access directly

- Main memory
- Registers built into each CPU core

❑ **Cache** is also added between the CPU and main memory

- Reason
 - Registers built into each CPU core are generally accessible within one cycle of the CPU clock
 - Completing a memory access may take many cycles of the CPU clock, and thus the CPU may need to **stall**



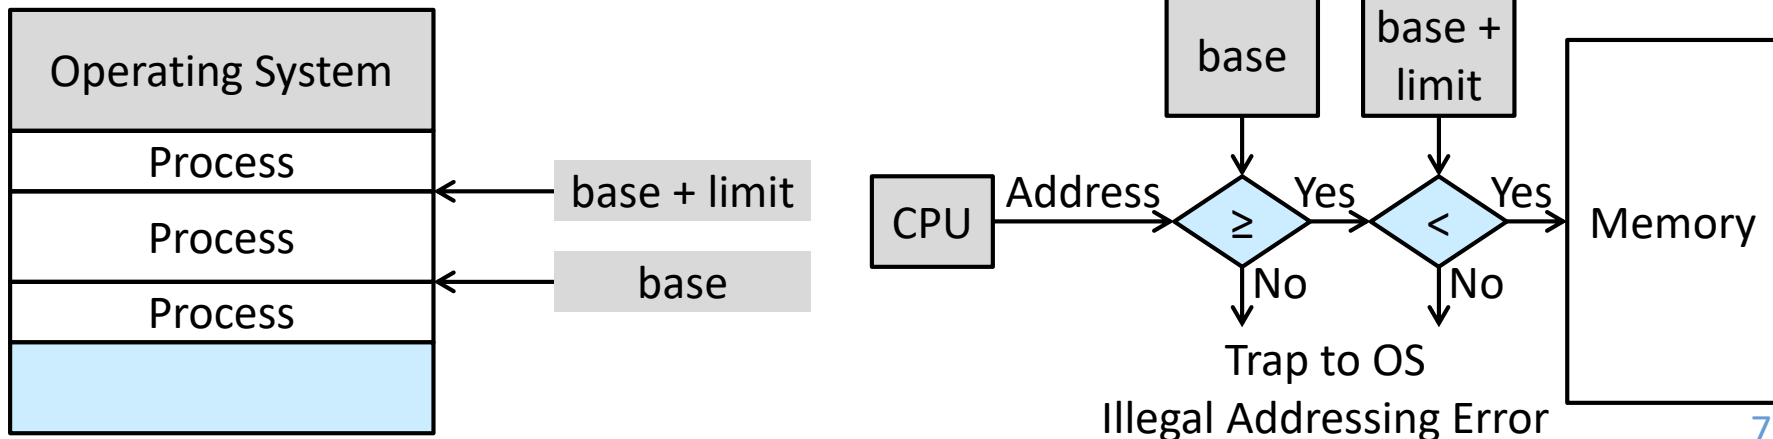
Address Protection

❑ This protection must be provided by the hardware

- An operating system does not usually intervene between the CPU and its memory accesses

❑ The protection compares every address generated in user mode with two registers

- The **base register** holds the smallest legal physical memory address
- The **limit register** specifies the size of the range
- The registers can be loaded only by the operating system which uses a special privileged instruction



Outline

☐ **Background**

- Basic Hardware
- **Address Binding**
- Logical Versus Physical Address Space
- Dynamic Loading
- Dynamic Linking and Shared Libraries

☐ Contiguous Memory Allocation

☐ Paging

☐ Structure of the Page Table

☐ Swapping

☐ Example: Intel 32- and 64-bit Architecture

☐ Example: ARMv8 Architecture

Address Binding (1/2)

❑ The steps before executing a user program

- Addresses in the source program are generally symbolic

- Example: a variable **count**

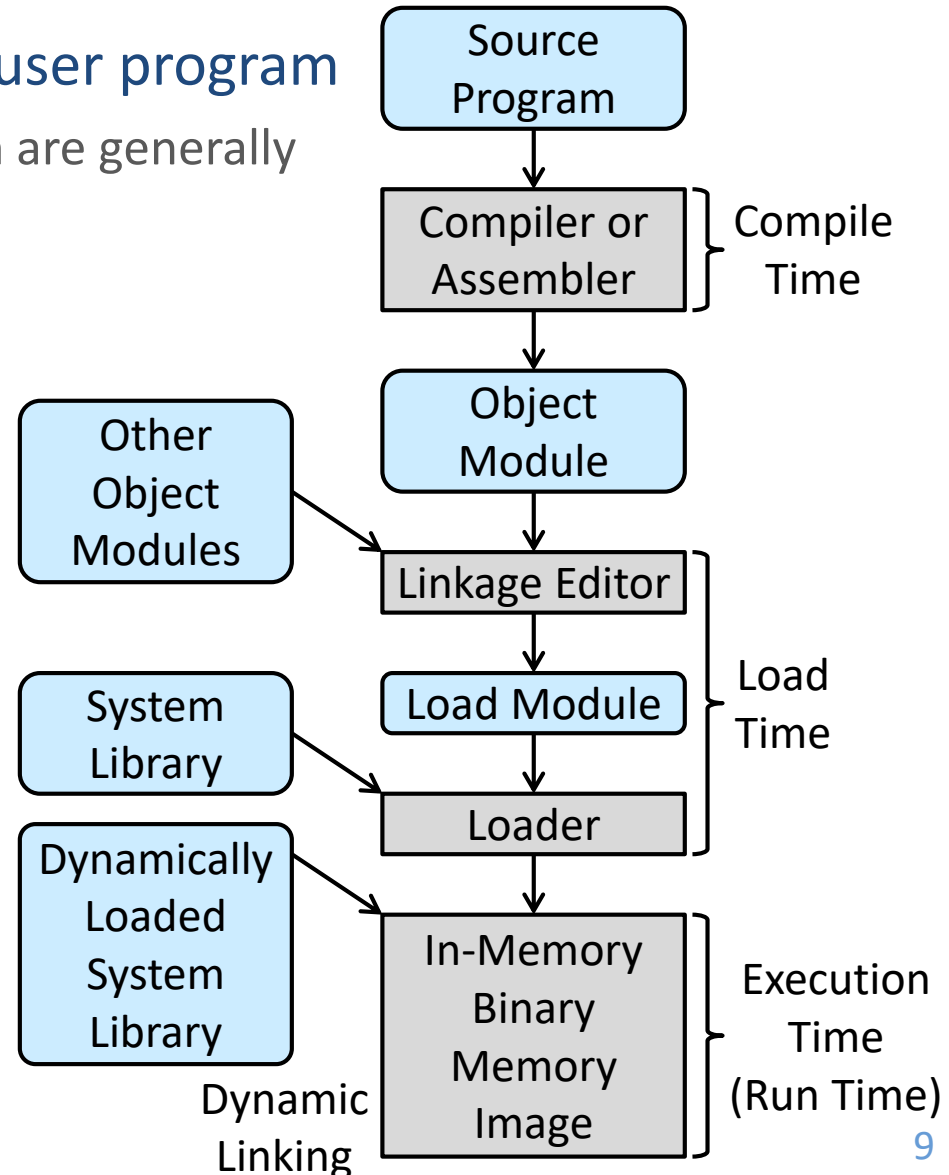
- A compiler typically **binds** the symbolic addresses to relocatable addresses

- Example: 14 bytes from the beginning of this module

- A linker or loader binds the relocatable addresses to absolute addresses

- Example: 74014

- Each binding is a mapping from one address space to another



Address Binding (2/2)

❑ The binding of instructions and data to memory addresses at

➤ Compile time

- If the location of a process in memory is known, then the compiler can generate absolute code
- If the starting location changes, then recompiling is needed

➤ Load time

- If the location of a process in memory is not known at compile time, then the compiler generates relocatable code
- The final binding is delayed until load time
- If the starting address changes, the reloading is needed

➤ Execution time

- If the process can be moved within memory during its execution, then the binding is delayed until run time
- Special hardware must be available (next slide)
- Most operating systems use this method

Outline

☐ **Background**

- Basic Hardware
- Address Binding
- **Logical Versus Physical Address Space**
- Dynamic Loading
- Dynamic Linking and Shared Libraries

☐ Contiguous Memory Allocation

☐ Paging

☐ Structure of the Page Table

☐ Swapping

☐ Example: Intel 32- and 64-bit Architecture

☐ Example: ARMv8 Architecture

Logical Versus Physical Address Space

❑ Logical address

- An address generated by the CPU

❑ Physical address

- An address seen by the memory unit, i.e., the one loaded into the memory-address register of the memory

❑ Binding addresses has

- Identical logical and physical addresses at either compile or load time
- Different logical and physical addresses at execution time
 - In this case, we usually refer to the logical address as a virtual address

❑ Logical address space


- Set of all logical addresses generated by a program

❑ Physical address space

- Set of all physical addresses corresponding to these logical addresses

Memory-Management Unit (MMU)

- ❑ Hardware device performing run-time mapping from virtual to physical addresses
 - The base register is now called a relocation register
 - The user program deals with logical addresses
 - MMU converts logical addresses into physical addresses



Outline

☐ **Background**

- Basic Hardware
- Address Binding
- Logical Versus Physical Address Space
- **Dynamic Loading**
- Dynamic Linking and Shared Libraries

☐ Contiguous Memory Allocation

☐ Paging

☐ Structure of the Page Table

☐ Swapping

☐ Example: Intel 32- and 64-bit Architecture

☐ Example: ARMv8 Architecture

Dynamic Loading

- ❑ A routine is not loaded until it is called for better memory-space utilization
 - All routines are kept on disk in a relocatable load format
 - The main program is first loaded into memory and executed
- ❑ This method is particularly useful when large amounts of code are needed to handle infrequently occurring cases
 - Examples: error routines
- ❑ Dynamic loading does not require special support from the operating system
 - It is the responsibility of the users to design their programs to take advantage of such a method
 - Operating systems may help the programmer by providing library routines to implement dynamic loading

Outline

☐ **Background**

- Basic Hardware
- Address Binding
- Logical Versus Physical Address Space
- Dynamic Loading
- **Dynamic Linking and Shared Libraries**

☐ Contiguous Memory Allocation

☐ Paging

☐ Structure of the Page Table

☐ Swapping

☐ Example: Intel 32- and 64-bit Architecture

☐ Example: ARMv8 Architecture

Dynamic Linking and Shared Libraries

❑ Static linking

- System libraries are treated like any other object module

❑ Dynamic linking

- Linking is postponed until execution time
 - Advantages: decrease the size of an executable image (save main memory), enable sharing among multiple processes, and enable updates

❑ **Dynamically linked libraries** (DLLs) are system libraries that are linked to user programs when programs are run

- Only one instance of the DLL is in main memory
 - For this reason, DLLs are also known as shared libraries
- DLLs are used extensively in Windows and Linux systems

❑ Dynamic linking and shared libraries generally require support from the operating system

Outline

- ❑ Background

- ❑ **Contiguous Memory Allocation**

- Memory Protection
- Memory Allocation
- Fragmentation

- ❑ Paging

- ❑ Structure of the Page Table

- ❑ Swapping

- ❑ Example: Intel 32- and 64-bit Architecture

- ❑ Example: ARMv8 Architecture

Contiguous Memory Allocation

❑ The memory is usually divided into two partitions

- One for the operating system
 - Many operating systems (including Linux and Windows) place the operating system in high memory
- One for the user processes

❑ Contiguous memory allocation (one early method)

- Each process is contained in a single section of memory that is contiguous to the section containing the next process

Outline

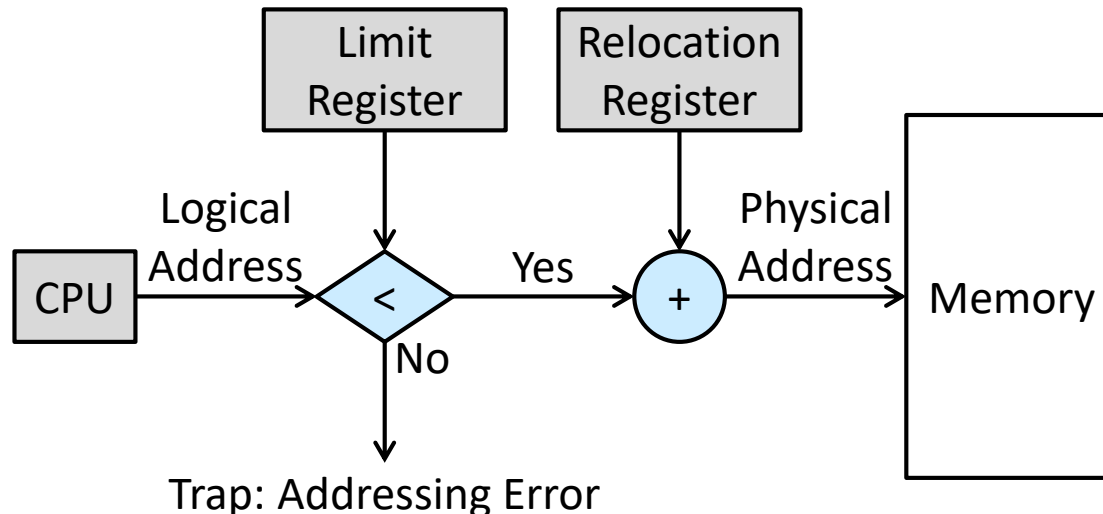
- ❑ Background
- ❑ **Contiguous Memory Allocation**
 - **Memory Protection**
 - Memory Allocation
 - Fragmentation
- ❑ Paging
- ❑ Structure of the Page Table
- ❑ Swapping
- ❑ Example: Intel 32- and 64-bit Architecture
- ❑ Example: ARMv8 Architecture

Memory Protection

❑ Combine two ideas previously discussed

- The limit register contains the range of logical addresses
- The relocation register contains the value of the smallest physical address

❑ The dispatcher loads the relocation and limit registers with the correct values as part of the context switch

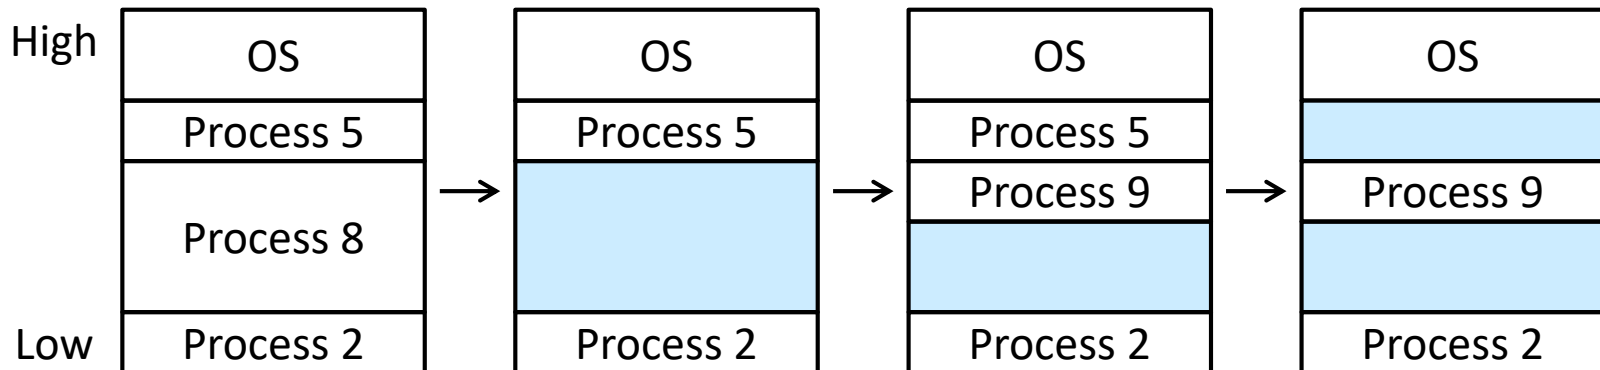


Outline

- ❑ Background
- ❑ **Contiguous Memory Allocation**
 - Memory Protection
 - **Memory Allocation**
 - Fragmentation
- ❑ Paging
- ❑ Structure of the Page Table
- ❑ Swapping
- ❑ Example: Intel 32- and 64-bit Architecture
- ❑ Example: ARMv8 Architecture

Memory Allocation

- ❑ Assign processes to variably sized partitions in memory, where each partition may contain exactly one process
 - A **variable-partition** scheme keeps a table indicating which parts of memory are available and which are occupied
 - Initially, all memory is available and considered one large block, a **hole**
 - When a process arrives, it is
 - Allocated space and loaded into memory
 - Rejected and provided an error message, or placed into a wait queue
 - When a process terminates, it releases its memory
 - Adjacent holes are merged



Dynamic Storage Allocation Problem

□ How to satisfy a request of size n from a list of free holes

- **First fit** allocates the first hole that is big enough
- **Best fit** allocates the smallest hole that is big enough
 - Search the entire list, unless the list is ordered by size
 - Produce the smallest leftover hole
- **Worst fit** allocates the largest hole
 - Search the entire list, unless the list is ordered by size
 - Produce the largest leftover hole
- Simulations have shown that both first fit and best fit are better than worst fit in terms of decreasing time and storage utilization

Outline

- ❑ Background

- ❑ **Contiguous Memory Allocation**

 - Memory Protection

 - Memory Allocation

 - **Fragmentation**

- ❑ Paging

- ❑ Structure of the Page Table

- ❑ Swapping

- ❑ Example: Intel 32- and 64-bit Architecture

- ❑ Example: ARMv8 Architecture

Internal Fragmentation

❑ Example

- A hole of 18,464 bytes
- A process requests 18,462 bytes
- If we allocate exactly the requested block, a hole of 2 bytes is left
 - The overhead to keep track of this hole will be substantially larger than the hole itself

❑ Internal fragmentation

- Break the physical memory into fixed-sized blocks and allocate memory in units based on block size, to avoid the problem above
 - The memory allocated to a process may be slightly larger than the requested memory
 - The difference between these two numbers is internal fragmentation (unused memory that is internal to a partition)

External Fragmentation

❑ External fragmentation

- There is enough total memory space to satisfy a request but the available spaces are not contiguous
 - Storage is fragmented into a large number of small holes
- 50-percent rule
 - Statistical analysis reveals that, even with some optimization, given N allocated blocks, another $0.5N$ blocks will be lost to fragmentation

❑ Solutions

- Compaction with dynamic relocation at execution time
 - Shuffle the memory contents to place all free memory together
- Paging (next section)
 - Permit the "physical" address space of processes to be noncontiguous
 - The most common memory-management technique for computer systems

Outline

- ❑ Background
- ❑ Contiguous Memory Allocation
- ❑ **Paging**
 - **Basic Method**
 - Hardware Support
 - Protection
 - Shared Pages
- ❑ Structure of the Page Table
- ❑ Swapping
- ❑ Example: Intel 32- and 64-bit Architecture
- ❑ Example: ARMv8 Architecture

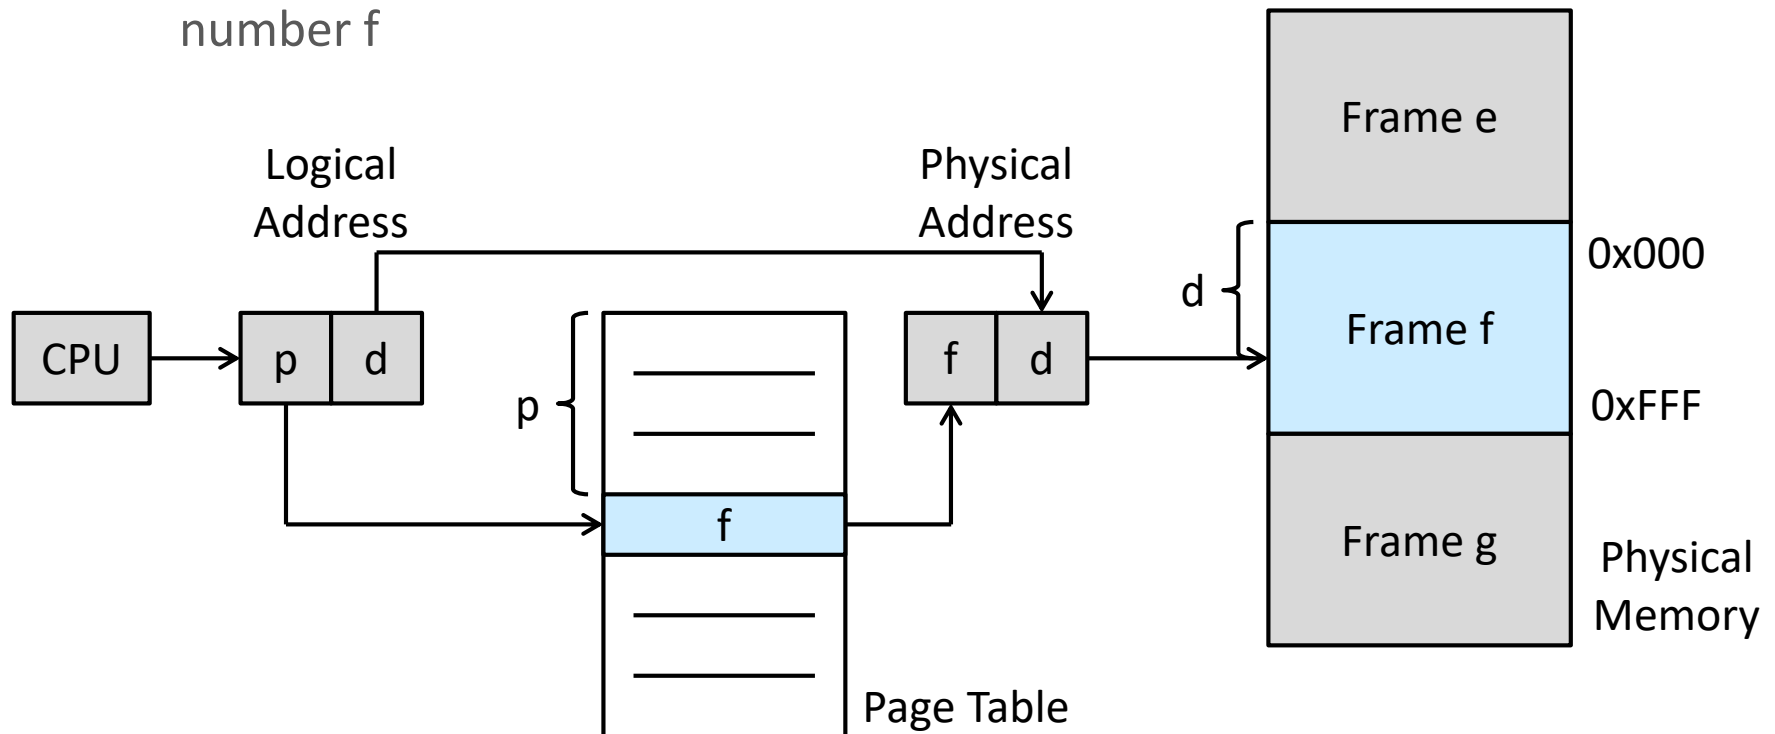
Pages and Frames

- ❑ When a process is to be executed, its pages are loaded into any available memory frames from their source
 - Pages: fixed-sized blocks of logical memory
 - Frames: fixed-sized blocks of physical memory
- ❑ Every address generated by the CPU is divided into two parts
 - A page number p : an index into a per-process page table
 - A page offset d : the location in the frame being referenced
- ❑ The page size (= the frame size) is defined by the hardware
 - The size is a power of 2, typically between 4 KB and 1 GB
 - If the size of the logical address space is 2^m , and a page size is 2^n bytes
 - The high-order $(m - n)$ bits of a logical address designate the page number
 - The low-order n bits of a logical address designate the page offset


Paging Hardware

❑ Steps taken by the MMU to translate a logical address generated by the CPU to a physical address

- Extract the page number p and use it as an index into the page table
- Extract the corresponding frame number f from the page table
- Replace the page number p in the logical address with the frame number f



Paging Model




Logical
Memory

Page Table

0	1
1	4
2	3
3	7

Frame
Number → 7



Physical
Memory

Paging Example

❑ $m = 4$

❑ $n = 2$

❑ Page size = 4 bytes

❑ # of frames = 8

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

Logical
Memory

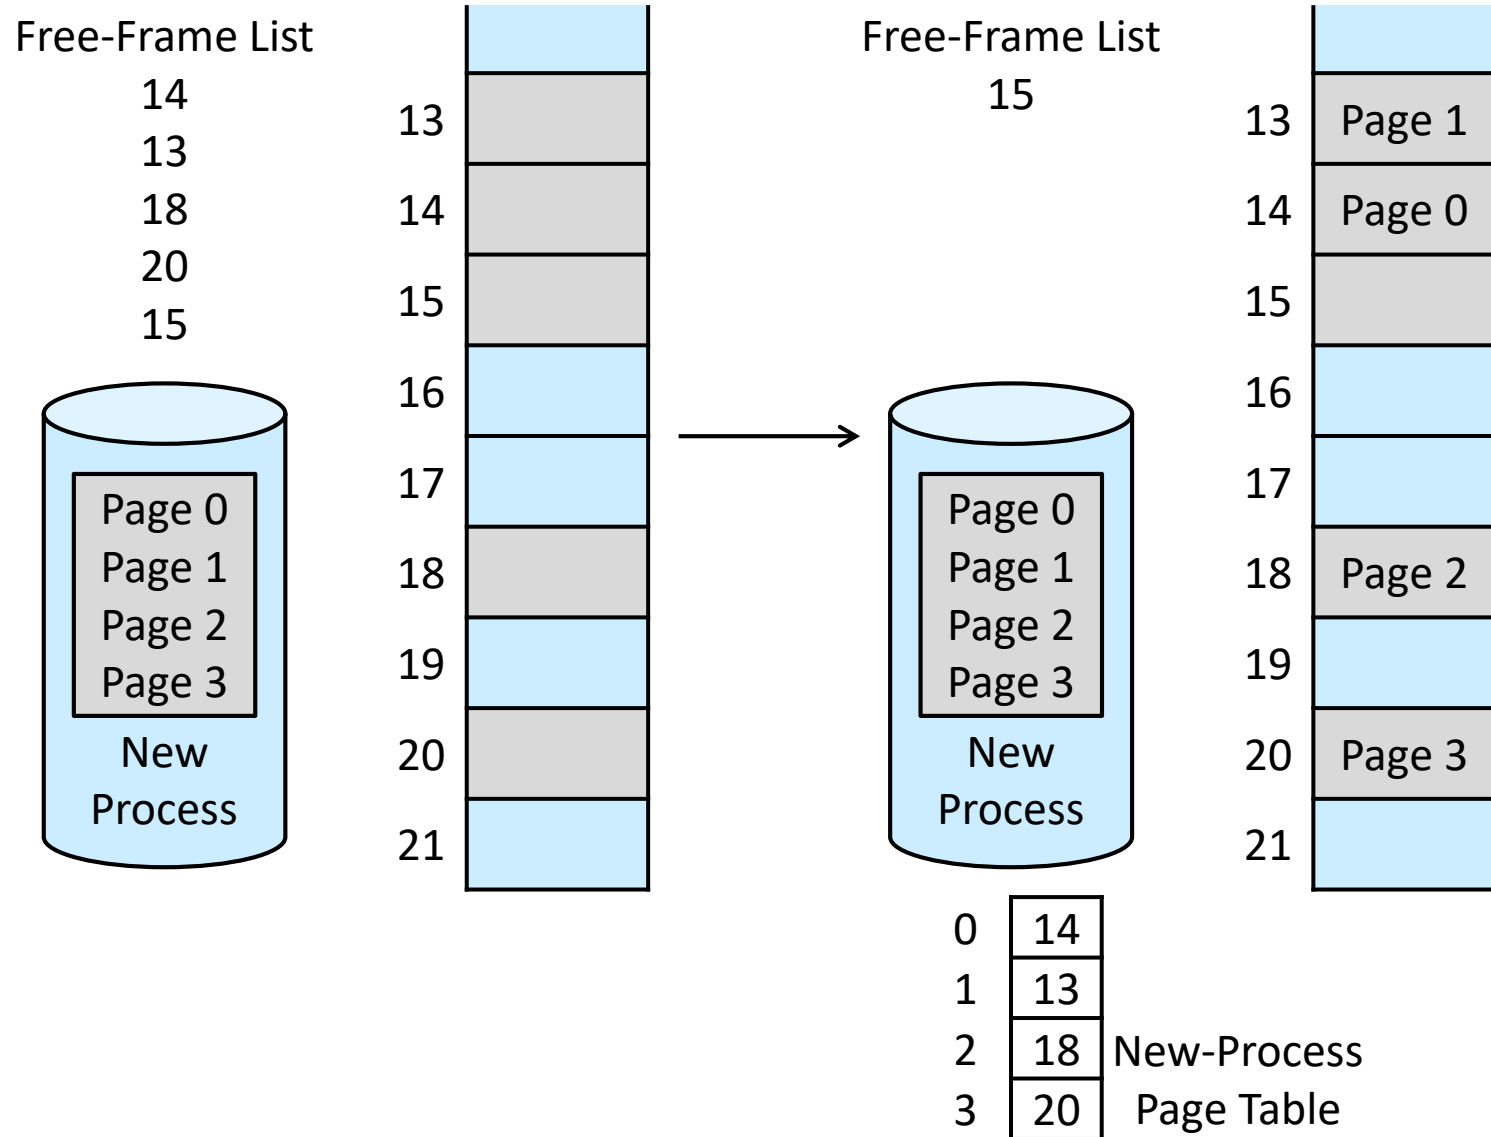
0	5
1	6
2	1
3	2

Page Table

0	
4	i j k l
8	m n o p
12	
16	
20	a b c d
24	e f g h
28	

Physical
Memory 32

Process Allocation



Paging's Internal Fragmentation

❑ Example

- Page size = 2,048 bytes
- Process size = 72,766 bytes = 35 pages + 1,086 bytes
- Internal fragmentation = 2,048 – 1,086 bytes = 962 bytes

❑ Worst-case fragmentation = 1 frame – 1 byte

❑ Average fragmentation = 0.5 frame

❑ Small page sizes are desirable?

- The overhead of page table entries decreases as the page size increases
- Disk I/O is more efficient when the amount of transferred data is larger

❑ Generally, page sizes have grown over time as processes, data sets, and main memory have become larger

- Today, pages are typically either 4 KB or 8 KB in size
- Some systems support even larger page sizes

Amount of Addressed Memory

- ❑ Frequently, on a 32-bit CPU, each page-table entry is 4-byte long, but that size can vary as well
 - A 4-byte entry can point to one of 2^{32} physical page frames
 - If the frame size is 4 KB (2^{12} bytes), then a system with 4-byte entries can address 16 TB (2^{44} bytes) of physical memory
- ❑ However, a system with 4-byte page-table entries may address less physical memory than the possible maximum
 - Other information that must be kept in the page-table entries
 - We will see in the later slides

Some Paging Facts

❑ An important aspect of paging

- The clear separation between the programmer's view of memory and the actual physical memory

❑ Frame table

- A single, system-wide data structure having one entry for each physical frame, indicating
 - Whether a frame is free or allocated
 - If it is allocated, to which page of which process (or processes)

❑ The operating system maintains a copy of the page table for each process

Outline

- ❑ Background
- ❑ Contiguous Memory Allocation
- ❑ **Paging**
 - Basic Method
 - **Hardware Support**
 - Protection
 - Shared Pages
- ❑ Structure of the Page Table
- ❑ Swapping
- ❑ Example: Intel 32- and 64-bit Architecture
- ❑ Example: ARMv8 Architecture

Page Table

❑ Page tables are per-process data structures

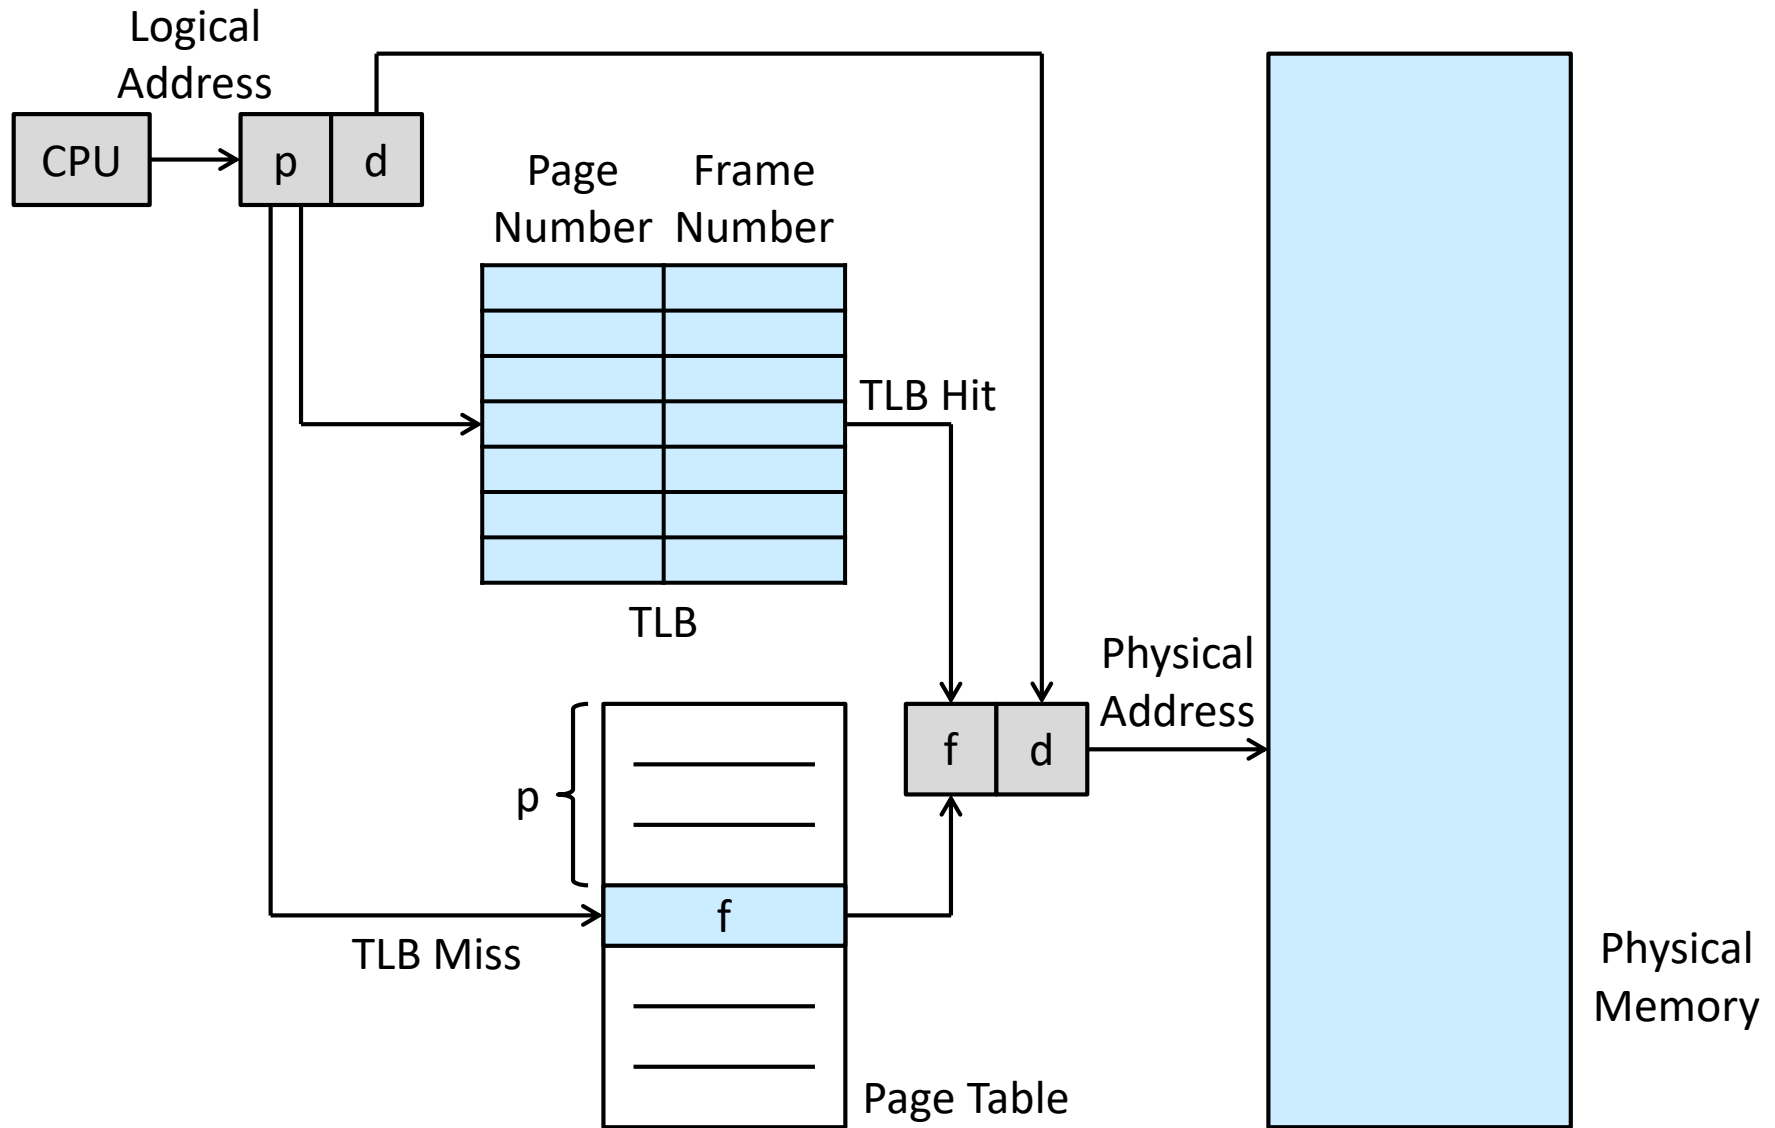
❑ Ways of implementation

- A set of dedicated high-speed hardware registers
 - Make the page-address translation very efficient
 - Increase context-switch time, as each one of these registers must be exchanged during a context switch
- A page-table base register (PTBR) stored in the process control block points to the page table stored in main memory
 - Change only the one register when changing page tables
 - Reduce context-switch time

Translation Look-Aside Buffer (1/3)

- ❑ Storing the page table in main memory yields faster context switches but results in slower memory access times
 - Two memory accesses are needed to access data (one for the page-table entry and one for the actual data)
- ❑ Standard solution: translation look-aside buffer (TLB), also called associative memory
 - A special, small, fast-lookup hardware cache
 - Each TLB entry consisting of two parts: a key (or tag) and a value
 - Typically number of TLB entries: 32 to 1,024 entries
- ❑ Systems have evolved from having no TLBs to having multiple levels of TLBs
 - Just as they have multiple levels of caches

Translation Look-Aside Buffer (2/3)



Translation Look-Aside Buffer (3/3)

- ❑ When a logical address is generated by the CPU, the MMU first checks if its page number is present in the TLB
 - If the page number is found, its frame number is used to access memory
 - Part of the instruction pipeline within the CPU, adding no performance penalty compared with a system that does not implement paging
 - If the page number is not in the TLB (known as a **TLB miss**), a memory reference to the page table must be made
 - Also add the page number and frame number to the TLB, so that they will be found quickly on the next reference
- ❑ If the TLB is full, an existing entry is selected for replacement
 - Examples: least recently used (LRU), round-robin, random, etc.
 - Some CPUs allow the OS to participate in LRU entry replacement
 - Some TLBs allow certain entries to be **wired down**
 - They (e.g., key kernel code) cannot be removed from the TLB

Address-Space Identifier

- ❑ Some TLBs store address-space identifiers (ASIDs) in each TLB entry
 - Identify each process
 - Provide address-space protection for that process
 - Allow the TLB to contain entries for several different processes simultaneously
- ❑ Without ASIDs, every time a new page table is selected (e.g., context switch), the TLB must be flushed (or erased)

Effective Memory-Access Time

❑ Hit ratio

- Percentage of times that the page number of interest is found in the TLB

❑ Example: 10 nanoseconds to access memory

- If the page number is in the TLB
 - A mapped-memory access takes 10 nanoseconds
- If the page number is not in the TLB
 - Access memory for the page table and frame number (10 nanoseconds)
 - Access memory for the desired byte (10 nanoseconds)
 - A total of 20 nanoseconds
- Effective memory-access time
 - 80-percent hit ratio: $0.8 * 10 + 0.2 * 20 = 12$ (nanoseconds)
 - 99-percent hit ratio: $0.99 * 10 + 0.01 * 20 = 10.1$ (nanoseconds)

❑ Calculation is more complicated with multiple levels of TLBs

Outline

- ❑ Background
- ❑ Contiguous Memory Allocation
- ❑ **Paging**
 - Basic Method
 - Hardware Support
 - **Protection**
 - Shared Pages
- ❑ Structure of the Page Table
- ❑ Swapping
- ❑ Example: Intel 32- and 64-bit Architecture
- ❑ Example: ARMv8 Architecture

Protection

❑ Memory protection is accomplished by protection bits associated with each frame

- Keep these bits in the page table
- Provide read-only, read-write, execute-only, or combined protection
- Trap illegal attempts to the operating system


❑ Valid-invalid bit

- When this bit is set to valid, the associated page is in the process's logical address space
- When the bit is set to invalid, the associated page is not in the process's logical address space

❑ Page-table length register (PTLR) to check the page-table size


- Indicate the size of the page table
- Check a logical address to verify that it is in the valid range

Valid-Invalid Bit



Frame Number	Valid-Invalid Bit	
0	2	v
1	3	v
2	4	v
3	7	v
4	8	v
5	9	v
6	0	i
7	0	i

Page Table



- What if a process uses only addresses 0 to 10468? A result of the 2-KB page size

Outline

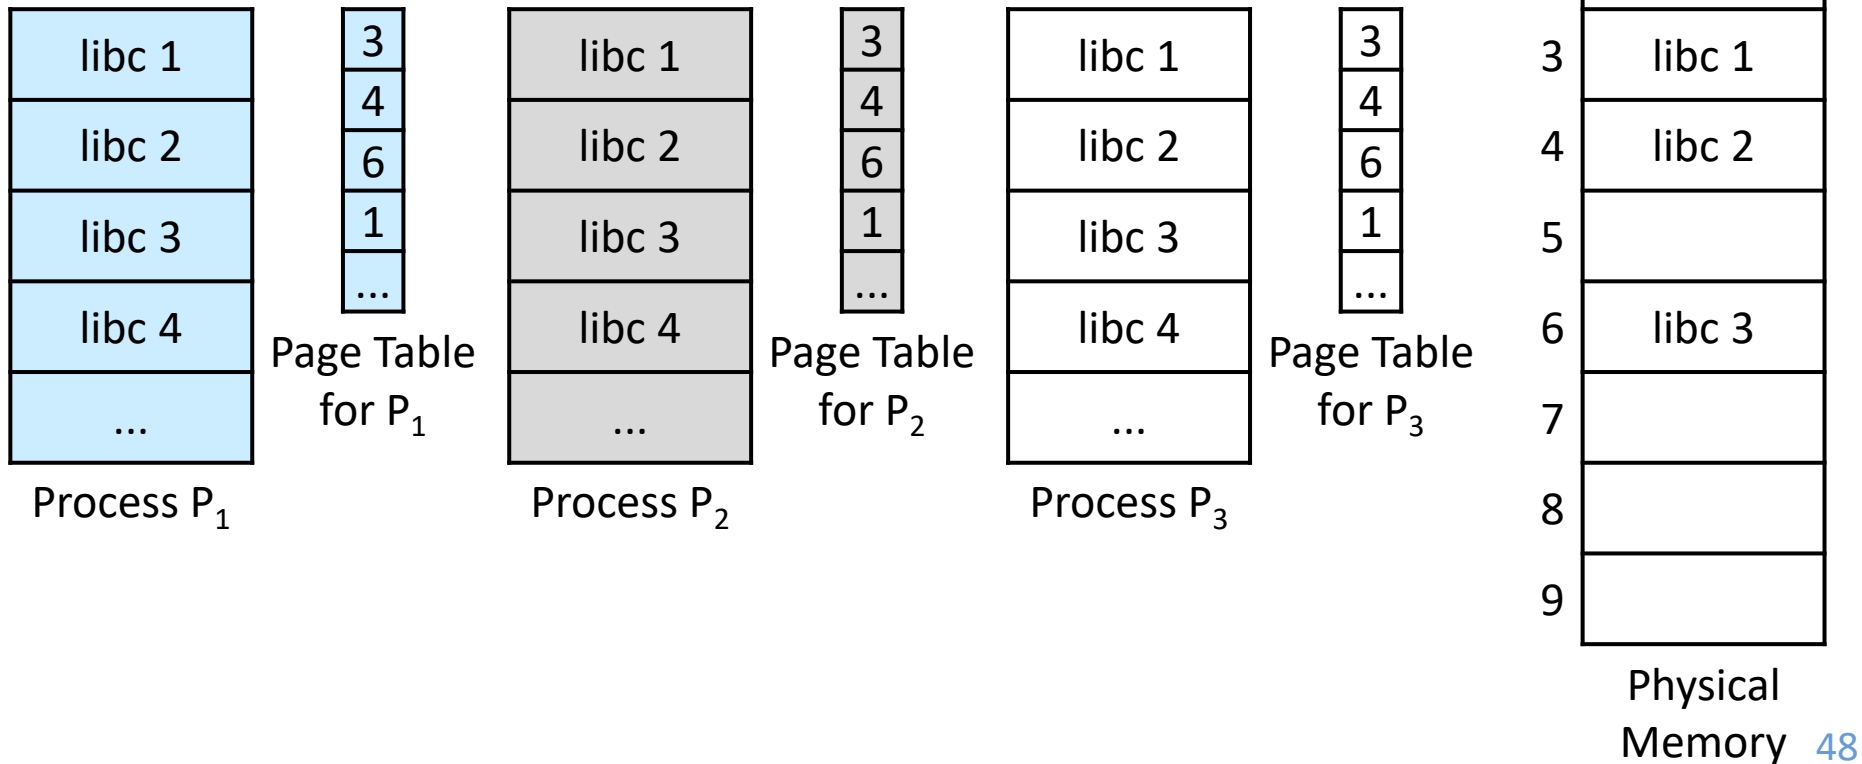
- ❑ Background
- ❑ Contiguous Memory Allocation
- ❑ **Paging**
 - Basic Method
 - Hardware Support
 - Protection
 - **Shared Pages**
- ❑ Structure of the Page Table
- ❑ Swapping
- ❑ Example: Intel 32- and 64-bit Architecture
- ❑ Example: ARMv8 Architecture

Shared Pages

❑ A paging advantage: the possibility of sharing common code

➤ **Reentrant code** is non-self-modifying code

- It never changes during execution
- It can be shared



Outline

- ❑ Background
- ❑ Contiguous Memory Allocation
- ❑ Paging
- ❑ **Structure of the Page Table**
 - **Hierarchical Paging**
 - Hashed Page Tables
 - Inverted Page Tables
 - Oracle SPARC Solaris
- ❑ Swapping
- ❑ Example: Intel 32- and 64-bit Architecture
- ❑ Example: ARMv8 Architecture

Size of Page Table

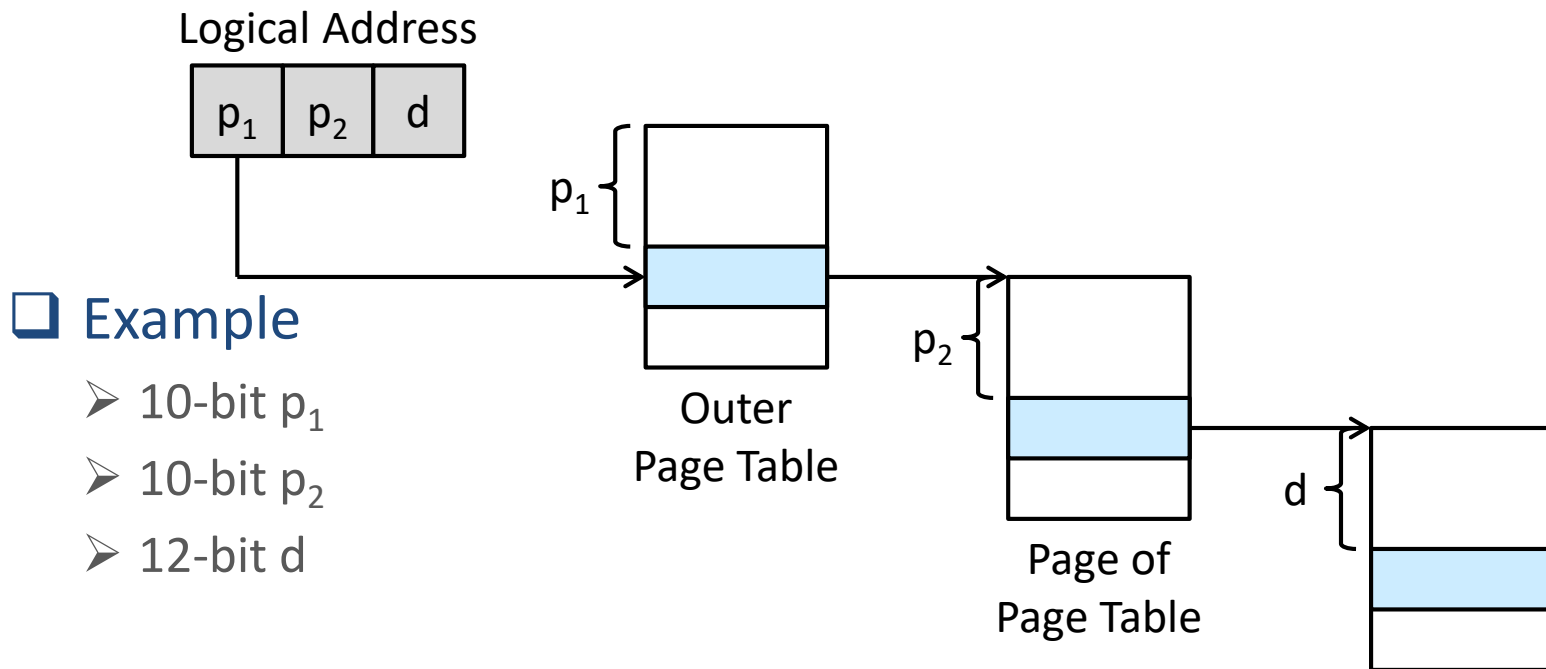
❑ The page table itself may become excessively large

- Most modern computer systems support a large logical address space (2^{32} to 2^{64} bytes)
- If the page size 4 KB (2^{12} bytes), then a page table consists of 2^{20} entries
 - $2^{20} = 2^{32}/2^{12}$
- If each entry consists of 4 bytes, then each process needs up to 4 MB of physical address space for the page table alone
 - We do not want to allocate the page table contiguously in main memory
 - One simple solution is to divide the page table into smaller pieces

Two-Level Paging

❑ Divide the page number and page the page table


❑ Forward-mapped page table



❑ Example

- 10-bit p_1
- 10-bit p_2
- 12-bit d

Two-Level Paging



Multi-Level Paging

- ❑ For a system with a 64-bit logical address space, a two-level paging scheme is no longer appropriate
- ❑ Example
 - If the page size 4 KB (2^{12} bytes), then a page table consists of 2^{52} entries
 - If we use a two-level paging scheme, then the inner page tables can conveniently be one page long, i.e., containing 2^{10} 4-byte entries
 - $p_1, p_2, d = 42, 10, 12$
 - $p_1, p_2, p_3, d = 32, 10, 10, 12$
 - ...
- ❑ The 64-bit UltraSPARC would require seven levels of paging
 - A prohibitive number of memory accesses

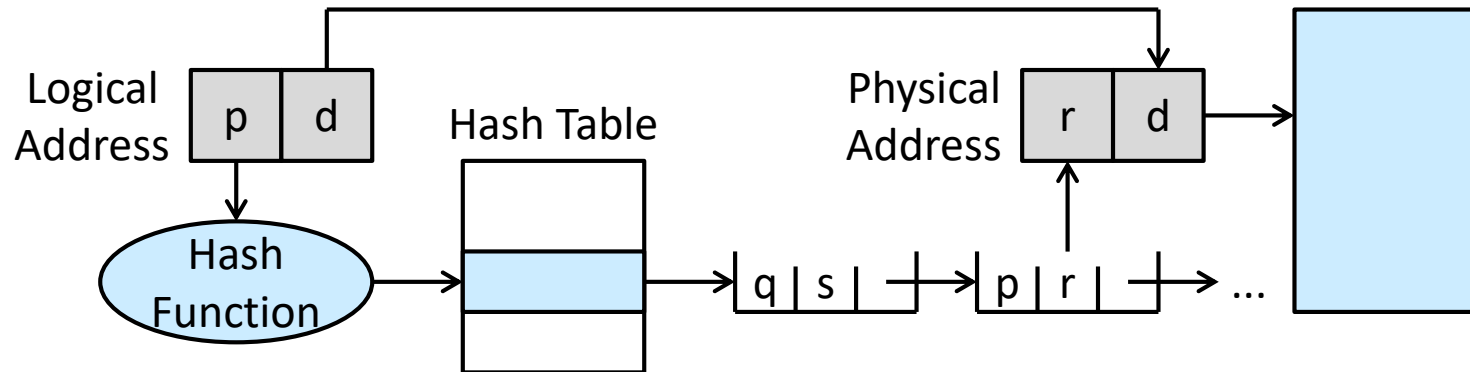
Outline

- ❑ Background
- ❑ Contiguous Memory Allocation
- ❑ Paging
- ❑ **Structure of the Page Table**
 - Hierarchical Paging
 - **Hashed Page Tables**
 - Inverted Page Tables
 - Oracle SPARC Solaris
- ❑ Swapping
- ❑ Example: Intel 32- and 64-bit Architecture
- ❑ Example: ARMv8 Architecture

Hashed Page Tables

- ❑ Each entry in the hash table contains a linked list of elements
- ❑ Each element consists of three fields

- The virtual page number
- The value of the mapped page frame
- A pointer to the next element in the linked list



❑ Variation: clustered page tables

- Similar to hashed page tables except that each entry in the hash table refers to several pages (such as 16) rather than a single page

Outline

- ❑ Background
- ❑ Contiguous Memory Allocation
- ❑ Paging
- ❑ **Structure of the Page Table**
 - Hierarchical Paging
 - Hashed Page Tables
 - **Inverted Page Tables**
 - Oracle SPARC Solaris
- ❑ Swapping
- ❑ Example: Intel 32- and 64-bit Architecture
- ❑ Example: ARMv8 Architecture

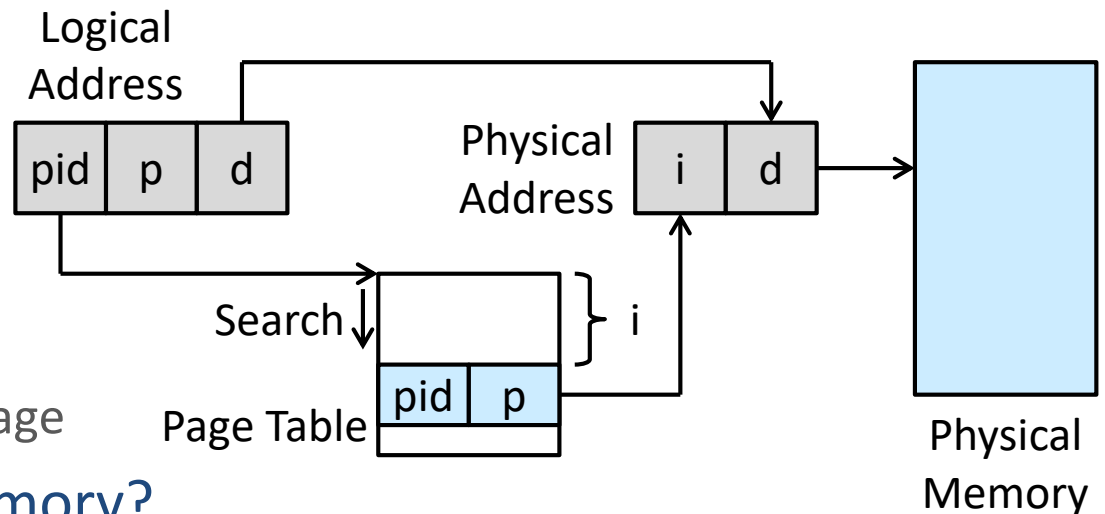
Inverted Page Tables

❑ One entry for each real page (or frame) of memory

- Decrease the amount of memory needed to store each page table
- Increase the amount of time needed to search the table
 - Use a hash table to limit the search to one or a few page-table entries
 - Use a TLB to accelerate

❑ Each entry consists of

- The virtual address of the page stored in that real memory location
- Information about the process that owns the page



❑ How about shared memory?

- Only one mapping of a virtual address to the shared physical address may occur at any given time

Outline

- ❑ Background
- ❑ Contiguous Memory Allocation
- ❑ Paging
- ❑ **Structure of the Page Table**
 - Hierarchical Paging
 - Hashed Page Tables
 - Inverted Page Tables
 - **Oracle SPARC Solaris**
- ❑ Swapping
- ❑ Example: Intel 32- and 64-bit Architecture
- ❑ Example: ARMv8 Architecture

Oracle SPARC Solaris


- ❑ **Solaris** on the **SPARC** CPU is a fully 64-bit operating system
 - Not use all of its physical memory by keeping multi-level page tables
- ❑ **Two hashed page tables**
 - One for the kernel and one for all user processes
 - Each table maps memory addresses from virtual to physical memory
 - Each entry represents a contiguous area of mapped virtual memory
 - More efficient than having a separate hash-table entry for each page
 - Each entry has a base address and a span indicating the number of pages the entry represents
- ❑ **Searching through a hash table would take too long**
 - A TLB holds translation table entries (TTEs)
 - A cache of these TTEs resides in a translation storage buffer (TSB)
 - **TLB walk**: search TLB and, if needed, TSB for a TTE

Outline

- ❑ Background
- ❑ Contiguous Memory Allocation
- ❑ Paging
- ❑ Structure of the Page Table
- ❑ **Swapping**
 - Standard Swapping
 - Swapping with Paging
 - Swapping on Mobile Systems
- ❑ Example: Intel 32- and 64-bit Architecture
- ❑ Example: ARMv8 Architecture

Swapping

- ❑ A process, or a portion of a process, can be
 - **Swapped** temporarily out of memory to a **backing store**
 - Brought back into memory for continued execution
- ❑ Swapping makes it possible for the total physical address space of all processes to exceed the real physical memory

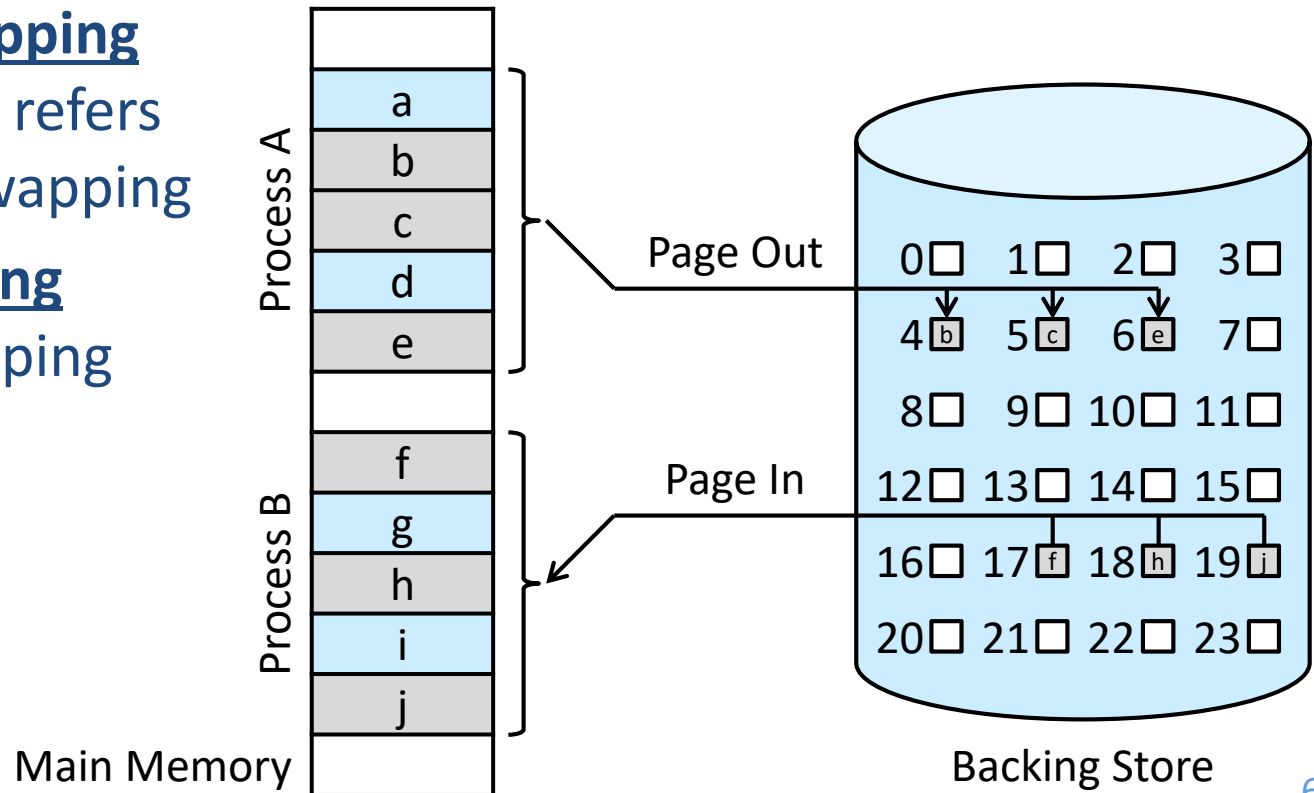


Standard Swapping

- ❑ Standard swapping involves moving entire processes between main memory and a backing store
 - The backing store is commonly fast secondary storage
 - When a process is swapped out, its data structures must be written to the backing store
 - For a multithreaded process, all per-thread data structures must be swapped as well
- ❑ Standard swapping allows physical memory to be oversubscribed
 - Idle or mostly idle processes are good candidates for swapping
 - If a process that has been swapped out becomes active, it must then be swapped back in

Swapping with Paging

- ❑ Most systems now (Linux and Windows) use a variation of swapping in which pages of a process can be swapped
 - Still allows physical memory to be oversubscribed
 - Does not incur the cost of swapping entire processes
- ❑ The term **swapping** now generally refers to standard swapping
- ❑ The term **paging** refers to swapping with paging
- ❑ Chapter 10



Swapping on Mobile Systems

- ❑ Mobile systems typically do not support swapping in any form
 - Generally use flash memory rather than more spacious hard disks for nonvolatile storage
 - Small amount of space
 - Limited # of writes that flash memory can tolerate before becoming unreliable
 - Poor throughput between main memory and flash memory
- ❑ When free memory falls below a certain threshold
 - iOS asks applications to voluntarily relinquish allocated memory
 - Read-only data are removed from main memory
 - Data that have been modified are never removed
 - Any application that fails to free up sufficient memory may be terminated
 - Android adopts a strategy similar to that used by iOS
 - Before terminating a process, Android writes its application state to flash memory

Outline

- ❑ Background
- ❑ Contiguous Memory Allocation
- ❑ Paging
- ❑ Structure of the Page Table
- ❑ Swapping
- ❑ **Example: Intel 32- and 64-bit Architecture**
- ❑ Example: ARMv8 Architecture

Intel 32- and 64-bit Architecture

❑ History

- 16-bit chips Intel 8086 and Intel 8088
- A series of 32-bit chips, IA-32, included the family of 32-bit Pentium processors
- A series of 64-bit chips based on x86-64 architecture

❑ Most current/popular PC operating systems run on Intel chips

- However, the dominance has not spread to mobile systems, where the ARM architecture currently enjoys considerable success

IA-32 Architecture

❑ The segmentation and paging units form the equivalent of the memory-management unit (MMU)

- The CPU generates logical addresses, which are given to the segmentation unit
- The segmentation unit produces a linear address for each logical address, which is given to the paging unit
- The paging unit generates the physical address in main memory

❑ IA-32 segmentation

- The logical address space of a process is divided into two partitions
 - The first partition consists of segments that are private to that process
 - The second partition consists of segments that are shared among all processes

❑ IA-32 paging

- A two-level paging scheme where $p_1, p_2, d = 10, 10, 12$

x86-64

❑ Intel adopted AMD's x86-64 architecture

- Historically, AMD had often developed chips based on Intel's architecture

❑ Support for a 64-bit address space yields 2^{64} bytes of addressable memory

❑ The x86-64 architecture currently provides a 48-bit virtual address

- Support for 52-bit physical addresses with page address extension
- Support for page sizes of 4 KB, 2 MB, or 1 GB
- Four levels of paging hierarchy

Outline

- ❑ Background
- ❑ Contiguous Memory Allocation
- ❑ Paging
- ❑ Structure of the Page Table
- ❑ Swapping
- ❑ Example: Intel 32- and 64-bit Architecture
- ❑ **Example: ARMv8 Architecture**

ARMv8 Architecture

- ❑ In addition to mobile devices, ARM also provides architecture designs for real-time embedded systems
- ❑ The ARMv8 has three different translation granules

➤ Each translation granule provides different page sizes, as well as larger sections of contiguous memory, known as regions

Translation Granule Size	Page Size	Region Size	Paging
4 KB	4 KB	2 MB, 1 GB	Up to 4 Levels
16 KB	16 KB	32 MB	Up to 4 Levels
64 KB	64 KB	512 MB	Up to 3 Levels

- ❑ The ARM architecture also supports two levels of TLBs

➤ At the inner level are two micro TLBs

- One for data and one for instructions; ASIDs supported

➤ At the outer level is a single main TLB

➤ Address translation begins at the micro-TLB level

Objectives

- ❑ Explain the difference between a logical and a physical address and the role of the memory management unit (MMU) in translating addresses
- ❑ Apply first-, best-, and worst-fit strategies for allocating memory contiguously
- ❑ Explain the distinction between internal and external fragmentation
- ❑ Translate logical to physical addresses in a paging system that includes a translation look-aside buffer (TLB)
- ❑ Describe hierarchical paging, hashed paging, and inverted page tables

Q&A

Exercises

□ 9.4

- Consider a logical address space of 64 pages of 1,024 words each, mapped onto a physical memory of 32 frames
 - How many bits are there in the logical address?
 - How many bits are there in the physical address?

□ 9.10

- Consider a computer system with a 32-bit logical address and 4-KB page size
- The system supports up to 512 MB of physical memory
- How many entries are there in each of the following?
 - A conventional, single-level page table
 - An inverted page table