

Operating Systems

[15. File-System Internals]

Chung-Wei Lin

cwlin@csie.ntu.edu.tw

CSIE Department

National Taiwan University

Objectives


- ❑ Delve into the details of file systems and their implementation
- ❑ Explore booting and file sharing
- ❑ Describe remote file systems, using NFS as an example

Outline

- ❑ **File Systems**
- ❑ File-System Mounting
- ❑ Partitions and Mounting
- ❑ File Sharing
- ❑ Virtual File Systems
- ❑ Remote File Systems
- ❑ Consistency Semantics
- ❑ NFS

File Systems (1/2)

- ❑ A general-purpose computer system can have multiple storage devices
- ❑ A device can be sliced up into partitions, which hold volumes, which in turn hold file systems
 - Depending on the volume manager, a volume may span multiple partitions as well



File Systems (2/2)

- ❑ Computer systems may also have varying numbers of file systems, and the file systems may be of varying types

➤ Example: Solaris

- tmpfs: a "temporary" file system that is created in volatile main memory
- objfs: a "virtual" file system that gives debuggers access to kernel symbols
- ctfs: a virtual file system that maintains "contract" information to manage which processes start when booting and must continue during operation
- lofs: a "loop back" file system that allows one file system to be accessed in place of another one
- procfs: a virtual file system that presents information on all processes
- ufs, zfs: general-purpose file systems

/	ufs
/devices	devfs
/dev	dev
/system/contract	ctfs
/proc	proc
/etc/mnttab	mntfs
/etc/svc/volatile	tmpfs
/system/object	objfs
/lib/libc.so.1	lofs
/dev/fd	fd
/var	ufs
/tmp	tmpfs
/var/run	tmpfs
/opt	ufs
/zpbge	zfs
/zpbge/backup	zfs
/export/home	zfs
/var/mail	zfs
/var/spool/mqueue	zfs
/zpbg	zfs
/zpbg/zones	zfs

Outline

- ❑ File Systems
- ❑ **File-System Mounting**
- ❑ Partitions and Mounting
- ❑ File Sharing
- ❑ Virtual File Systems
- ❑ Remote File Systems
- ❑ Consistency Semantics
- ❑ NFS

File-System Mounting (1/2)

- ❑ A file system must be mounted before it can be available to processes on the system
 - The OS is given the device name and the mount point, the location within the file structure where the file system is to be attached
 - Some require that a file-system type be provided
 - Others inspect the structures and determine the file-system type
 - The OS verifies that the device contains a valid file system
 - Ask the device driver to read the device directory
 - Verify that the directory has the expected format
 - The OS notes in its directory structure that a file system is mounted at the specified mount point
 - Enable the operating system to traverse its directory structure and switch among file systems

File-System Mounting (2/2)


❑ Systems impose semantics to clarify functionality

➤ Example 1


- Disallow a mount over a directory that contains files
- Obscure the directory's existing files until the file system is unmounted

➤ Example 2


- Allow multiple mounts (at different mount points) per file system
- Only allow one mount per file system



Existing System



Unmounted Volume



Volume mounted at /users

Outline

- ❑ File Systems
- ❑ File-System Mounting
- ❑ **Partitions and Mounting**
- ❑ File Sharing
- ❑ Virtual File Systems
- ❑ Remote File Systems
- ❑ Consistency Semantics
- ❑ NFS

Partitions and Mounting (1/2)

- ❑ Each partition can be either raw (having no file system) or cooked (having a file system)
 - Raw disk is used where no file system is appropriate
- ❑ If a partition contains a file system that is bootable, then the partition also needs boot information
 - The information has its own format because the system does not have the file-system code loaded at boot time
 - Usually a sequential series of blocks loaded as an image into memory
 - The image, the bootstrap loader, knows enough about the file-system structure to be able to find and load the kernel and start executing
 - Many systems can be dual-booted, allowing us to install multiple operating systems on a single system

Partitions and Mounting (2/2)

❑ The root partition selected by the boot loader is mounted at boot time

- Contain the operating-system kernel and sometimes other system files

❑ Mount operation

- The operating system verifies that the device contains a valid file system
 - Ask the device driver to read the device directory
 - Verify that the directory has the expected format
- If the format is invalid, the partition must have its consistency checked and possibly corrected
- The operating system notes in its in-memory mount table that a file system is mounted, along with the type of the file system

Outline

- ❑ File Systems
- ❑ File-System Mounting
- ❑ Partitions and Mounting
- ❑ **File Sharing**
 - Multiple Users
- ❑ Virtual File Systems
- ❑ Remote File Systems
- ❑ Consistency Semantics
- ❑ NFS

File Sharing

- ❑ To implement sharing and protection, the system must maintain more file and directory attributes
 - Most systems have evolved to use the concepts of file (or directory) owner (or user) and group
 - The owner is the user who can change attributes and grant access and who has the most control over the file
 - The group attribute defines a subset of users who can share access to the file
 - The owner and group IDs of a given file (or directory) are stored with the other file attributes
 - When a user requests an operation on a file, the user ID can be compared to determine if the requesting user is the owner of the file
 - Likewise, the group IDs can be compared
 - The result indicates which permissions are applicable
 - The system then applies those permissions to the requested operation

Outline


- ❑ File Systems
- ❑ File-System Mounting
- ❑ Partitions and Mounting
- ❑ File Sharing
- ❑ **Virtual File Systems**
- ❑ Remote File Systems
- ❑ Consistency Semantics
- ❑ NFS

Multiple Types of File Systems

- ❑ An obvious but suboptimal method is to write directory and file routines for each type
- ❑ Most operating systems use object-oriented techniques to simplify, organize, and modularize the implementation
 - They allow very dissimilar file-system types to be implemented within the same structure, including network file systems, such as NFS
 - Users can access files contained within multiple file systems on the local drive or even on file systems available across the network

Virtual File Systems

- ❑ File-system interface based on the **open()**, **read()**, **write()**, and **close()** calls and on file descriptors
- ❑ Virtual file system (VFS) layer
 - Separate file-system-generic operations from their implementation by defining a clean VFS interface
 - Provide a mechanism for uniquely representing a file throughout a network
 - The VFS is based on a file-representation structure, called a vnode
 - vnode contains a numerical designator for a network-wide unique file
 - This networkwide uniqueness is required for support of network file systems



Virtual File Systems: Linux

❑ Four object types

- **inode object**: represents an individual file
- **File object**: represent an open file
- **Superblock object**: represent an entire file system
- **Dentry object**: represent an individual directory entry

❑ For each type, the VFS defines a set of operations

- Every object of one of these types contains a pointer to a function table
- The function table lists the addresses of the actual functions that implement the defined operations for that particular object
- Examples for the file object
 - `int open(...)`: open a file
 - `int close(...)`: close an already-open file
 - `ssize_t read(...)` / `write(...)`: read/write from a file
 - `int mmap(...)`: memory-map a file

Outline

- ❑ File Systems
- ❑ File-System Mounting
- ❑ Partitions and Mounting
- ❑ File Sharing
- ❑ Virtual File Systems
- ❑ **Remote File Systems**
 - The Client-Server Model
 - Distributed Information Systems
 - Failure Models
- ❑ Consistency Semantics
- ❑ NFS

Remote File Systems

❑ Networking allows the sharing of data in the form of files

- The first implemented method is **ftp**
 - It involves manually transferring files between machines via programs
 - **ftp** is used for both anonymous and authenticated access
- The second major method uses a **distributed file system** (DFS)
 - Remote directories are visible from a local machine
 - Its involves a much tighter integration between the machine that is accessing the remote files and the machine providing the files
- The third method is the **World Wide Web**
 - A browser is needed to gain access to the remote files
 - Separate operations (essentially a wrapper for **ftp**) are used to transfer files

Outline

- ❑ File Systems
- ❑ File-System Mounting
- ❑ Partitions and Mounting
- ❑ File Sharing
- ❑ Virtual File Systems
- ❑ **Remote File Systems**
 - **The Client-Server Model**
 - Distributed Information Systems
 - Failure Models
- ❑ Consistency Semantics
- ❑ NFS

The Client-Server Model

- ❑ Remote file systems allow a computer to mount one or more file systems from one or more remote machines
 - The machine containing the files is the server
 - The machine seeking access to the files is the client
- ❑ Security results in many challenges
 - Example: a client can be specified by a network name or other identifiers, but these can be spoofed or imitated
 - In the case of UNIX and its network file system (NFS), authentication takes place via the client networking information, by default
- ❑ File operation requests are sent on behalf of the user across the network to the server via the DFS protocol
 - The server determines if the user has credentials to access the file in the mode requested

Outline

- ❑ File Systems
- ❑ File-System Mounting
- ❑ Partitions and Mounting
- ❑ File Sharing
- ❑ Virtual File Systems
- ❑ **Remote File Systems**
 - The Client-Server Model
 - **Distributed Information Systems**
 - Failure Models
- ❑ Consistency Semantics
- ❑ NFS

Distributed Information Systems

- ❑ Provide unified access to the information needed for remote computing, also known as **distributed naming services**
 - The **domain name system** (DNS) provides host-name-to-network-address translations
 - The **network information service** (NIS) provides user-name, password, user ID, group ID space for a distributed facility
 - In Microsoft's **common Internet file system** (CIFS), network information is used with user authentication to create a network login request
 - **Active directory** as a distributed naming structure
 - **Kerberos** network authentication protocol
 - **Lightweight directory-access protocol** (LDAP) as a secure distributed naming mechanism

Outline

- ❑ File Systems
- ❑ File-System Mounting
- ❑ Partitions and Mounting
- ❑ File Sharing
- ❑ Virtual File Systems
- ❑ **Remote File Systems**
 - The Client-Server Model
 - Distributed Information Systems
 - **Failure Models**
- ❑ Consistency Semantics
- ❑ NFS

Failure Models

❑ Local file systems can fail for a variety of reasons

- Failure of the drive containing the file system
- Corruption of the directory structure or other disk-management information (metadata)
- Disk-controller, cable, or host adapter failure
- User or system-administrator failure

❑ Remote file systems have even more failure modes

- Network or server failure, networking implementation issues

❑ If both server and client maintain state information, then they can seamlessly recover from a failure

- NFS v3 implements a stateless DFS, carrying all the information needed, which is resilient, rather easy to implement, and unsecure
- NFS v4 is made stateful to improve its security, performance, and functionality

Outline

- ❑ File Systems
- ❑ File-System Mounting
- ❑ Partitions and Mounting
- ❑ File Sharing
- ❑ Virtual File Systems
- ❑ Remote File Systems
- ❑ **Consistency Semantics**
 - UNIX Semantics
 - Session Semantics
 - Immutable-Shared-Files Semantics
- ❑ NFS

Consistency Semantics

- ❑ An important criterion for evaluating any file system that supports file sharing
 - Specify how multiple users of a system are to access a shared file simultaneously
 - Specify when modifications of data by one user will be observable by other users
- ❑ These semantics are typically implemented as code with the file system
- ❑ The series of accesses between the **open ()** and **close ()** operations makes up a **file session**

Outline

- ❑ File Systems
- ❑ File-System Mounting
- ❑ Partitions and Mounting
- ❑ File Sharing
- ❑ Virtual File Systems
- ❑ Remote File Systems
- ❑ **Consistency Semantics**
 - **UNIX Semantics**
 - Session Semantics
 - Immutable-Shared-Files Semantics
- ❑ NFS

UNIX Semantics

❑ The UNIX file system uses the following consistency semantics

- Writes to an open file by a user are visible immediately to other users who have this file open
- One mode of sharing allows users to share the pointer of current location into the file
 - The advancing of the pointer by one user affects all sharing users
 - A file has a single image that interleaves all accesses, regardless of their origin

❑ A file is associated with a single physical image that is accessed as an exclusive resource

- Contention for this single image causes delays in user processes

Outline

- ❑ File Systems
- ❑ File-System Mounting
- ❑ Partitions and Mounting
- ❑ File Sharing
- ❑ Virtual File Systems
- ❑ Remote File Systems
- ❑ **Consistency Semantics**
 - UNIX Semantics
 - **Session Semantics**
 - Immutable-Shared-Files Semantics
- ❑ NFS

Session Semantics

- ❑ The Andrew file system (OpenAFS) uses the following consistency semantics
 - Writes to an open file by a user are not visible immediately to other users that have the same file open
 - Once a file is closed, the changes made to it are visible only in sessions starting later
 - Already open instances of the file do not reflect these changes
- ❑ A file may be associated temporarily with several images at the same time
 - Multiple users are allowed to perform both read and write accesses concurrently on their images of the file, without delay

Outline

- ❑ File Systems
- ❑ File-System Mounting
- ❑ Partitions and Mounting
- ❑ File Sharing
- ❑ Virtual File Systems
- ❑ Remote File Systems
- ❑ **Consistency Semantics**
 - UNIX Semantics
 - Session Semantics
 - **Immutable-Shared-Files Semantics**
- ❑ NFS

Immutable-Shared-Files Semantics

- ❑ Once a file is declared as shared by its creator, it cannot be modified
 - An immutable file has two key properties
 - Its name may not be reused
 - Its contents may not be altered
 - The name of an immutable file signifies that the contents of the file are fixed
- ❑ The implementation of these semantics in a distributed system is simple
 - The sharing is disciplined (read-only)

Outline


- ❑ File Systems
- ❑ File-System Mounting
- ❑ Partitions and Mounting
- ❑ File Sharing
- ❑ Virtual File Systems
- ❑ Remote File Systems
- ❑ Consistency Semantics
- ❑ **NFS (Network File System)**
 - The Mount Protocol
 - The NFS Protocol
 - Path-Name Translation
 - Remote Operations

NFS (1/2)

- ❑ An implementation and a specification of a software system for accessing remote files across LANs (or WANs)
 - NFS views a set of interconnected workstations as a set of independent machines with independent file systems
 - Sharing is based on a client-server relationship

❑ Mounting

- Diskless workstations can even mount their own roots from servers
- Cascading mounts are also permitted in some NFS implementations



NFS (2/2)

- ❑ NFS is to operate in a heterogeneous environment of different machines, operating systems, and network architectures
 - The NFS specification is independent of these media
 - This independence is achieved through the remote procedure call (RPC) primitives built on top of an external data representation (XDR) protocol
- ❑ The NFS specification distinguishes between
 - The services provided by a mount mechanism
 - A mount protocol
 - The actual remote-file-access services
 - A protocol for remote file accesses, the NFS protocol

Outline

- ❑ File Systems
- ❑ File-System Mounting
- ❑ Partitions and Mounting
- ❑ File Sharing
- ❑ Virtual File Systems
- ❑ Remote File Systems
- ❑ Consistency Semantics
- ❑ **NFS**
 - **The Mount Protocol**
 - The NFS Protocol
 - Path-Name Translation
 - Remote Operations

The Mount Protocol

- ❑ Establish the initial logical connection between a server and a client
 - A mount operation includes
 - The name of the remote directory to be mounted
 - The name of the server machine storing it
 - The mount operation changes only the user's view and does not affect the server side
 - The mount request is mapped to the corresponding RPC and is forwarded to the mount server running on the specific server machine
 - The server maintains an export list that specifies local file systems that it exports for mounting, along with names of permitted machines
 - If request conforms to the export list, the server returns to the client a file handle for accesses to files within the mounted file system
 - The file handle contains all the information that the server needs to distinguish an individual file it stores

Outline

- ❑ File Systems
- ❑ File-System Mounting
- ❑ Partitions and Mounting
- ❑ File Sharing
- ❑ Virtual File Systems
- ❑ Remote File Systems
- ❑ Consistency Semantics
- ❑ **NFS**
 - The Mount Protocol
 - **The NFS Protocol**
 - Path-Name Translation
 - Remote Operations

The NFS Protocol

❑ Provide a set of RPCs for remote file operations

- Search for a file within a directory
- Read a set of directory entries
- Manipulate links and directories
- Access file attributes
- Read and write files

❑ These procedures can be invoked only after a file handle for the remotely mounted directory has been established

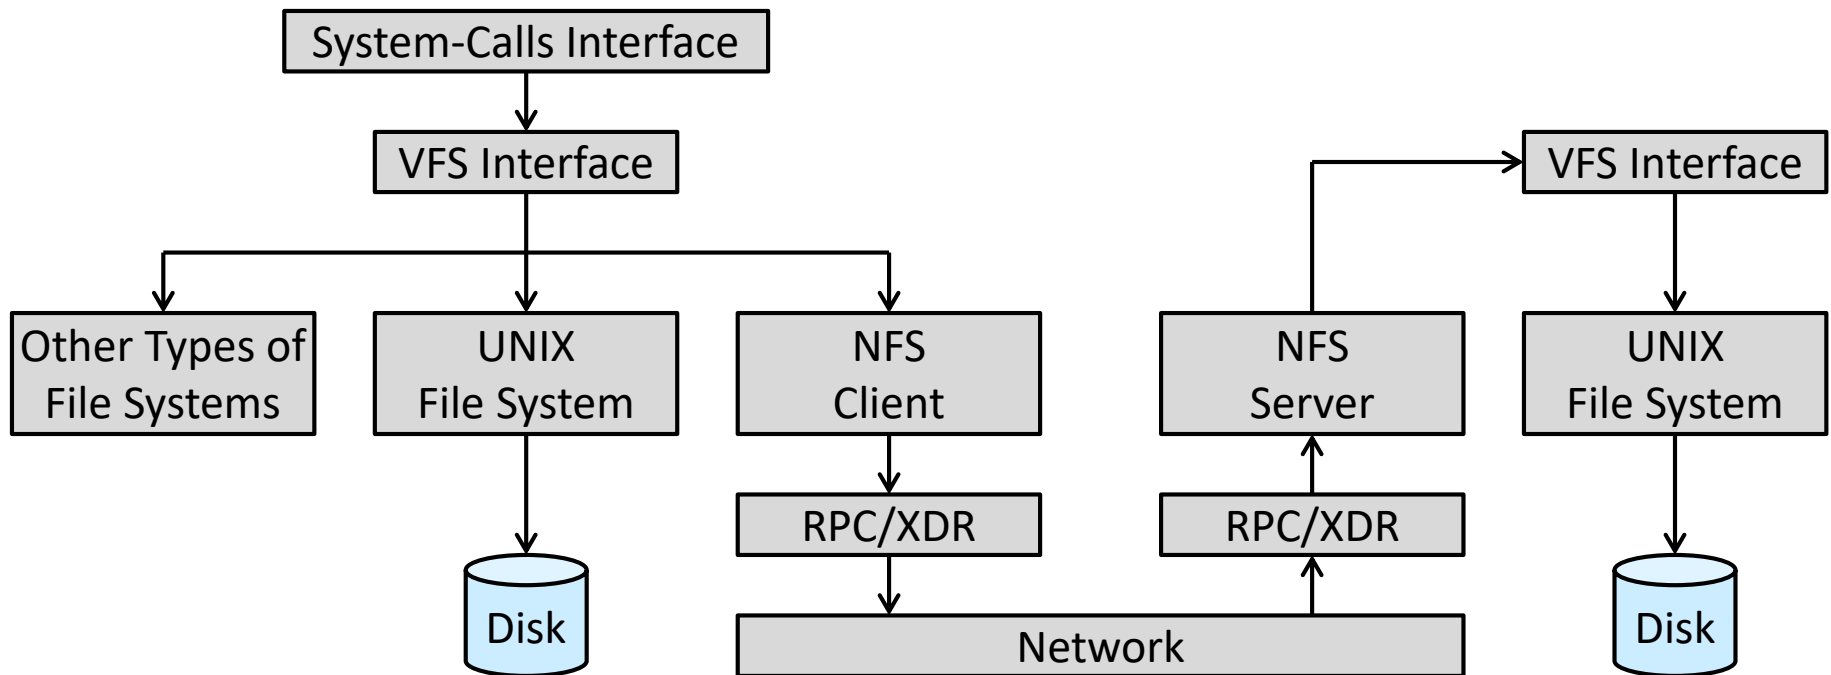
❑ NFS v3 is stateless (NFS v4 is stateful)

- Servers do not maintain information about clients from one access to another
- Each request has to provide a full set of arguments
- Modified data must be committed to the server's disk before returned

Schematic View of NFS Architecture

❑ NFS is integrated into the operating system via a VFS

➤ How an operation on an already-open remote file is handled



Outline

- ❑ File Systems
- ❑ File-System Mounting
- ❑ Partitions and Mounting
- ❑ File Sharing
- ❑ Virtual File Systems
- ❑ Remote File Systems
- ❑ Consistency Semantics
- ❑ **NFS**
 - The Mount Protocol
 - The NFS Protocol
 - **Path-Name Translation**
 - Remote Operations

Path-Name Translation

- ❑ Break the path name into the names of separate directory entries or components
 - Perform a separate NFS **lookup** call for every pair of component name and directory vnode
- ❑ Speed up references to files with the same initial path name
 - A directory-name-lookup cache on the client side holds the vnodes for remote directory names

Outline

- ❑ File Systems
- ❑ File-System Mounting
- ❑ Partitions and Mounting
- ❑ File Sharing
- ❑ Virtual File Systems
- ❑ Remote File Systems
- ❑ Consistency Semantics
- ❑ **NFS**
 - The Mount Protocol
 - The NFS Protocol
 - Path-Name Translation
 - **Remote Operations**

Remote Operations

- ❑ Almost one-to-one correspondence between the regular UNIX system calls for file operations and the NFS protocol RPCs
 - Except opening and closing files
- ❑ In practice, buffering and caching are employed for performance
 - When a file is opened, the kernel checks with the remote server to determine whether to fetch or revalidate the cached attributes
 - The file-attribute (inode-information) cache is updated whenever new attributes arrive from the server
 - The file-blocks cache is used only if the cached attributes are up to date
- ❑ Both read-ahead and delayed-write techniques are used
 - Clients do not free delayed-write blocks until the server confirms that the data have been written to disk

Objectives

- ❑ Delve into the details of file systems and their implementation
- ❑ Explore booting and file sharing
- ❑ Describe remote file systems, using NFS as an example

Q&A